

CSCI 1470/2470  
Spring 2023

Ritambhara Singh

March 01, 2023  
Wednesday

# Deep Learning






# Roadmap



Machine Learning Concepts

Perceptron



Fully Connected Neural Networks




Convolutional Neural Networks



# New data type: sequences

- Audio



- DNA



- Stock market



- Weather



What is the data property here  
that we could leverage?

# Natural Language

*“language that has developed naturally in use”*

# Natural Language

*“language that has developed naturally in use”*

Compare to *constructed* or *formal* language

- code: `for i in range(50):`
- math:  $52 + 94 = 147$
- logic:  $A \wedge B \rightarrow C$  (if A and B, then C)

# Natural Language

In this class: **sequence of words**

*“They went to the grocery store and bought bread,  
peanut butter, and jam.”*

# Natural Language: Prediction tasks?

Example of prediction?



Input: X

I do not want sour  
cream in my  
burrito



Function:  $f$



Output: Y

No quiero crema  
agrea en mi  
burrito

# Natural Language: Prediction tasks?

Example of classification?

Input: X

“The story telling was erratic and, at times, slow”

“Loved the diverse cast of this movie”



Function:  $f$



Output: Y

“Good review?”





# Natural Language: Prediction tasks?

Example of prediction?

*“They went to the grocery store and bought... bread?  
milk?  
rock?”*

**Generating artificial sentences:** Here each word is a discrete unit; predicting the next part of the sequence means predicting words

# Language models

Definition: Probability distribution over strings in a language.

Exponentially-many strings means each string has very low probability

Relative probabilities are meaningful:

$P(\text{“}i\text{they went to the store}\text{”}) \gg P(\text{“}i\text{butter dancing rock}\text{”})$

# Language models logic: leverage sentence structure

$P(\text{any sequence})$  is determined by  $P(\text{the words in the sequence})$ .

Said differently, we can represent a sequence as  $w_1, w_2, \dots w_n$ , and

$$P(w_1, w_2, \dots w_n) = P(w_1) * P(w_2|w_1) * P(w_3|w_1, w_2) * \dots P(w_n|w_1 \dots w_{n-1})$$

$$P(\text{"they went to the store"}) = P(\text{"they"}) * P(\text{"went"} | \text{"they"}) * P(\text{"to"} | \text{"they went"}) * \dots$$

*"The probability of a sentence is the product of the probabilities of each word given the previous words"*

This is an application of the **chain rule for probabilities**

# Language models: weird & cool!

Model trained on the King James Bible, Structure and Interpretation of Computer Programs, and some of Eric S. Raymond's writings:

- *The righteous shall inherit the land, and leave it for an inheritance unto the children of Gad according to the number of steps that is linear in  $b$ .*
- *And this I pray, that your love may abound yet more and more like a controlled use of shared memory.*

(King James Programming)

<https://kingjamesprogramming.tumblr.com/>

Any questions?



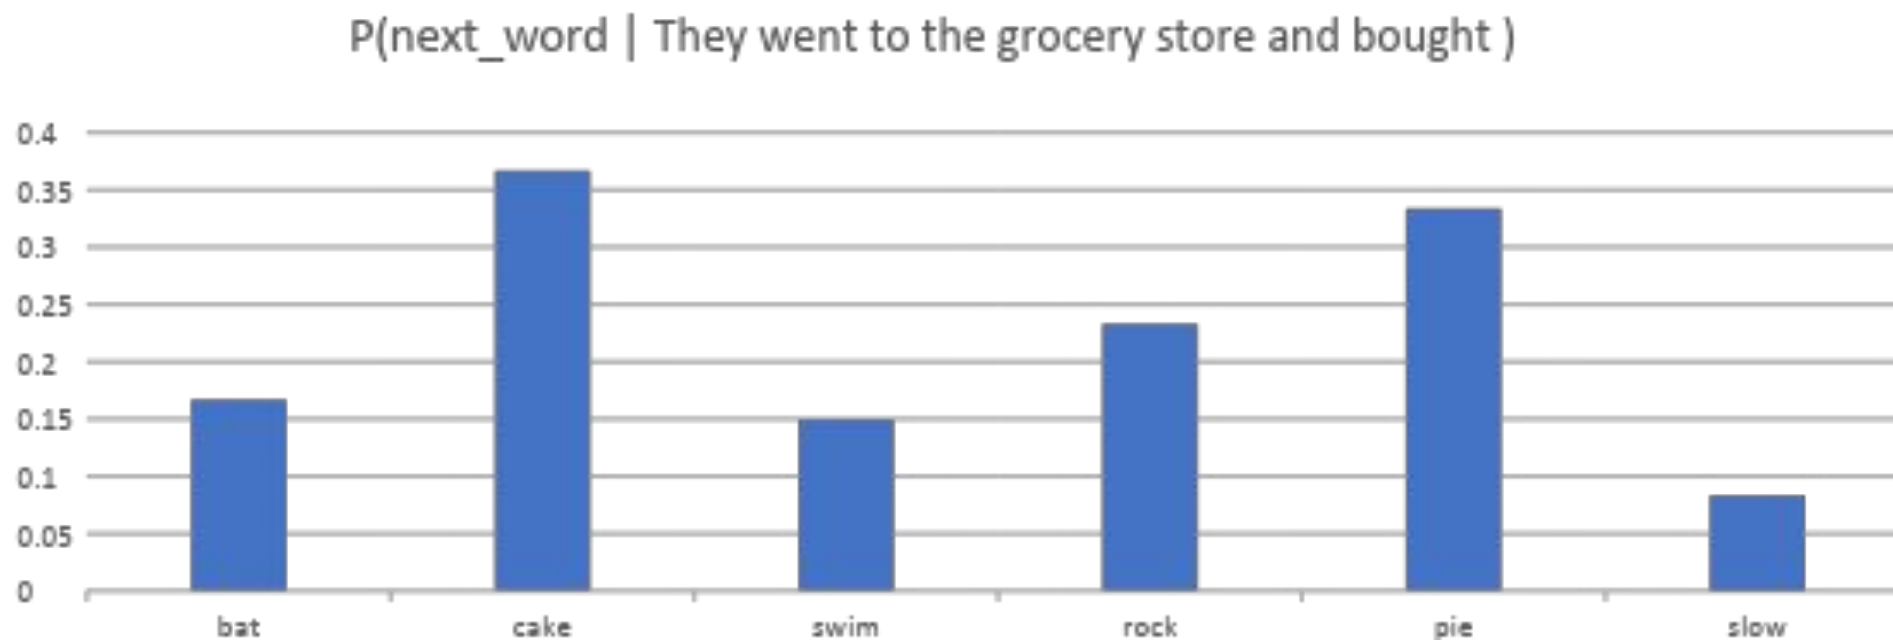


# Language models: the math

But first, how do we represent sentence?

At each step, we look at a probability distribution for what the *next* word might be.

*They went to the grocery store and bought ..*



# Natural language: tokenization

*“They went to the grocery store and bought bread, peanut butter, and jam.”*



[“they”, “went”, “to”, “the”,  
“grocery”, “store”, “and”,  
“bought”, “bread”, “peanut”,  
“butter”, “and”, “jam”]

# Natural language: tokenization

*“They went to the grocery store and bought bread, peanut butter, and jam.”*

- Consistent casing
- Strip punctuation
- One word is one token
- Split on spaces

[“they”, “went”, “to”, “the”,  
“grocery”, “store”, “and”,  
“bought”, “bread”, “peanut”,  
“butter”, “and”, “jam”]

# Aside: Tokenization itself can be challenging...

- A lot easier in English than other languages (e.g. Chinese)
  - Chinese is character-based; words & phrases have different character lengths
  - No spaces




# Language models: the math

How do we know which words to calculate probabilities for?

At each step, we look at a probability distribution for what the *next* word might be.

*They went to the grocery store and bought ..*



# Vocabularies: Defining a finite set of words

Vocabularies: the set of all words “known” to the model

Why?

- We need a finite set of words in order to define a discrete distribution over it.

How?

- Choose a hyperparameter **vocab\_size** for how many words the model should know
- Keep only the **vocab\_size** with most frequent words – replace everything else with “**UNK**”

# Vocabularies: how

- Original sentence:

- *“They galloped to the Ratty for dinner, and ate exactly seventy-three waffle fries and chocolate peamilk.”*

# Vocabularies: how

- Original sentence:

- *“They galloped to the Ratty for dinner, and ate exactly seventy-three waffle fries and chocolate peamilk.”*

- Tokenized:

- [“they”, “galloped”, “to”, “the”, “ratty”, “for”, “dinner”, “and”, “ate”, “exactly”, “seventy-three”, “waffle”, “fries”, “and”, “chocolate”, “peamilk”]



# Vocabularies: how

- Original sentence:

- *“They galloped to the Ratty for dinner, and ate exactly seventy-three waffle fries and chocolate peamilk.”*

- Tokenized:

- [“they”, “galloped”, “to”, “the”, “ratty”, “for”, “dinner”, “and”, “ate”, “exactly”, “seventy-three”, “waffle”, “fries”, “and”, “chocolate”, “peamilk”]

# Vocabularies: how

- Original sentence:

- *“They galloped to the Ratty for dinner, and ate exactly seventy-three waffle fries and chocolate peamilk.”*

- Tokenized:

- [“they”, “galloped”, “to”, “the”, “ratty”, “for”, “dinner”, “and”, “ate”, “exactly”, “seventy-three”, “waffle”, “fries”, “and”, “chocolate”, “peamilk”]

- UNKed:

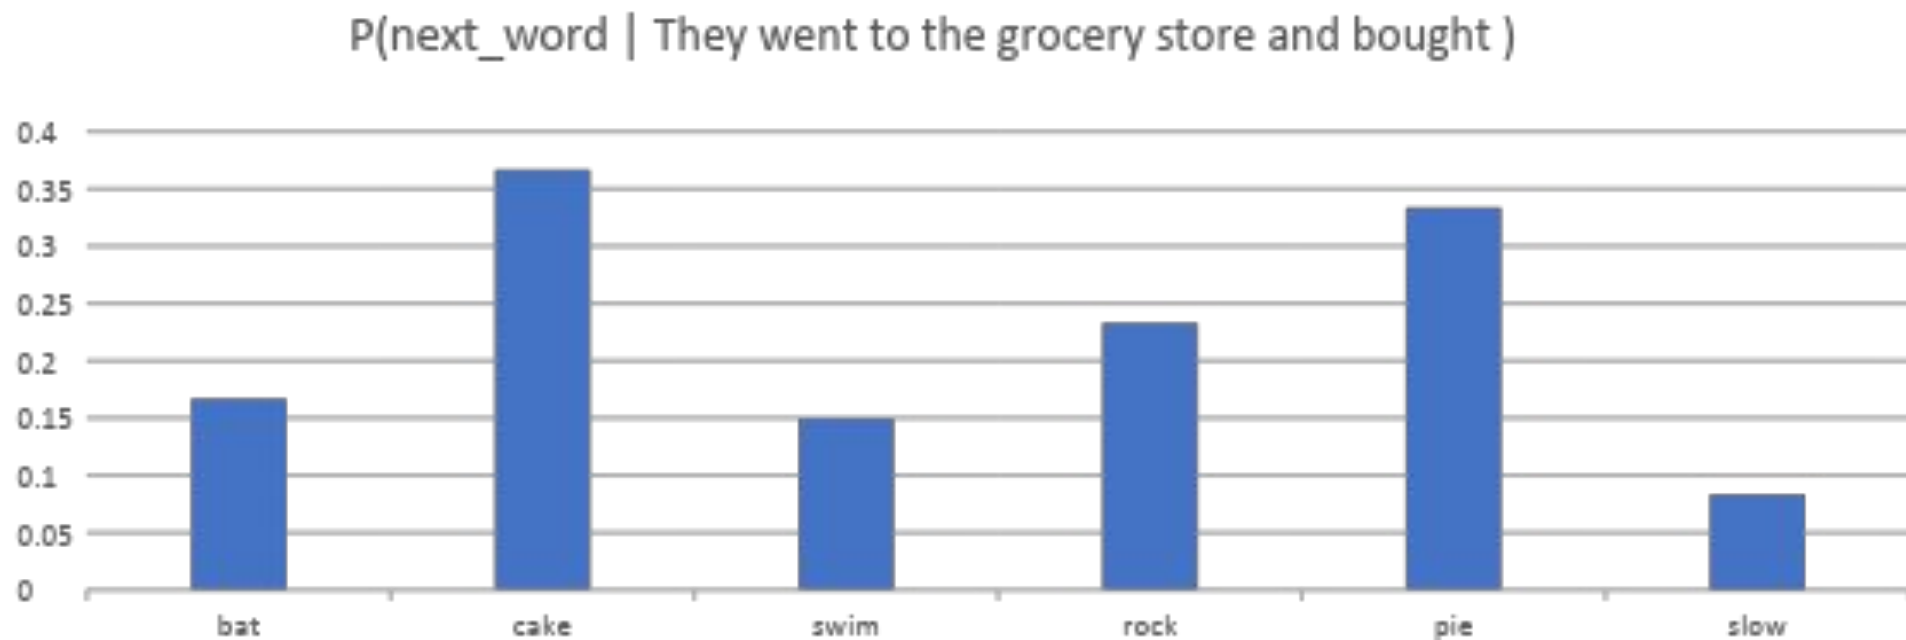
- [“they”, “UNK”, “to”, “the”, “UNK”, “for”, “dinner”, “and”, “ate”, “exactly”, “UNK”, “waffle”, “fries”, “and”, “chocolate”, “UNK”]

# Language models: the math

How to calculate the probability for words in our vocabulary?

At each step, we look at a probability distribution for what the *next* word might be.

*They went to the grocery store and bought ..*



# LM implementation: counting

- Goal: predict next word given a preceding sequence

- $$P(\mathbf{word}_n | word_1, word_2, \dots word_{n-1}) = \frac{Count(word_1, word_2, \dots word_{n-1}, \mathbf{word}_n)}{Count(word_1, word_2, \dots word_{n-1})}$$



# LM implementation: counting

- - Goal: predict next word given a preceding sequence
  - $P(\mathbf{word}_n | word_1, word_2, \dots word_{n-1}) = \frac{Count(word_1, word_2, \dots word_{n-1}, \mathbf{word}_n)}{Count(word_1, word_2, \dots word_{n-1})}$
  - Example task: predict the next word
    - *he danced*

# LM implementation: counting

- Goal: predict next word given a preceding sequence
  - $P(\mathbf{word}_n | word_1, word_2, \dots word_{n-1}) = \frac{Count(word_1, word_2, \dots word_{n-1}, \mathbf{word}_n)}{Count(word_1, word_2, \dots word_{n-1})}$
  - Example task: predict the next word
    - **he danced**
  - Strategy: iterate through all words in vocabulary, and calculate  $\frac{Count(he\ danced\ <word>)}{Count(he\ danced)}$  for each word

# LM implementation: counting

- Our training sentences were:

- “She danced happily”
- “They sang beautifully”
- “He danced energetically”
- “He sang happily”
- “She danced gracefully”

- “He danced \_ \_ \_”

- “He danced *happily*”

Has 0 probability

$$\frac{\text{Count}(\text{he danced } < \text{word} > )}{\text{Count}(\text{he danced})}$$

**Why doesn't this work?**

This strategy depends on having instances of sentence prefixes.

# LM implementation: N-gram counting

Improvement: **N-gram** model – only look at **N** words at a time

# LM implementation: N-gram counting

Improvement: **N-gram** model – only look at **N** words at a time  
(in this case, **bigrams** look at **2** words at a time)

- “*She danced happily*”
- “*They sang beautifully*”
- “*He danced energetically*”
- “*He sang happily*”
- “*She danced gracefully*”

# LM implementation: N-gram counting

Improvement: **N-gram** model – only look at **N** words at a time  
(in this case, **bigrams** look at **2** words at a time)

- “*danced happily*”
- “*sang beautifully*”
- “*danced energetically*”
- “*sang happily*”
- “*danced gracefully*”

“*He danced happily*” now has 1/3 probability!

But what if the answer was “*He danced beautifully*” ?



# LM implementation

Problem: it's impossible for the training set to have *every possible valid* sequence of words!

Let's try to learn a better **numerical** representation

*What is the simplest thing you can think of?*

# LM implementation: Simple approach

- “She danced happily”
- “They sang beautifully”
- “He danced energetically”
- “He sang happily”
- “She danced gracefully”

vocab\_sz

**Any potential issues with this?**

“They danced happily”

⋮	⋮	⋮	⋮
⋮	0	0	0
they	1	0	0
danced	0	1	0
sang	0	0	0
happily	0	0	1
⋮	⋮	⋮	⋮

# LM implementation

Problem: it's impossible for the training set to have *every possible valid* sequence of words!

Can we learn a better numerical representation **which associates *related* words with one another?**

# Embedding matrix

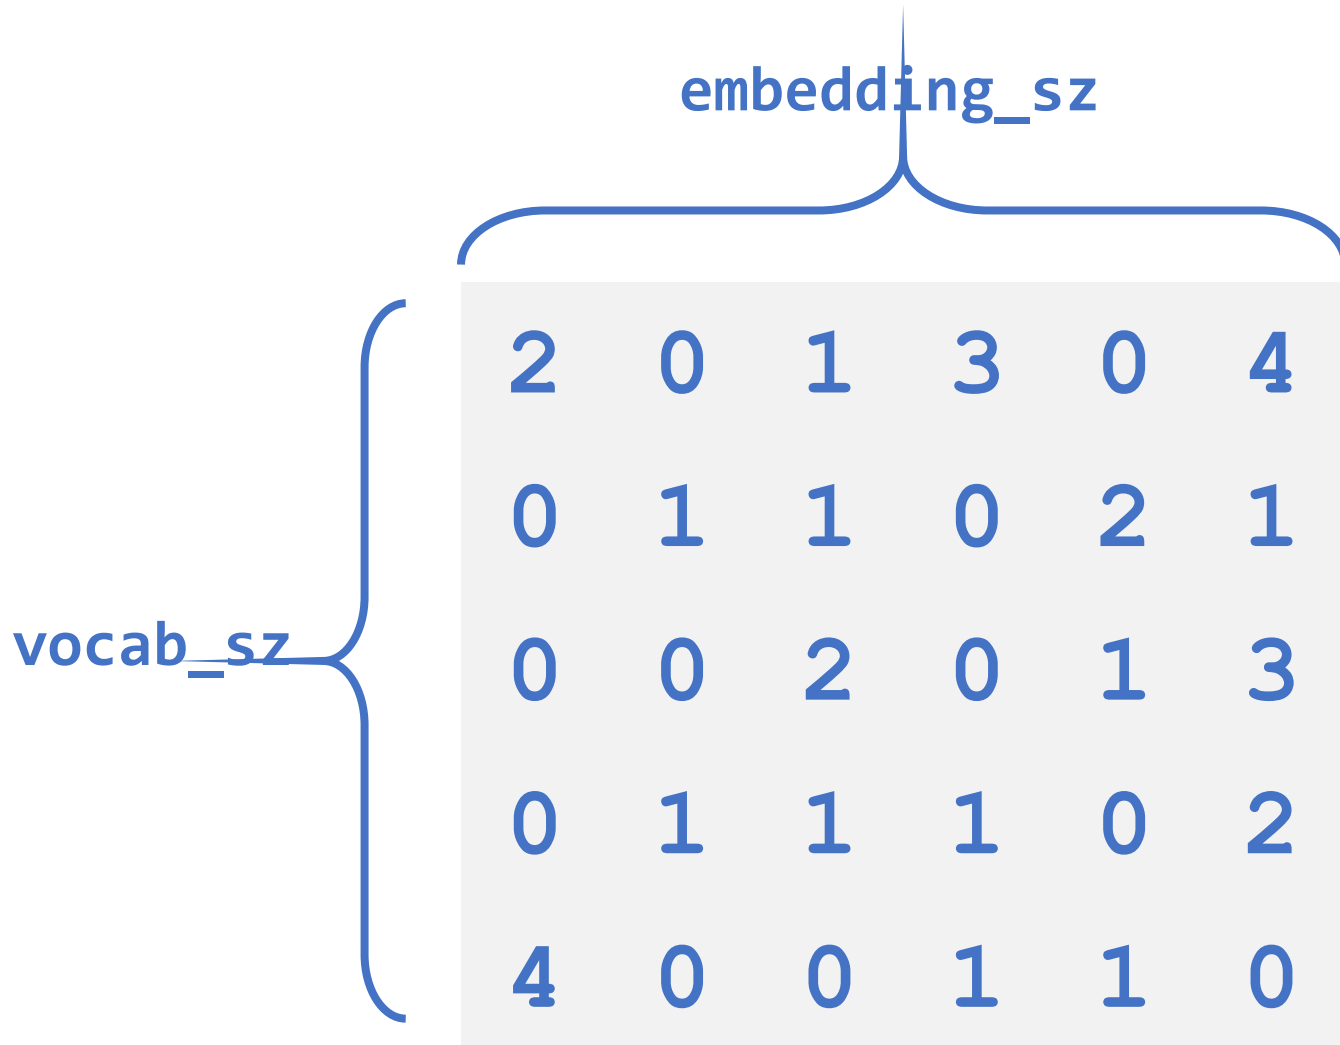


The diagram illustrates an embedding matrix. On the left, a blue bracket groups the words in the first column, with the label `vocab_sz` positioned next to it. The words are listed vertically: *they*, *danced*, *sang*, *happily*, and *gleefully*, with vertical ellipses above and below. To the right of these words is a 5x6 grid of numbers, each in a blue font, representing the embedding vectors. The grid is highlighted with a light gray background.

<i>they</i>	2	0	1	3	0	4
<i>danced</i>	0	1	1	0	2	1
<i>sang</i>	0	0	2	0	1	3
<i>happily</i>	0	1	1	1	0	2
<i>gleefully</i>	4	0	0	1	1	0



# Embedding matrix



The diagram shows a 5x6 matrix of integers. A blue curly brace on the left side, labeled 'vocab\_sz', spans all five rows. A blue curly brace on the top side, labeled 'embedding\_sz', spans all six columns.

2	0	1	3	0	4
0	1	1	0	2	1
0	0	2	0	1	3
0	1	1	1	0	2
4	0	0	1	1	0

- 2d matrix: **vocab\_sz**  
**x embedding\_sz**

# Embedding matrix



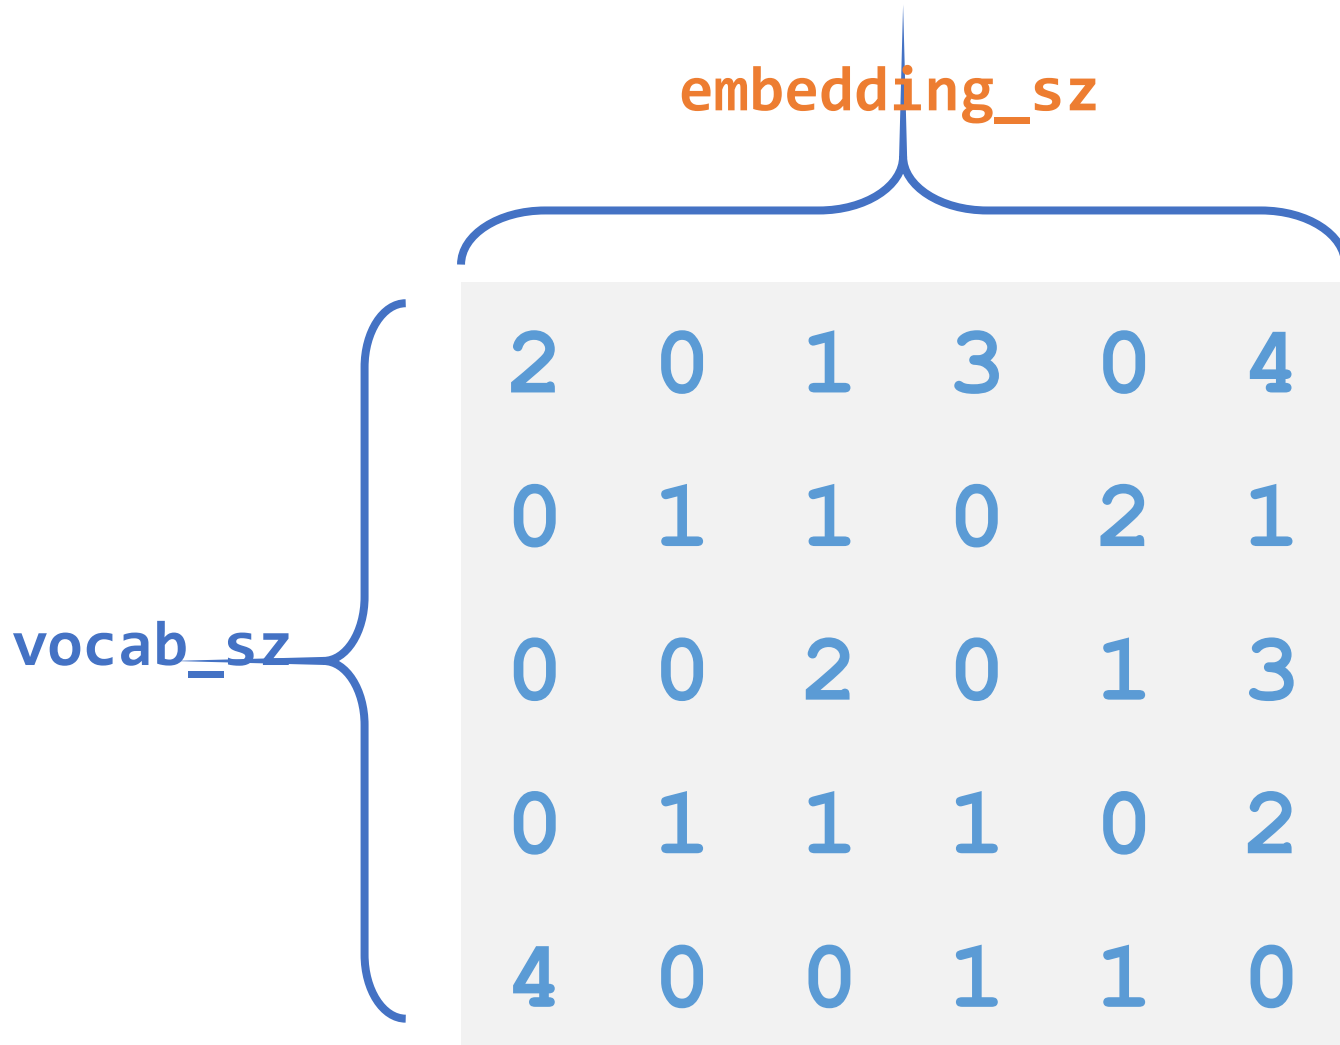
The diagram shows a 5x6 matrix of integers. A blue bracket on the left side, labeled 'vocab\_sz', spans all five rows. A blue bracket on the top, labeled 'embedding\_sz', spans all six columns. The matrix contains the following values:

2	0	1	3	0	4
0	1	1	0	2	1
0	0	2	0	1	3
0	1	1	1	0	2
4	0	0	1	1	0

- 2d matrix: **vocab\_sz** x **embedding\_sz**
- each word corresponds to an index, or word ID – hence the **vocab\_sz** dimension

# Embedding matrix

*How to build this embedding matrix?*



The diagram shows a 5x6 matrix of integers. A blue bracket on the left side, labeled 'vocab\_sz', spans all five rows. A blue bracket on the top side, labeled 'embedding\_sz', spans all six columns. The matrix contains the following values:

2	0	1	3	0	4
0	1	1	0	2	1
0	0	2	0	1	3
0	1	1	1	0	2
4	0	0	1	1	0

- 2d matrix: **vocab\_sz** x **embedding\_sz**
- each word corresponds to an index, or word ID – hence the **vocab\_sz** dimension
- **embedding\_sz** is a hyperparameter



# LM implementation: deep learning

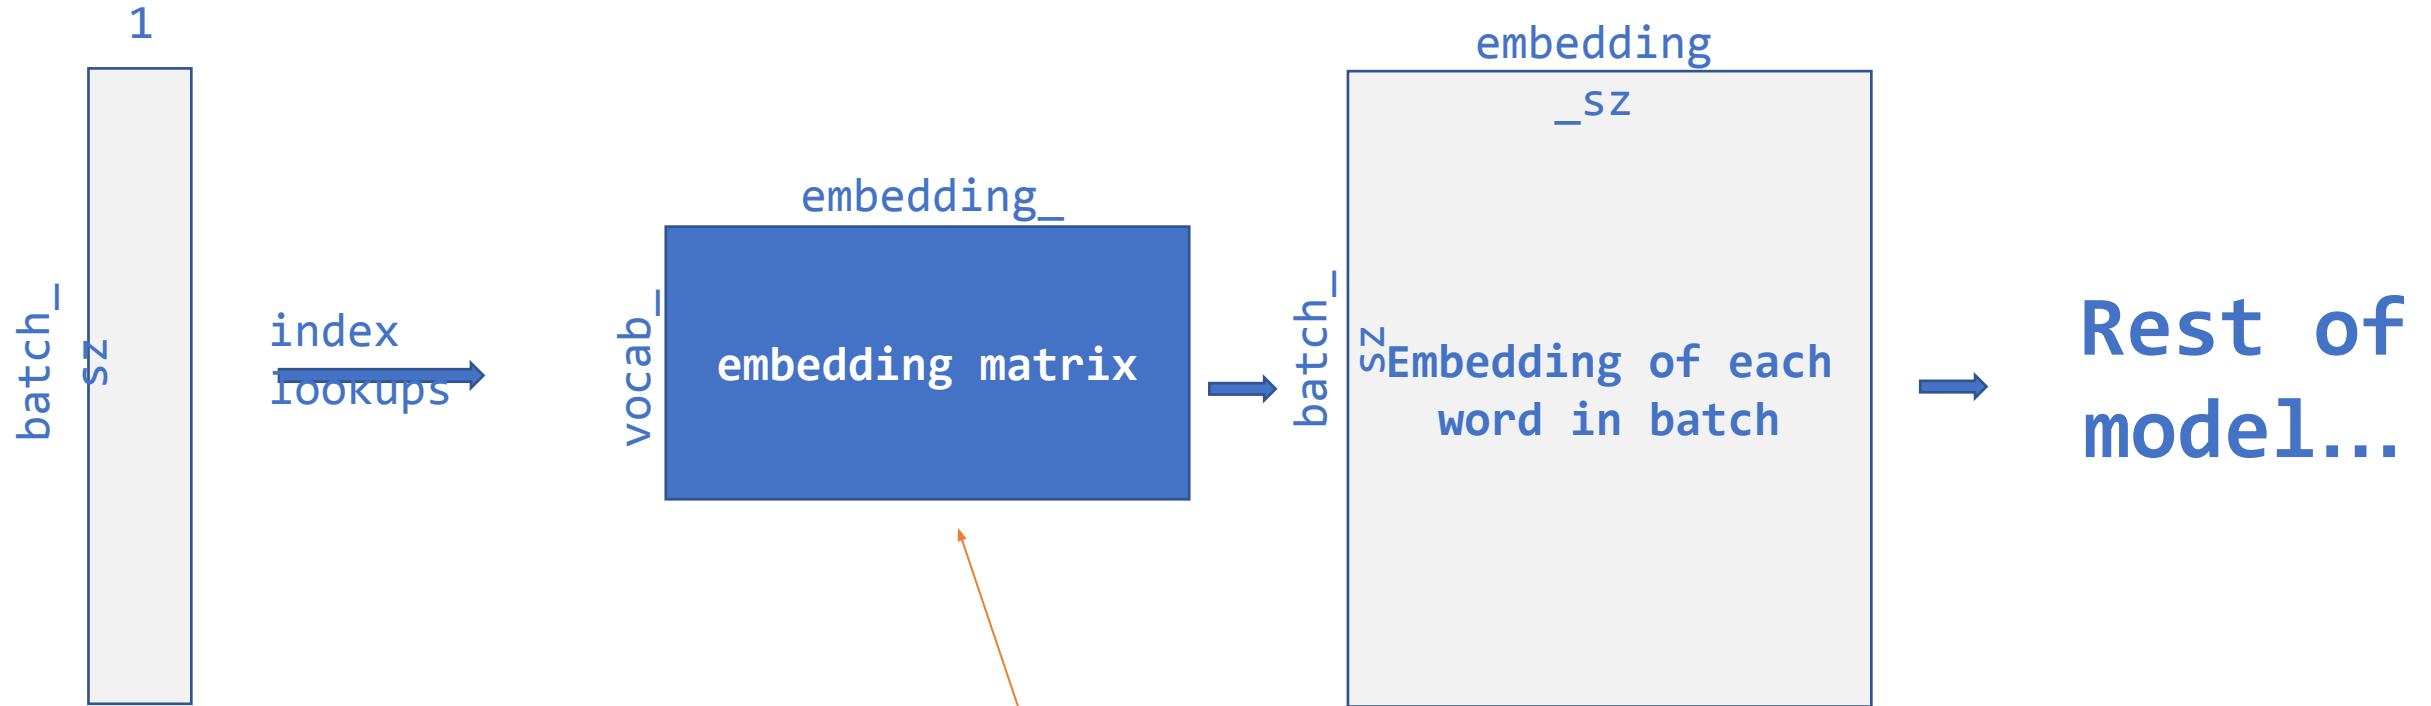
Deep learning helps solve this! **How?**

We can learn an ***embedding matrix*** that associates *related* words with one another for solving a prediction task.

# Using the Embedding Matrix in a Network

If you want to input a [batch of] words into a neural net, this is how:

*they, danced, happily*



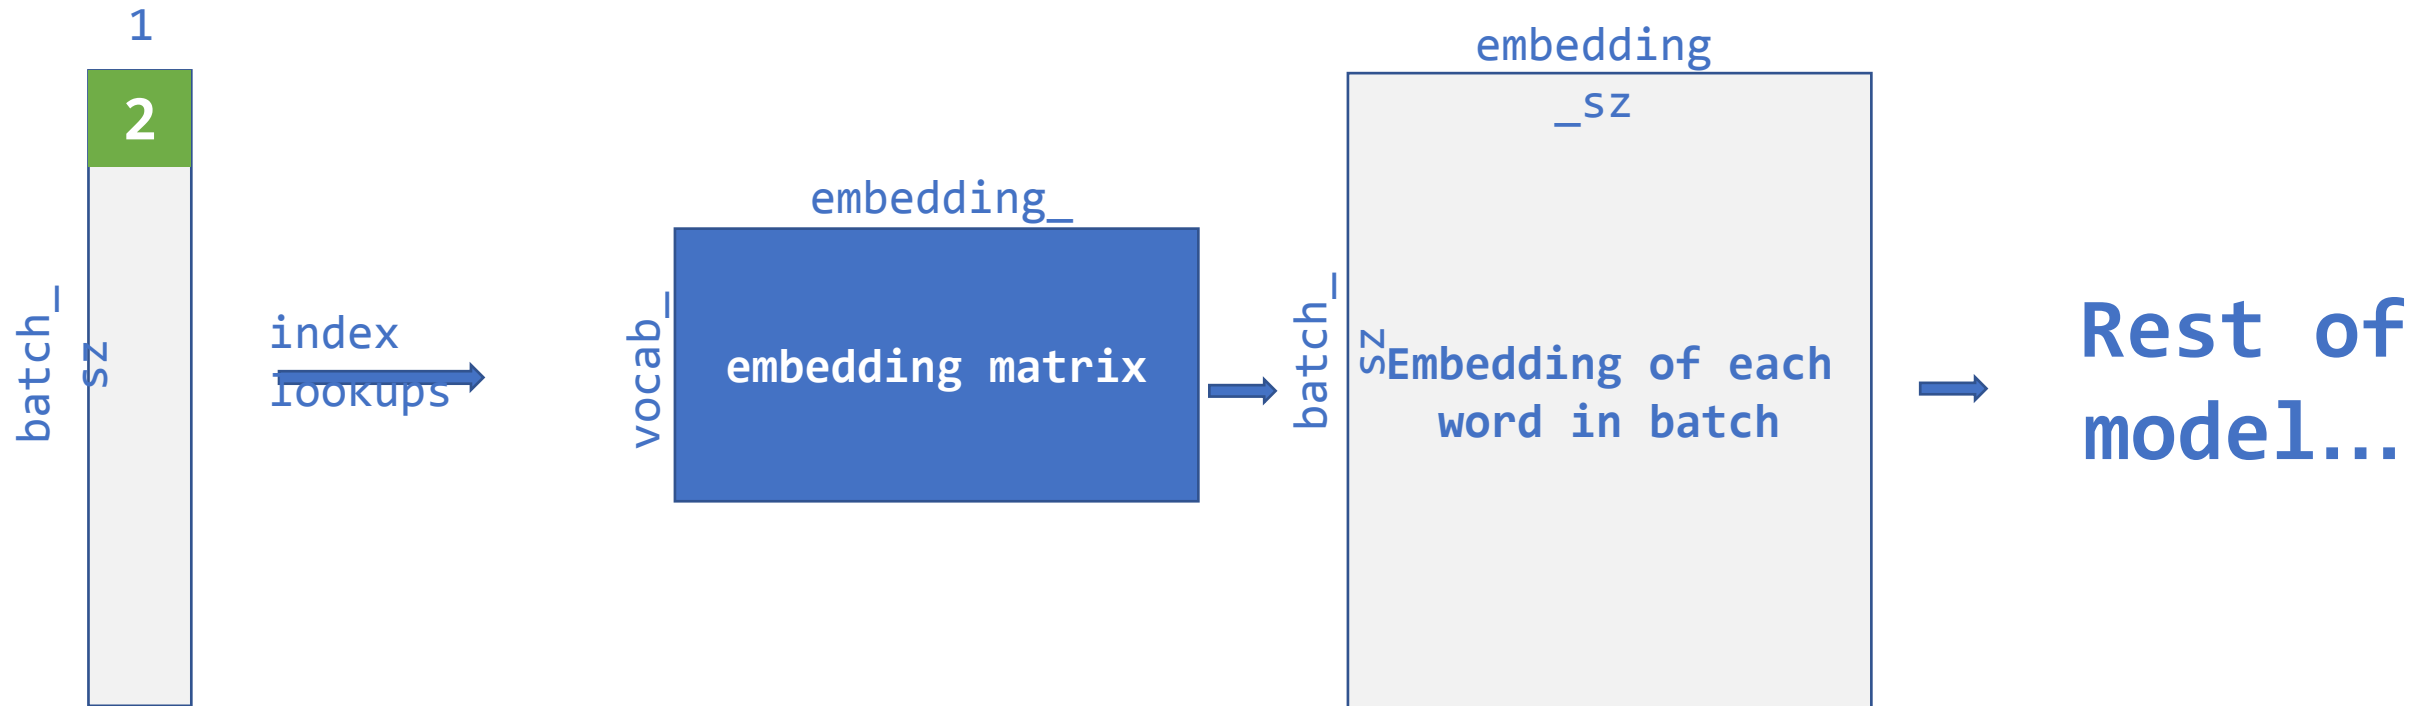
Critical bit: the entries of this matrix can be **learned**!

The network learns what word embeddings are most effective for performing its task

# Using the Embedding Matrix in a Network

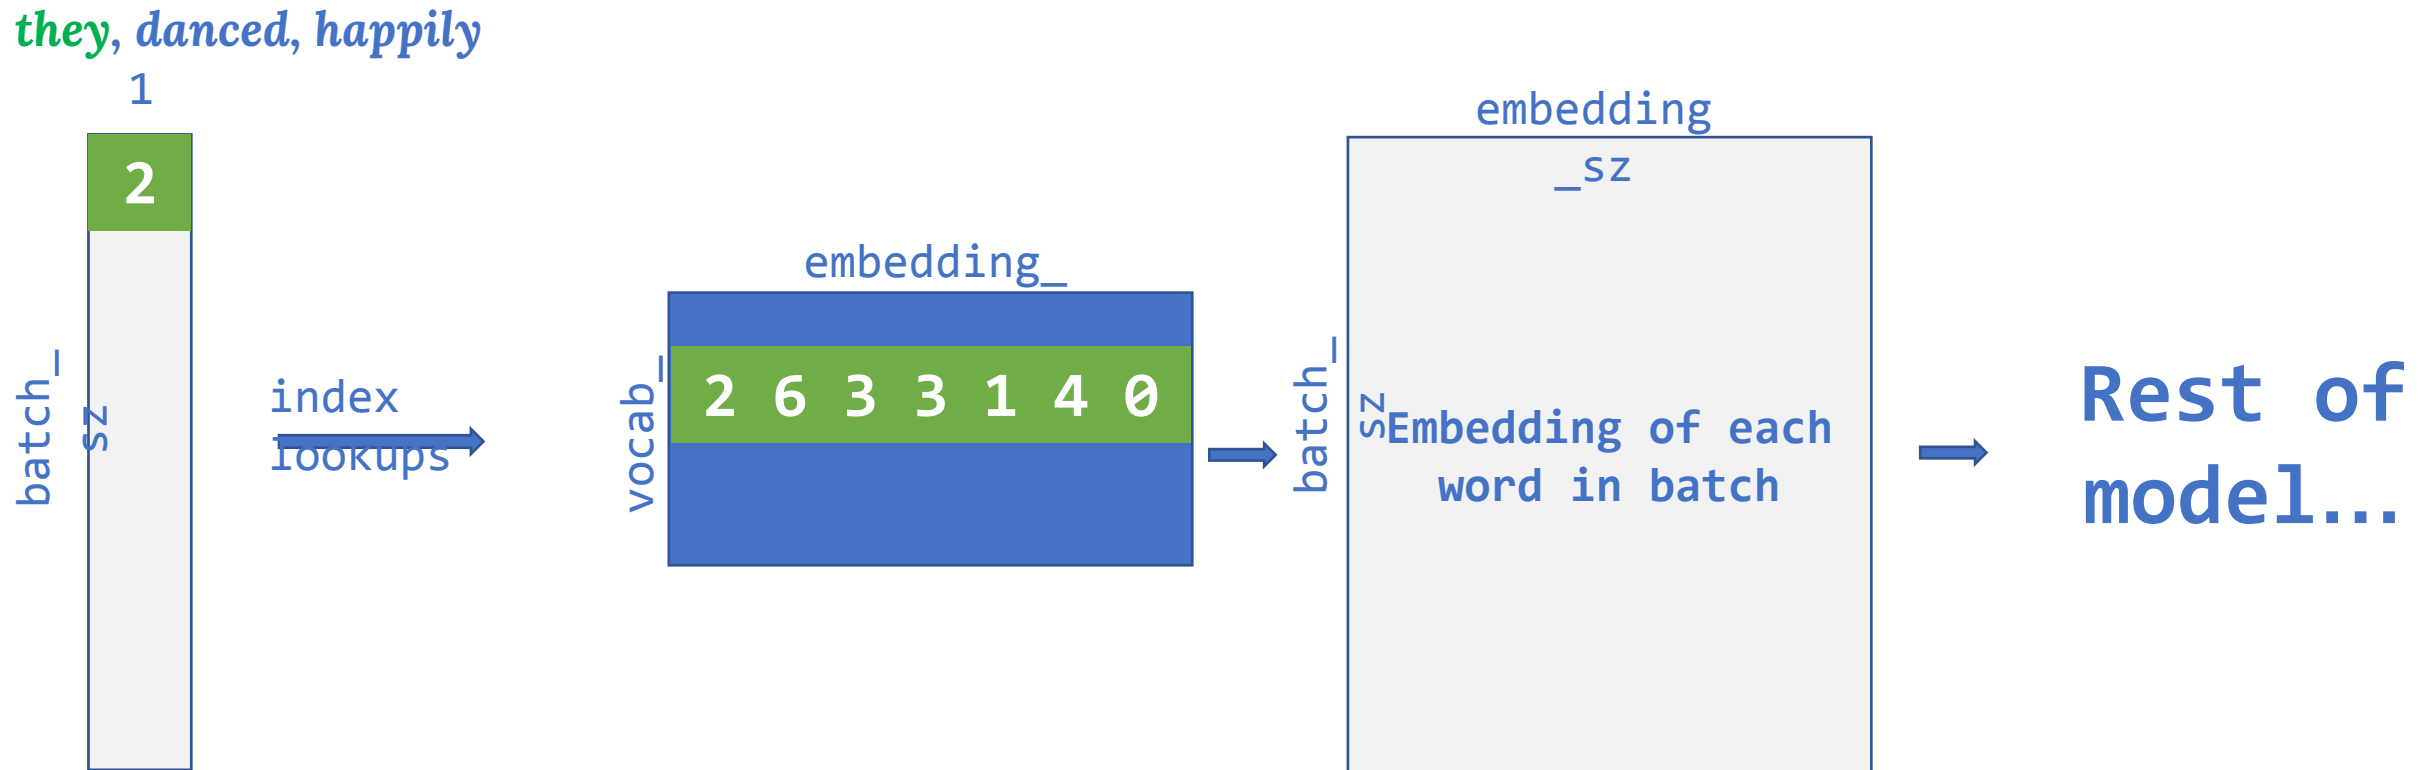
Let's look at the 0<sup>th</sup> word in this batch; its ID in the vocab is 2.

*they, danced, happily*



# Using the Embedding Matrix in a Network

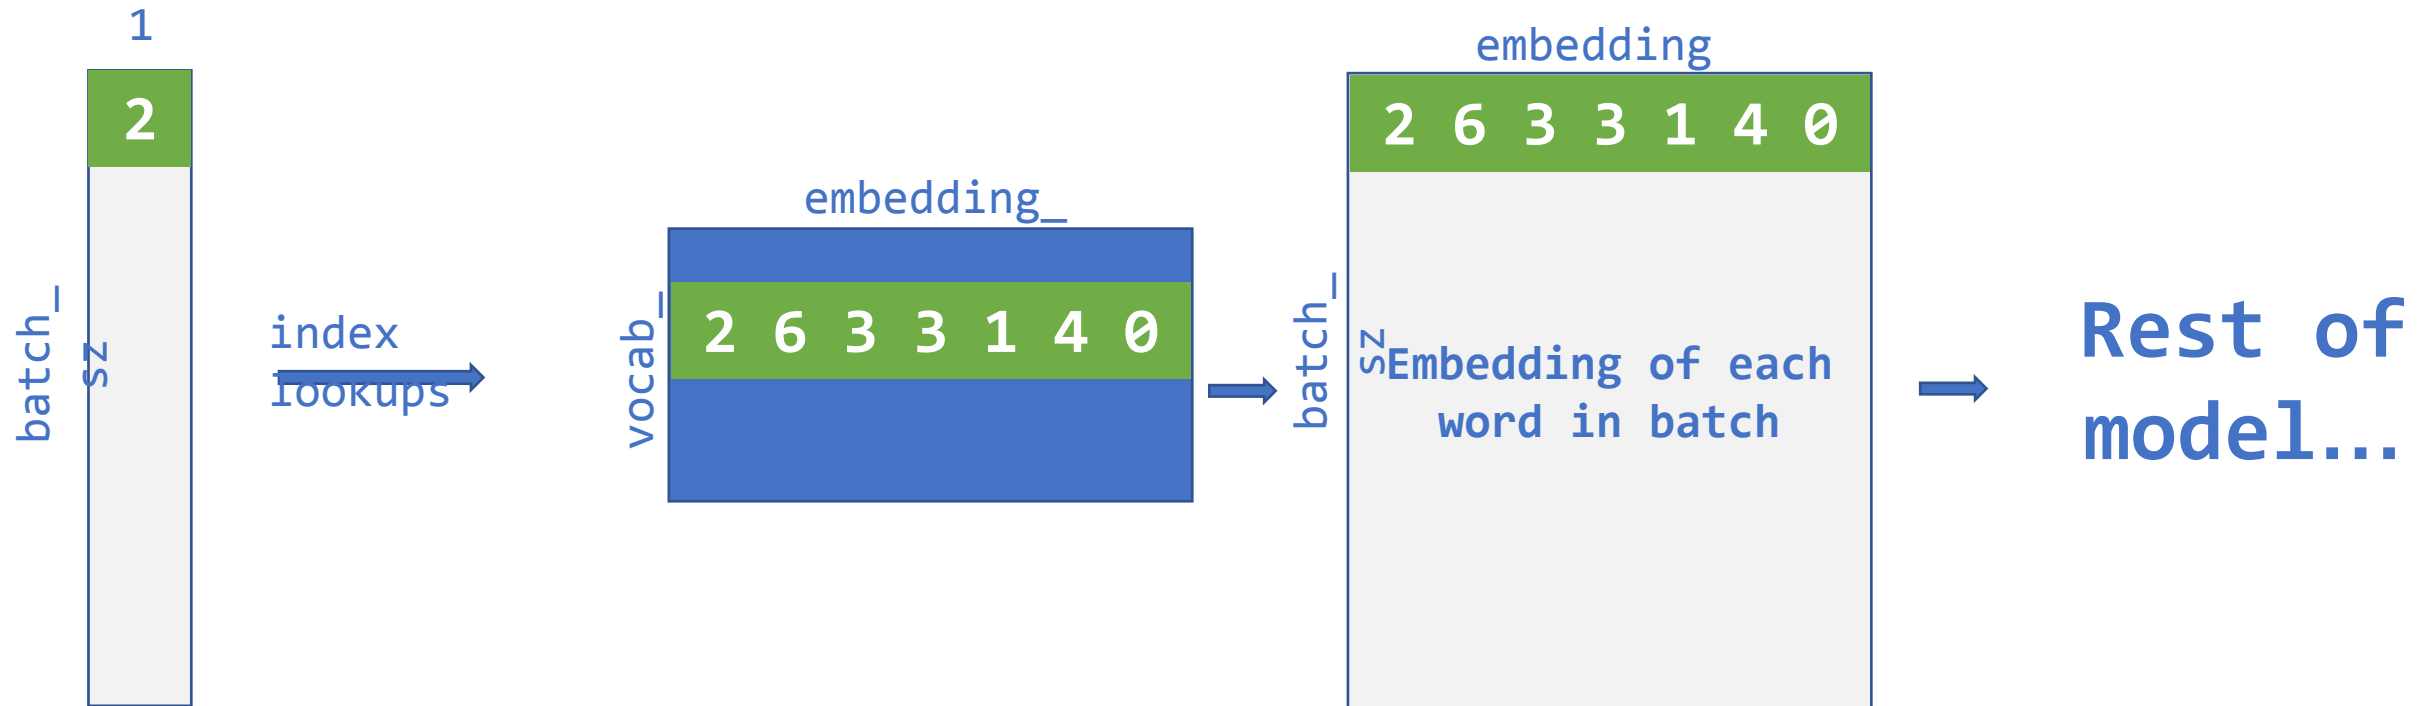
So we look at row 2 of the embedding matrix.



# Using the Embedding Matrix in a Network

We can then pull out this embedding so we can use it in the rest of the model!

*they, danced, happily*

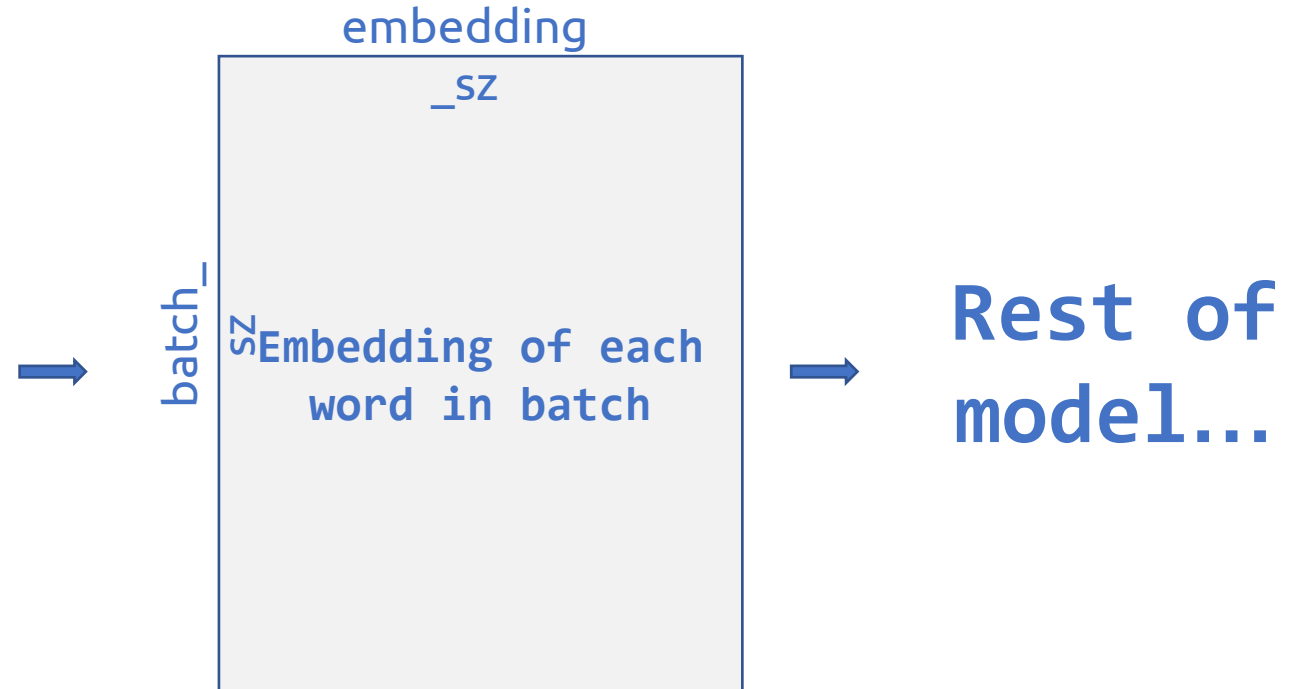


# Using the Embedding Matrix in a Network

In tensorflow, we can use

**`tf.nn.embedding_lookup`**


which takes in an embedding matrix and a list of indices, and returns the embedding corresponding to each index.



# What does the embedding matrix represent?

- Each row in the matrix can be viewed as a vector in vector space

Example 2-D  
vector space:




Vocab size: 3  
Embed size: 2

1	3
2	1
3	2

# What does the embedding matrix represent?

- Each row in the matrix can be viewed as a vector in vector space
- “Embedding”: We’re **embedding** a non-Euclidian entity [a word] into Euclidian space

Example 2-D  
vector space:



Vocab size: 3  
Embed size: 2


1	3
2	1
3	2



# What does the embedding matrix represent?

- Each row in the matrix can be viewed as a vector in vector space
- “Embedding”: We’re **embedding** a non-Euclidian entity [a word] into Euclidian space
- Each row represents the “embedding” for a single word

Example 2-D  
vector space:




Vocab size: 3  
Embed size: 2

1	3
2	1
3	2

# What does the embedding matrix represent?

- Each row in the matrix can be viewed as a vector in vector space
- “Embedding”: We’re **embedding** a non-Euclidian entity [a word] into Euclidian space
- Each row represents the “embedding” for a single word
- This has pretty remarkable properties!

Example 2-D  
vector space:


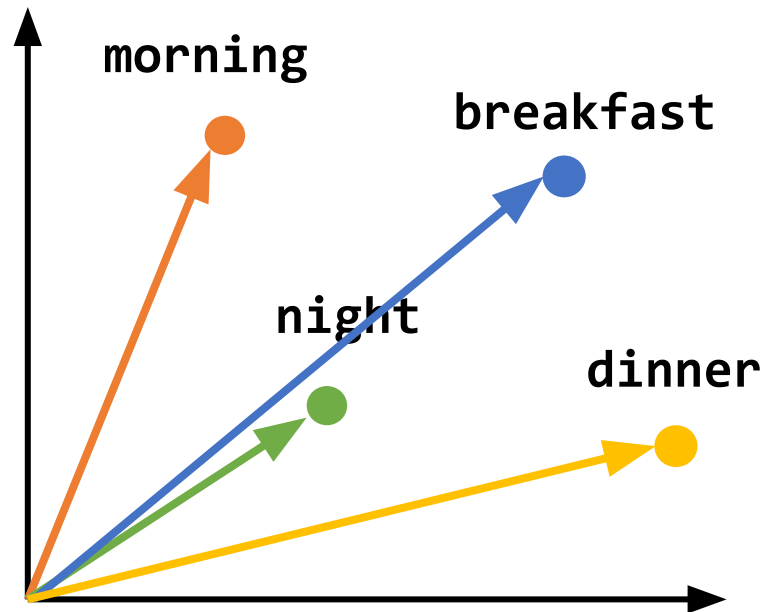


Vocab size: 3  
Embed size: 2


1	3
2	1
3	2

# Vector arithmetic in the embedding matrix

Demo [here](#)



# More 'semantic directions' in embedding space




Male-Female

$$E(\text{queen}) - E(\text{king}) \approx$$


$$E(\text{woman}) - E(\text{man})$$

# More 'semantic directions' in embedding space



Male-Female

$$\begin{aligned} E(\text{queen}) - E(\text{king}) &\approx \\ E(\text{woman}) - E(\text{man}) \end{aligned}$$




Verb tense

$$\begin{aligned} E(\text{walked}) - E(\text{walking}) &\approx \\ E(\text{swam}) - E(\text{swimming}) \end{aligned}$$


# More 'semantic directions' in embedding space

Any questions?




Male-Female

$$\begin{aligned} E(\text{queen}) - E(\text{king}) &\approx \\ E(\text{woman}) - E(\text{man}) \end{aligned}$$



Verb tense

$$\begin{aligned} E(\text{walked}) - E(\text{walking}) &\approx \\ E(\text{swam}) - E(\text{swimming}) \end{aligned}$$

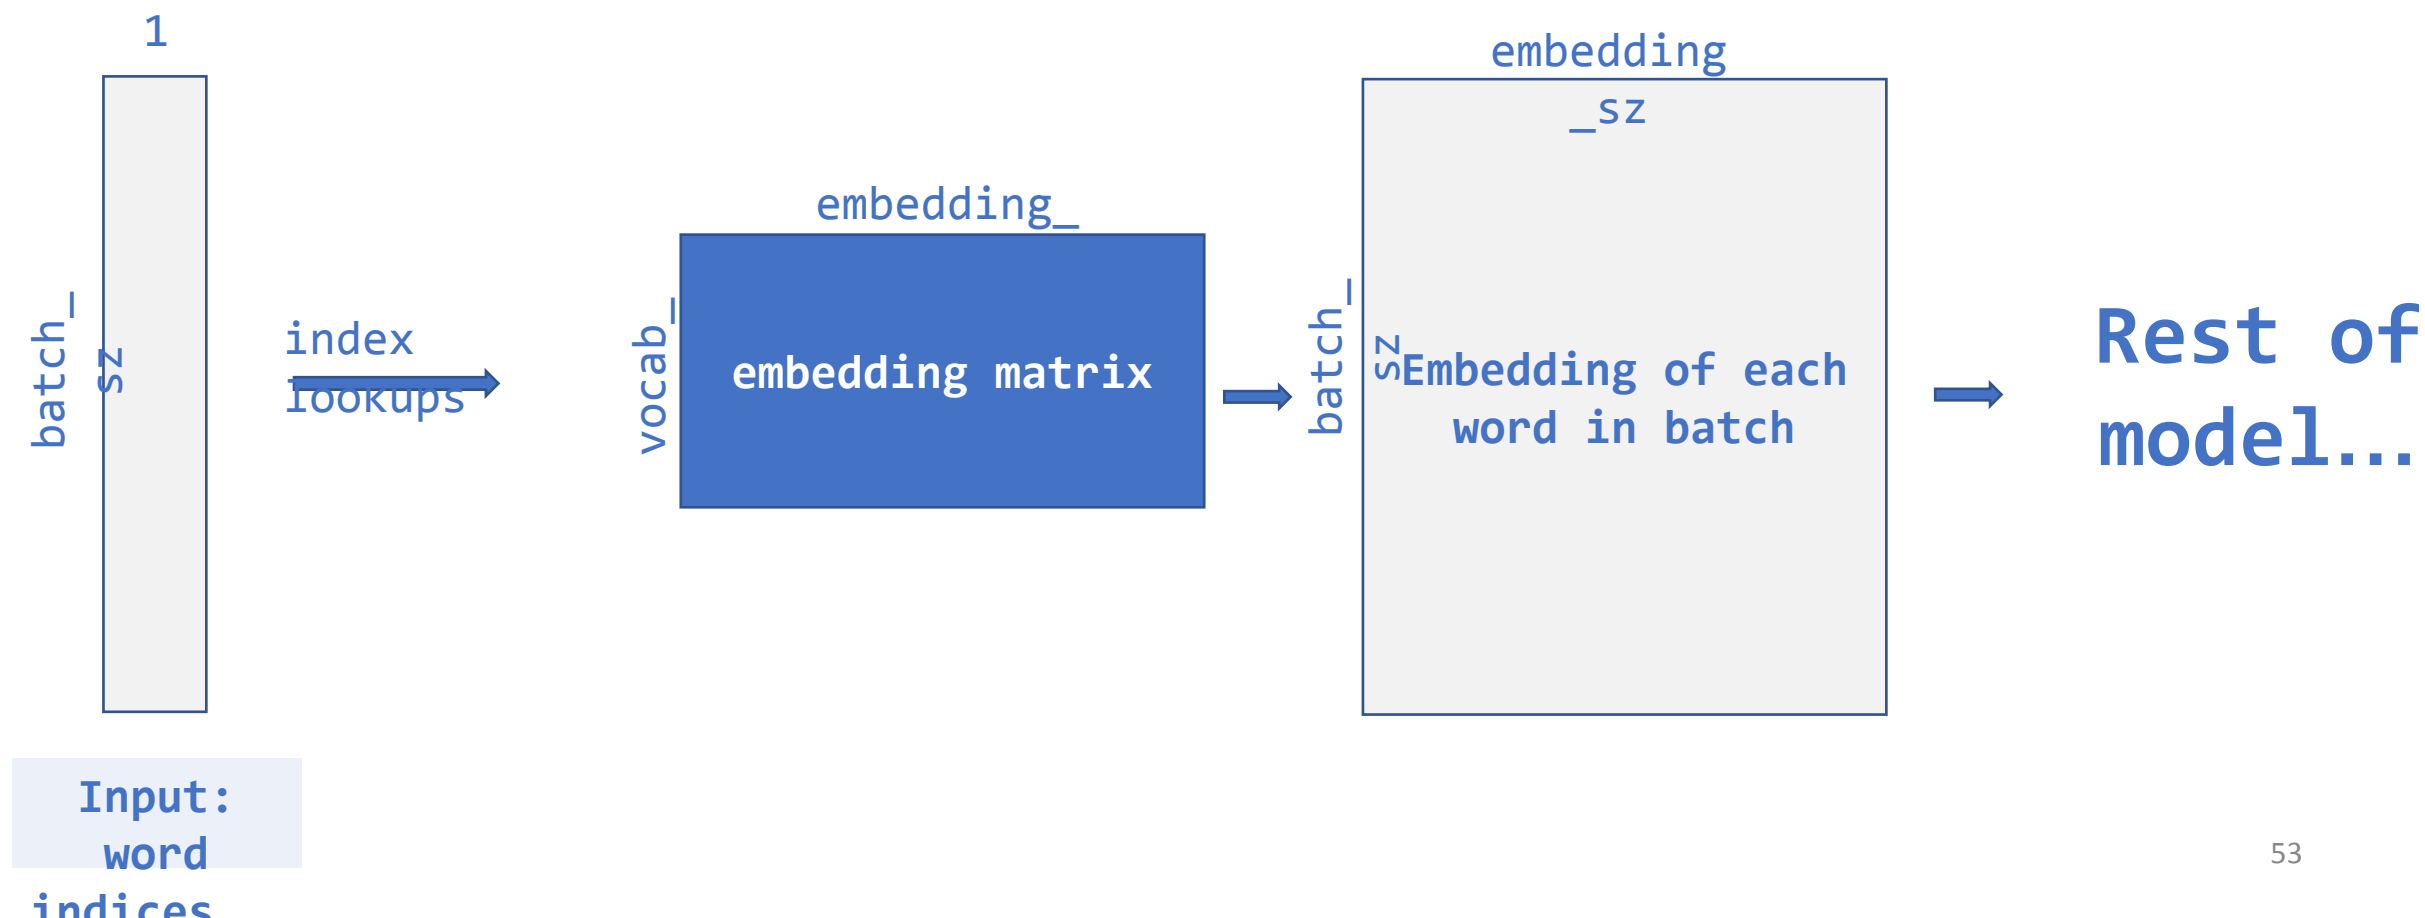


Country-Capital

$$\begin{aligned} E(\text{Spain}) - E(\text{Madrid}) &\approx \\ E(\text{Vietnam}) - E(\text{Hanoi}) \end{aligned}$$

# Why do embedding matrices work like this?

- When the language model is trained, it's incentivized to put words with similar context near each other in the embedding space.



# Why do embedding matrices work like this?

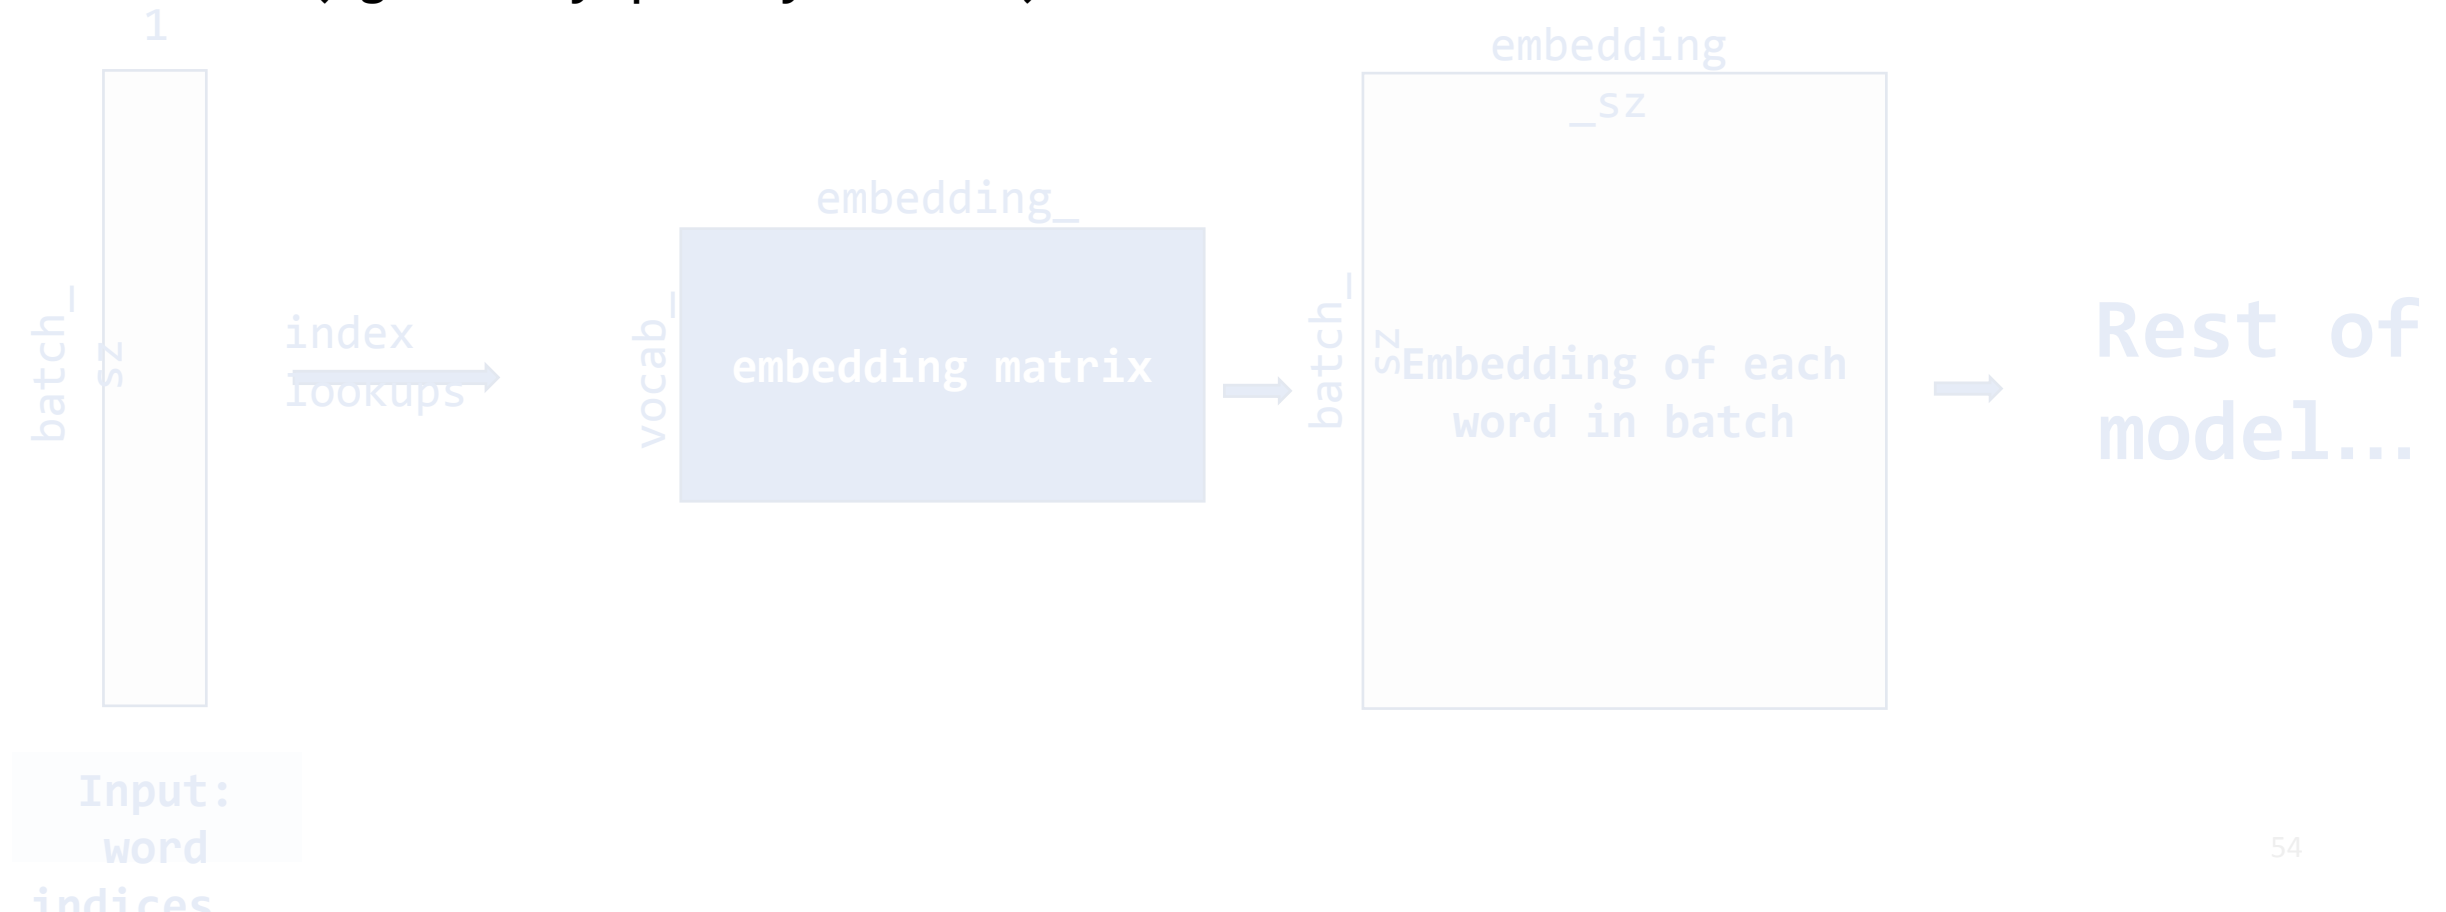
Let's say in the middle of training...

Then, the model sees a lot of "danced gleefully"

$P(\text{"happily"} | \text{"they danced"}) = \text{high}$

How do we increase  $P(\text{"gleefully"} | \text{"they danced"})$ ?

$P(\text{"gleefully"} | \text{"they danced"}) = \text{low}$





# Why do embedding matrices work like this?

Let's say in the middle of training...

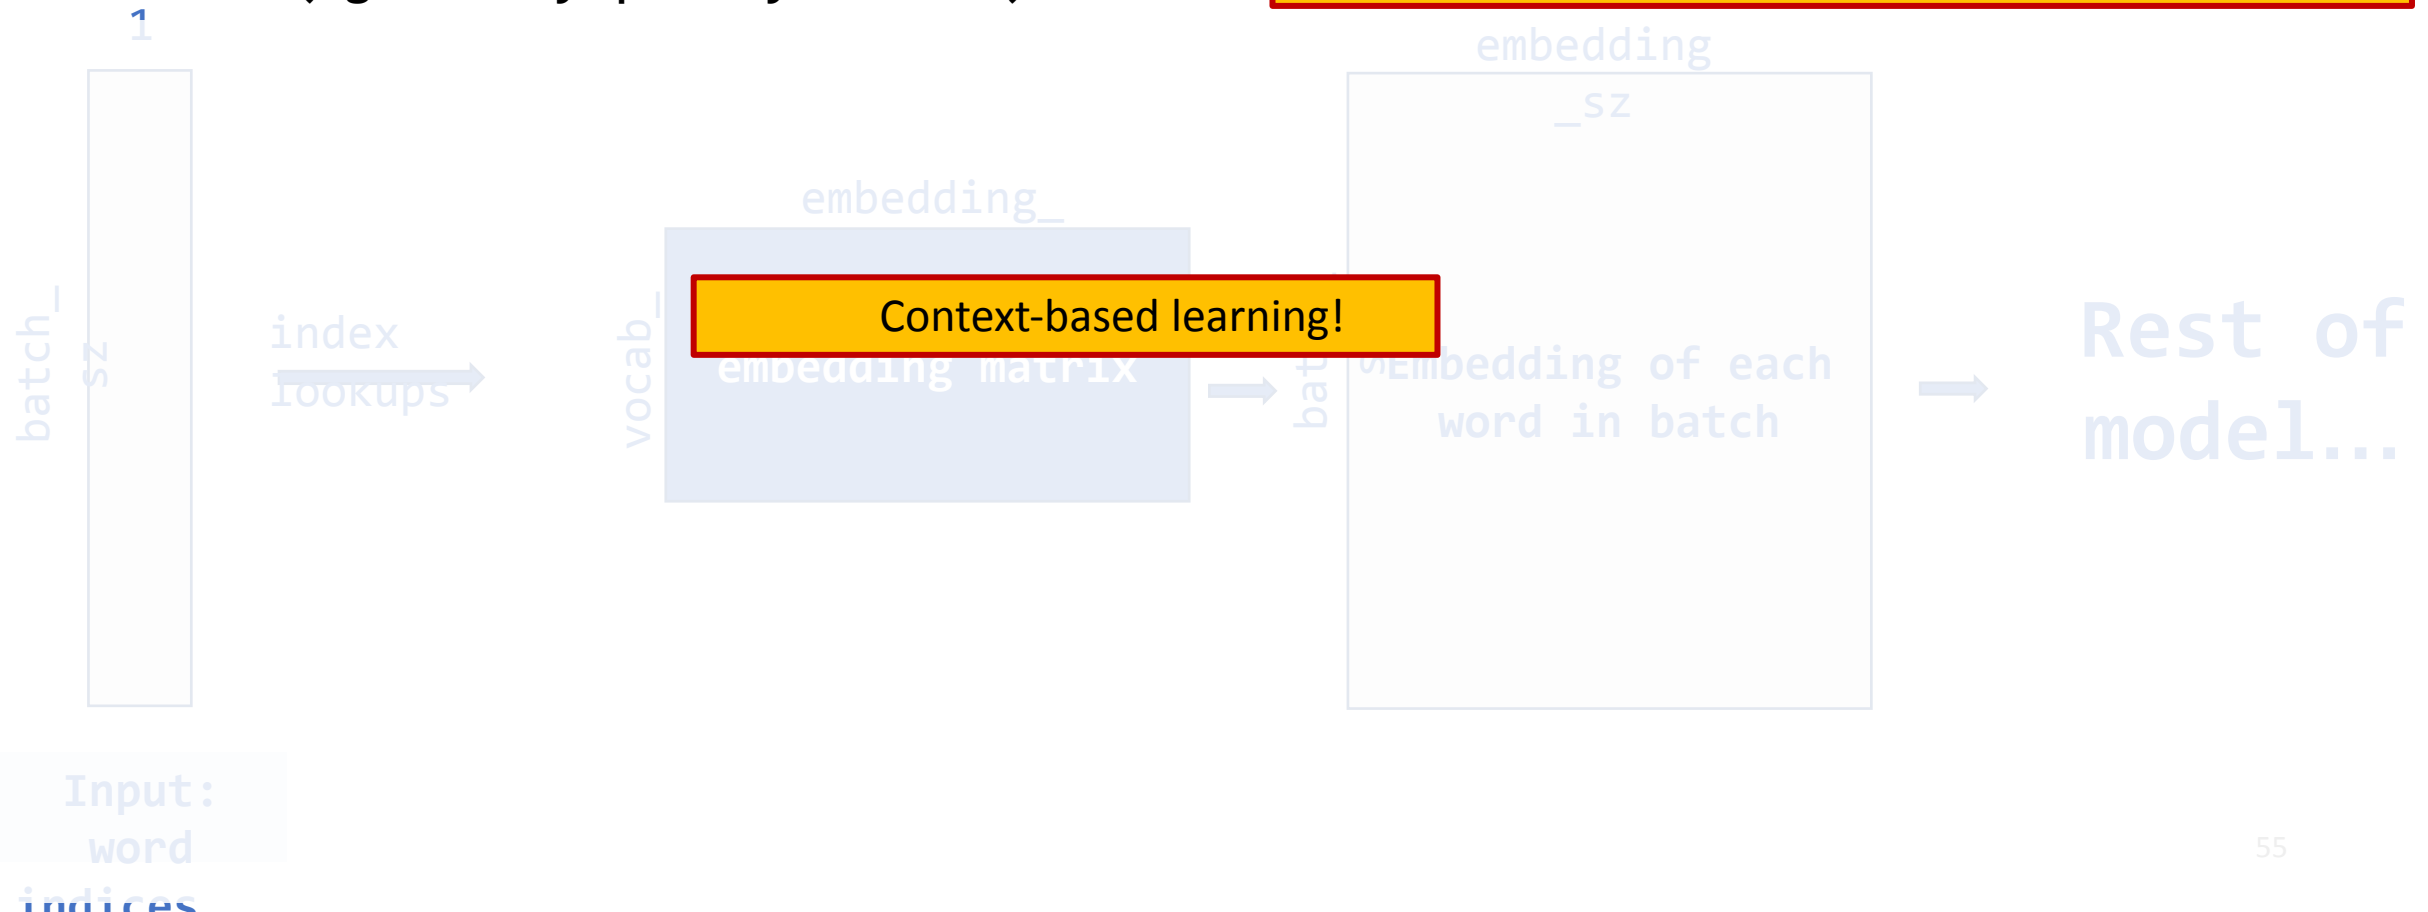
$P(\text{"happily"} | \text{"they danced"}) = \text{high}$

$P(\text{"gleefully"} | \text{"they danced"}) = \text{low}$

Since probability is calculated based on the embedding matrix...


Modify the embedding of **"gleefully"** so that it's similar to the embedding of **"happily"**!

Context-based learning!



# Quantifying “similarity”

- $$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{||A|| ||B||} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$



$$\begin{aligned}\cos(0^\circ) &= 1 \\ \cos(90^\circ) &= -0.448 \\ \cos(180^\circ) &= -0.598\end{aligned}$$

# Limitations of the context-based approach

- Context is correlated with meaning, but context  $\neq$  meaning
- Synonyms typically have similar context:
  - $P(\text{"happily"} \mid \text{"they danced"})$
  - $P(\text{"gleefully"} \mid \text{"they danced"})$
- ...but often antonyms do, too:
  - $P(\text{"happily"} \mid \text{"they danced"})$
  - $P(\text{"unwillingly"} \mid \text{"they danced"})$
- "happily" and "unwillingly" might be used in similar contexts, but have the **opposite** meaning  $\square$  a language model might (erroneously) give them similar embeddings

# Other failure modes are even more dire

What happens when your dataset reflects historical / societal biases?

# Other failure modes are even more dire

What happens when your dataset reflects historical / societal biases?

## **Google News word2vec:**

- Large set of *pretrained* word embeddings, published 2013
- Dataset: news articles aggregated by Google News (100 billion words)

# Other failure modes are even more dire


What happens when your dataset reflects historical / societal biases?

## **Google News word2vec:**

- Large set of *pretrained* word embeddings, published 2013
- Dataset: news articles aggregated by Google News (100 billion words)

**What kinds of relationships do these embeddings contain?**

# Google News word2vec



# Google News word2vec

## - Why did this happen?

- The training dataset (news articles) was biased.
- The news cycle **over-represents** crimes by black perpetrators
  - (Entman 94, Gilliam et.al. 96, Dixon 08, Dixon 15) – this is true over time as well
- Viewers respond more strongly to news stories about crimes by black perpetrators.
  - (Dixon and Maddox 06, Dixon and Azocar 07, Hurley et.al. 15)
  - (News outlets optimize for clicks, therefore report crime by black people more)





why are black women so



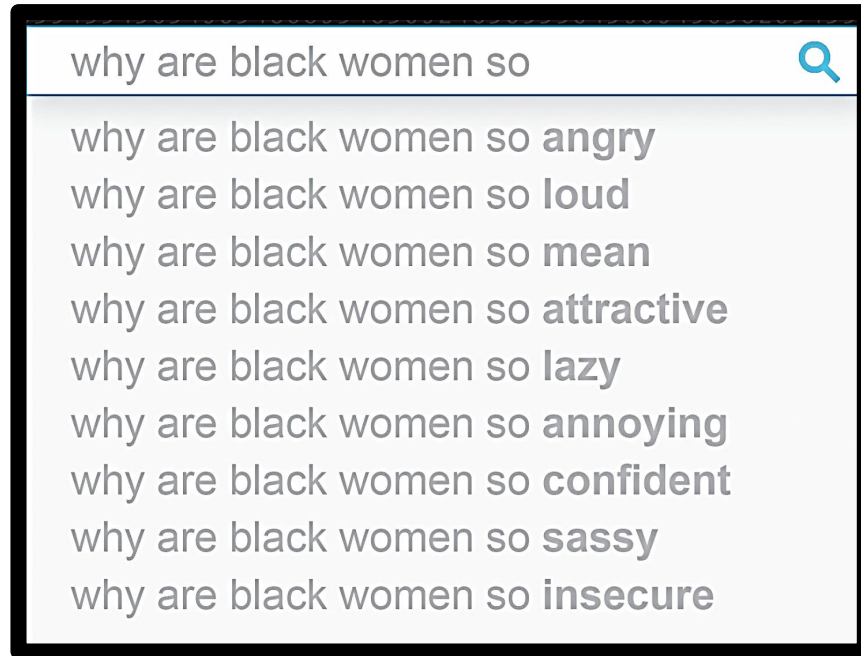
why are black women so **angry**  
why are black women so **loud**  
why are black women so **mean**  
why are black women so **attractive**  
why are black women so **lazy**  
why are black women so **annoying**  
why are black women so **confident**  
why are black women so **sassy**  
why are black women so **insecure**

# ALGORITHMS OF OPPRESSION

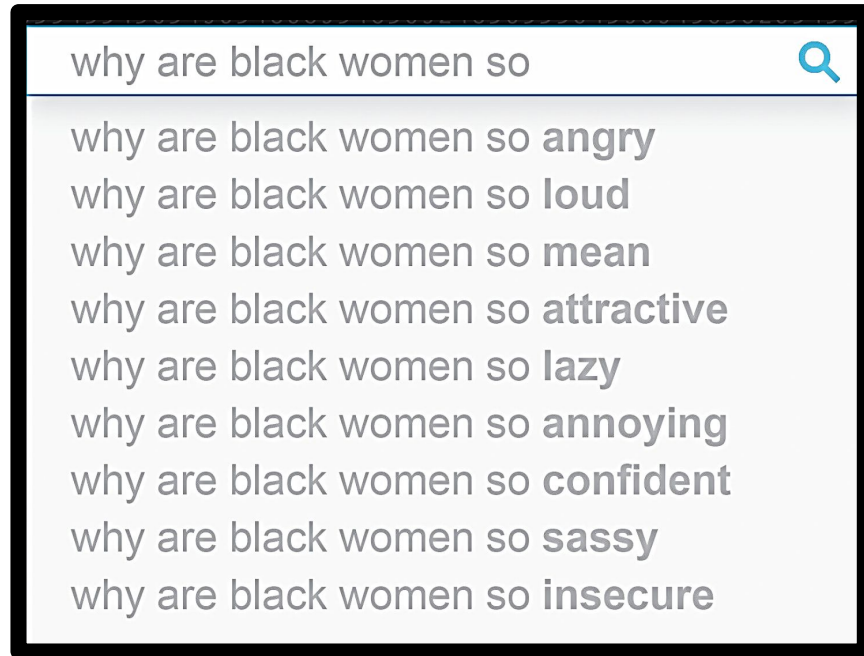
HOW SEARCH ENGINES  
REINFORCE RACISM

SAFIYA UMOJA NOBLE






- In ~2010, when Noble started working on this book, these were the real Google autocomplete suggestions
- ***Takeaway: language models reproduce the biases of the data on which they are trained***
  - ...unless special care is taken—we have an upcoming lab on this!



- Think about the algorithms behind autocomplete, or ad recommendation...
- ***The math might be cool, but there's more to algorithms than math. It is important to consider their potential ethical and social implications once deployed***

# Recap

Language modeling



Language modeling using Deep Learning

Natural Language

Counting/N-grams

Limitations of traditional methods

Learning embedding matrix

Useful properties of embeddings

Limitations of context-based learning

*And this I pray, that your love may abound yet more and more like a controlled use of shared memory.*

