

CSCI 1470/2470
Spring 2023

Ritambhara Singh

April 05, 2023
Wednesday

Deep Learning



Review: Supervised v/s Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation, image
captioning, etc.

Unsupervised Learning

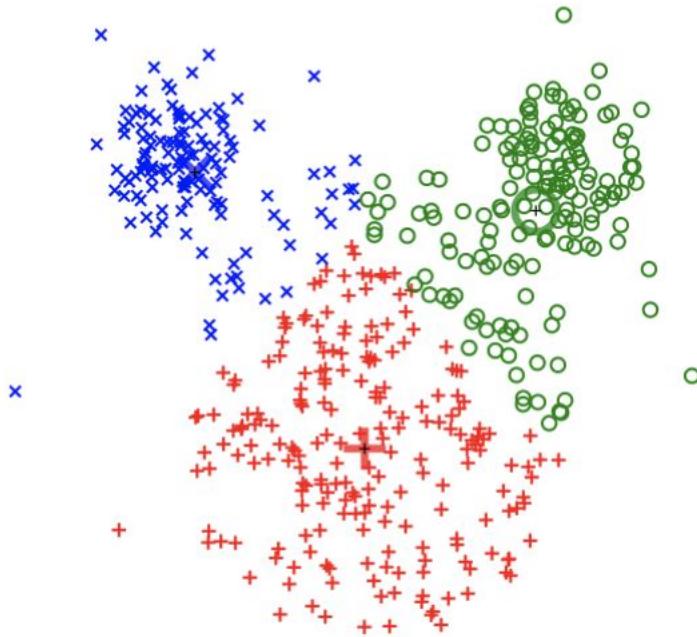
Data: x

Just data, no labels!

Goal: Learn some underlying hidden
structure of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

Review: Unsupervised Learning



k-means clustering

[This image](#) is CC0 public domain

Unsupervised Learning

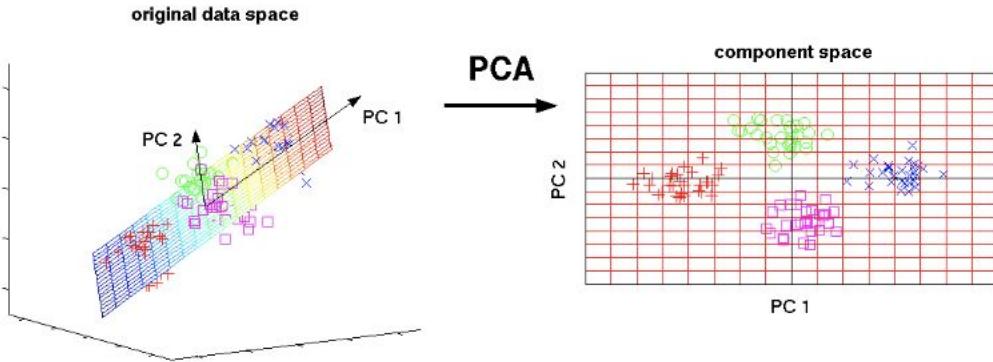
Data: x

Just data, no labels!

Goal: Learn some underlying hidden
structure of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

Review: Unsupervised Learning



dimensionality reduction

[This image](#) is CC0 public domain

Unsupervised Learning

Data: x

Just data, no labels!

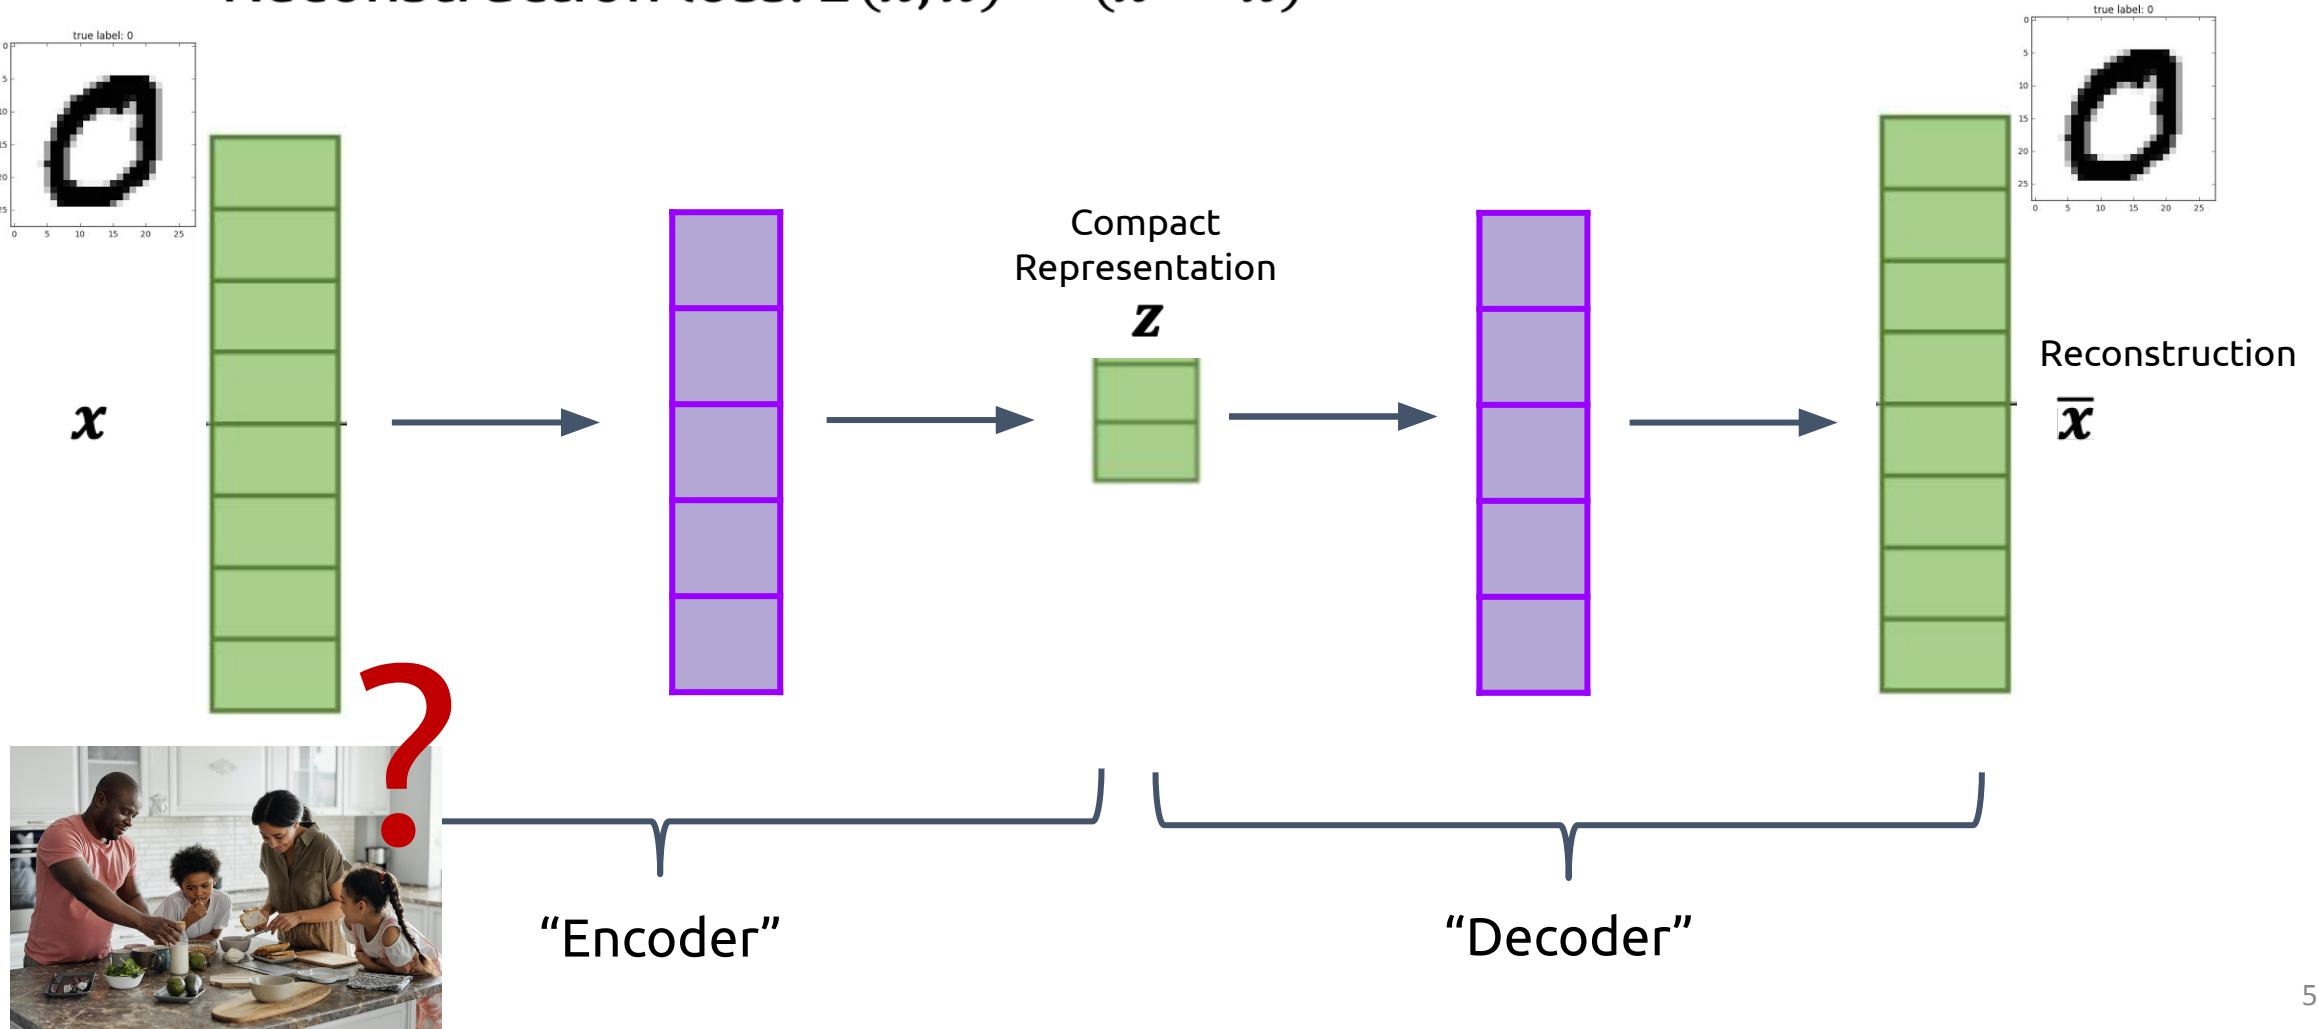
Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering,

dimensionality reduction, feature learning, density estimation, etc.

Review: Autoencoder

- Reconstruction loss: $L(x, \bar{x}) = (x - \bar{x})^2$



Today's goal – learn about variational
autoencoders (VAEs)

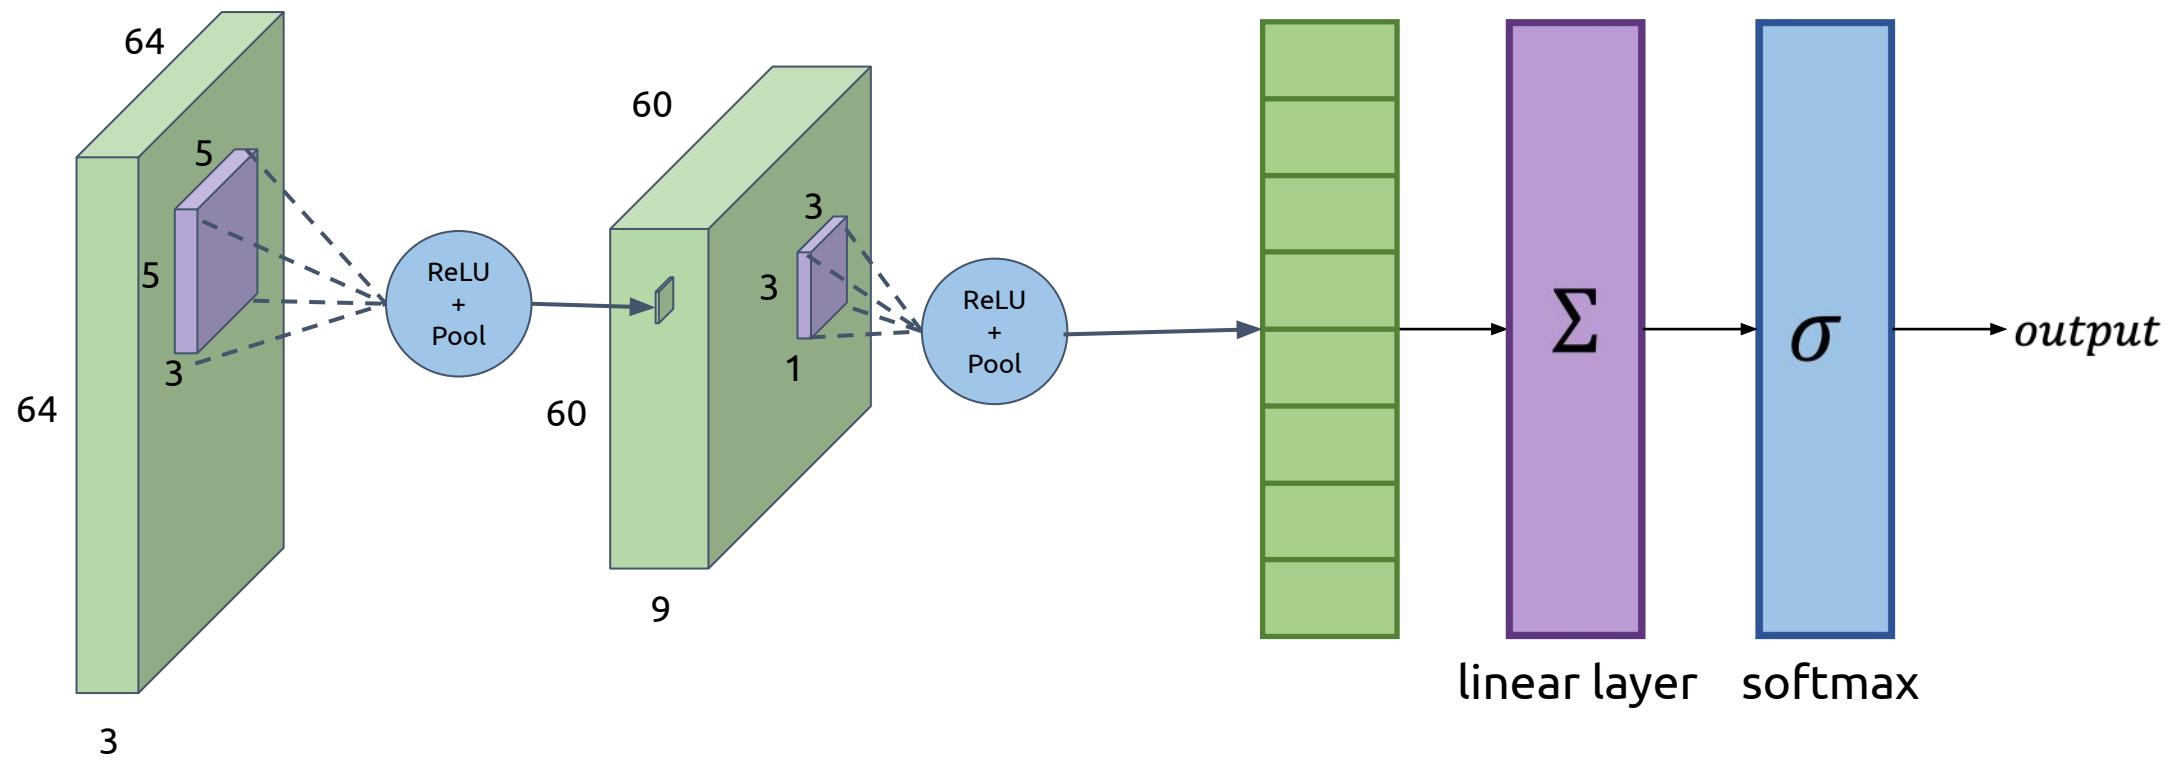
(1) Convolutional AEs

(2) Generative models

(3) Variational Autoencoders (VAEs)

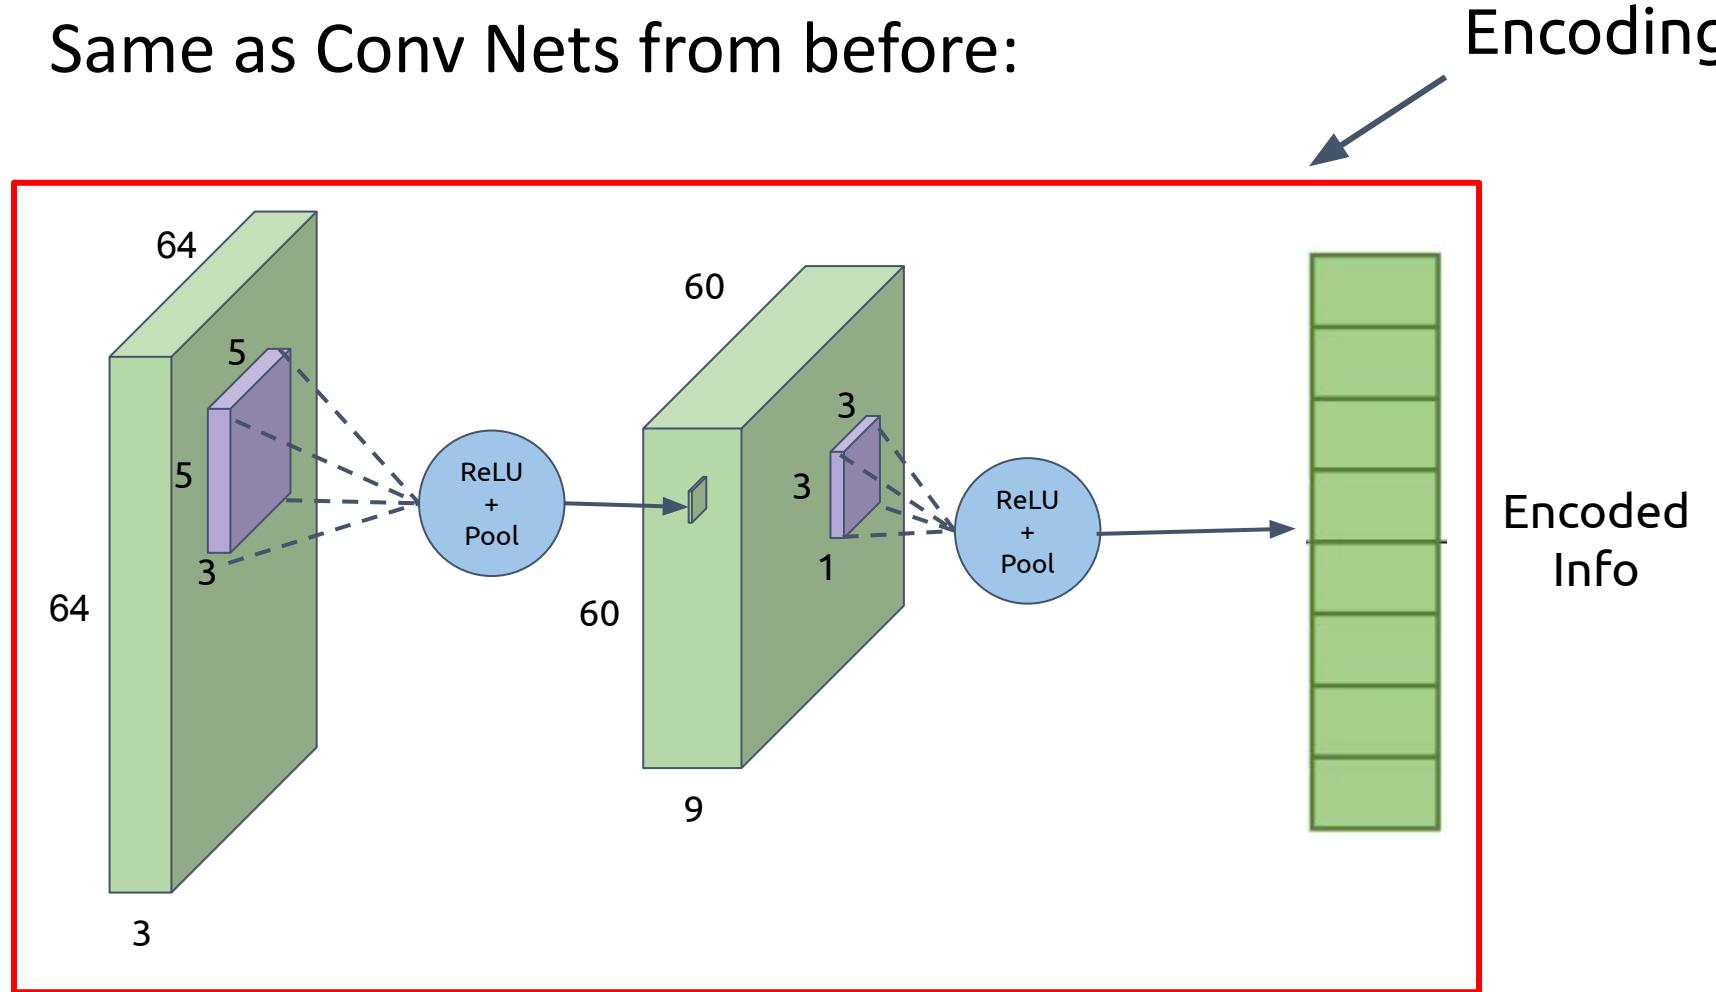
Convolutional Autoencoders

- CNNs are great for image processing in Neural Networks
- How can we build a *convolutional* autoencoder?



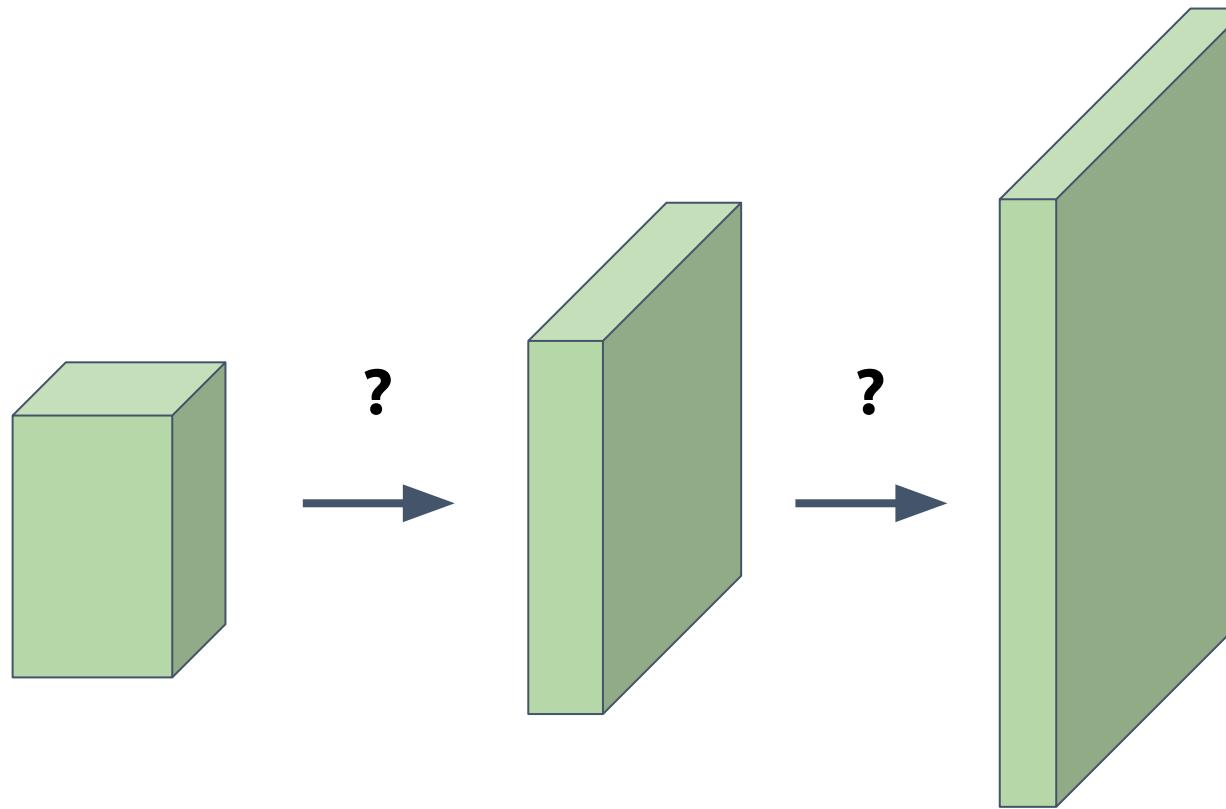
Convolutional Autoencoders: Encoding

Same as Conv Nets from before:



Autoencoders: Decoding

- Convolution as we know it only keeps resolution same or decreases it
- How do we go up in resolution?



Autoencoders: Transpose Convolution

- Convolution can be viewed as a matrix multiplication
- How do we represent it this way?

2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

Input

?

1	2	3
4	5	6
7	8	9

Kernel

=

57	60
66	61

Output

Autoencoders: Transpose Convolution

Step 1: Flatten the image into a column vector

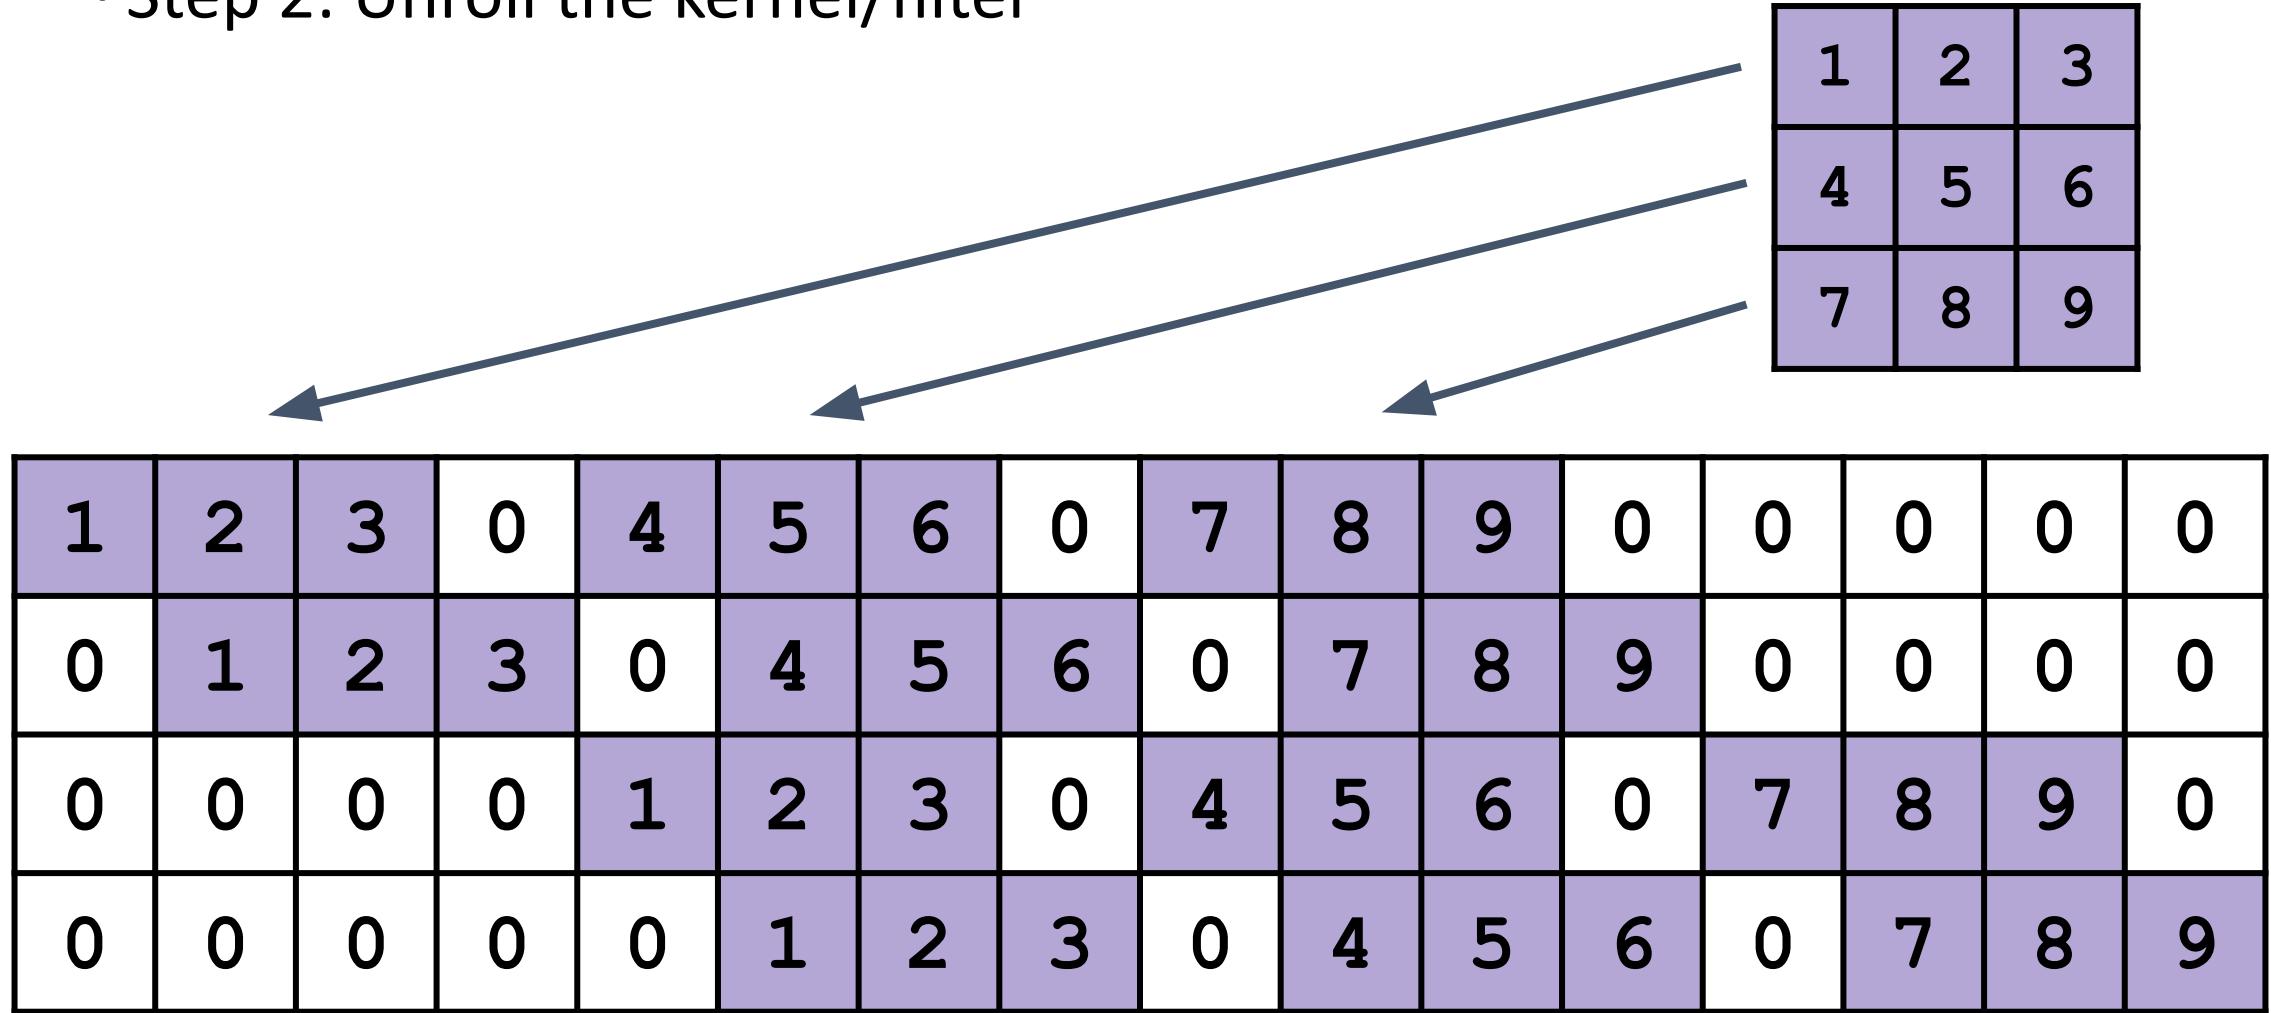
2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1



2
1
0
3
0
0
1
2
3
1
2
0
0
2
2
1

Autoencoders: Transpose Convolution

- Step 2: Unroll the kernel/filter



Autoencoders: Transpose Convolution

Step 3: Matrix multiply
unrolled kernel with flattened
image

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9



2
1
0
3
0
0
1
2
3
1
2
0
0
2
2
1

=

57
50
66
61

Autoencoders: Transpose Convolution

Each row of the convolution matrix corresponds to a dot product between filter and image patch:

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

Autoencoders: Transpose Convolution

Each row of the convolution matrix corresponds to a dot product between filter and image patch:

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

Autoencoders: Transpose Convolution

Each row of the convolution matrix corresponds to a dot product between filter and image patch:

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

Autoencoders: Transpose Convolution

Each row of the convolution matrix corresponds to a dot product between filter and image patch:

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9



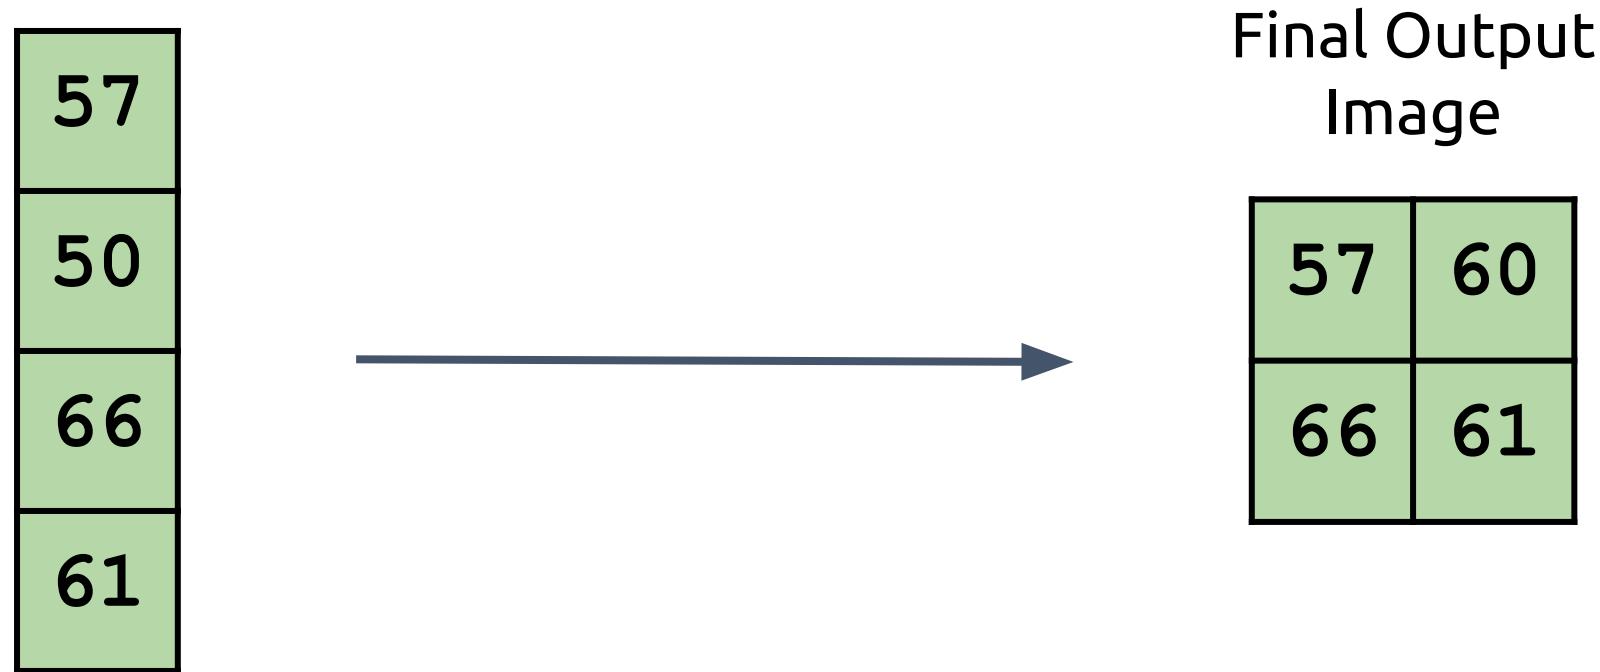
2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1



2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

Autoencoders: Transpose Convolution

Step 4: Finally reshape the output back into a grid



Autoencoders: Transpose Convolution

1	2	3	0	4	5	6	0	7	8	9	0	0	0	0	0
0	1	2	3	0	4	5	6	0	7	8	9	0	0	0	0
0	0	0	0	1	2	3	0	4	5	6	0	7	8	9	0
0	0	0	0	0	1	2	3	0	4	5	6	0	7	8	9



2
1
0
3
0
0
1
2
3
1
2
0
0
2
2
1

=

57
50
66
61

Autoencoders: Transpose Convolution

To upsample an image, we just do the inverse of this operation.

What matrix do we use?

The **transpose** of the big convolution matrix



1	0	0	0
2	1	0	0
3	2	0	0
0	3	0	0
4	0	1	0
5	4	2	1
6	5	3	2
0	6	0	3
7	0	4	0
8	7	5	4
9	8	6	5
0	9	0	6
0	0	7	0
0	0	8	7
0	0	9	8
0	0	0	9

Input image
flattened to
column vector



1
0
2
1



1
2
3
0
6
10
14
3
15
22
26
6
14
23
26
9
20

Autoencoders: Transpose Convolution

Finally, reshape the output vector into a grid to get the final output image:

1
2
3
0
6
10
14
3
15
22
26
6
14
23
26
9

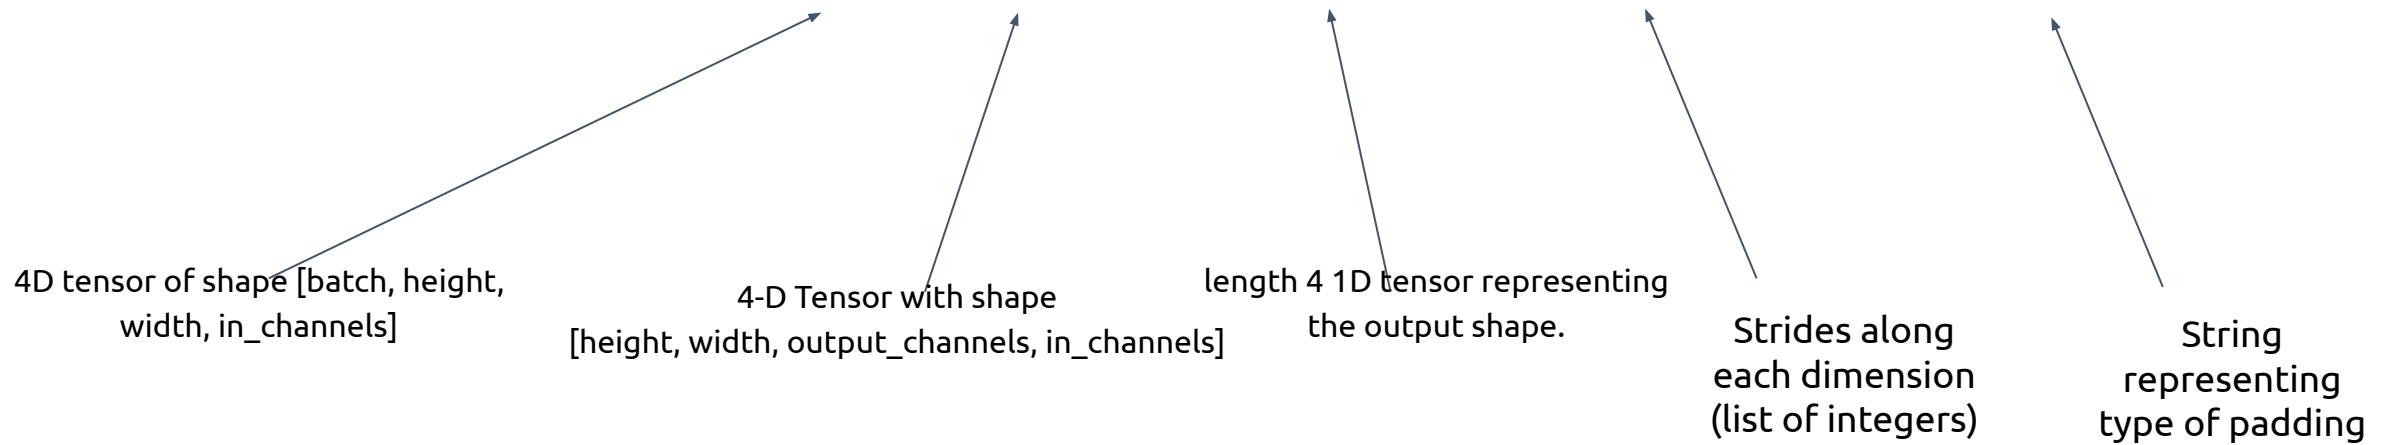


Final output image

1	2	3	4
6	10	14	3
15	1	22	26
14	23	26	9

Transpose Convolution in Tensorflow

```
tf.nn.conv2d_transpose(input, filters, output_shape, strides, padding='SAME')
```



Documentation here: https://www.tensorflow.org/versions/r2.0/api_docs/python/tf/nn/conv2d

Transpose Convolution in Tensorflow

```
tf.nn.conv2d_transpose(input, filters, output_shape, strides, padding='SAME')
```

Why do we need to specify output size?

Specifying Output Size

- An image can be the result of the same convolution on images of different resolution
- We need to specify which one we want.

57	60
66	61

1	2	3
4	5	6
7	8	9

Kernel



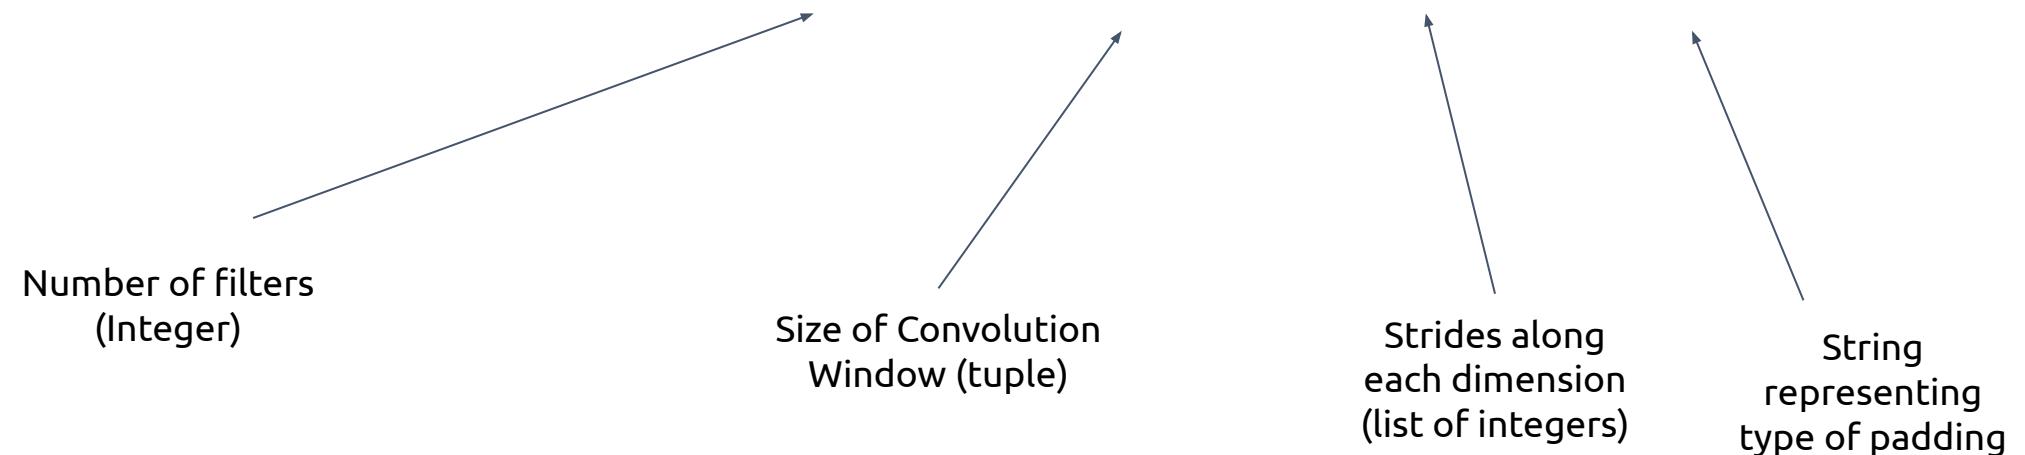
2	1	0	3
0	0	1	2
3	1	2	0
0	2	2	1

2	1	0	3	0
0	0	1	2	0
3	1	2	0	0
0	2	2	1	0
0	0	0	0	0



Transpose Convolution in Keras

```
tf.keras.layers.Conv2DTranspose(filters, kernel_size, strides, padding='SAME')
```



Note: Output Shape is inferred, but can be specified via the “output_padding” parameter

Documentation here: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2DTranspose

Today's goal – learn about variational
autoencoders (VAEs)

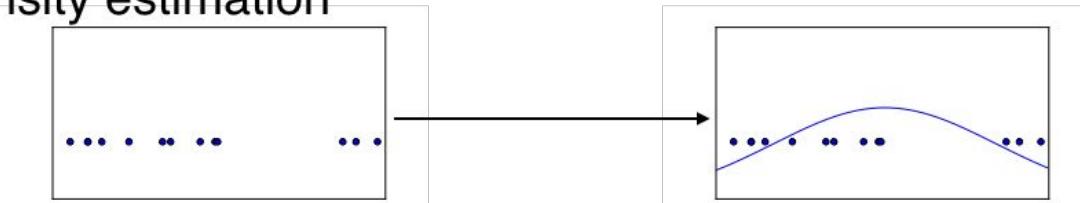
(1) Convolutional AEs

(2) Generative models

(3) Variational Autoencoders (VAEs)

Unsupervised Learning

- Density estimation



Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering,
dimensionality reduction, feature
learning, **density estimation**, etc.

Unsupervised Learning

Generative models

- Density estimation



- Sample generation



Training examples

Model samples

Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, **density estimation**, etc.

Discriminative v/s Generative models

Discriminative Model:

Learn a probability distribution $p(y|x)$

Generative Model:

Learn a probability distribution $p(x)$

Data: x



Label: y

Cat

Discriminative v/s Generative models

Discriminative Model:

Learn a probability distribution $p(y|x)$

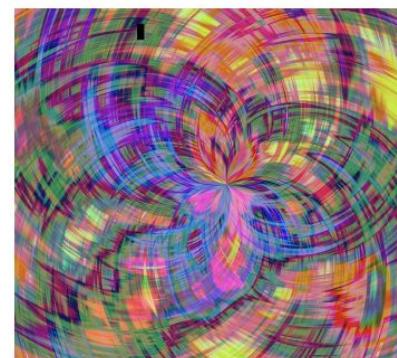


$$P(\text{cat} | \text{monkey})$$

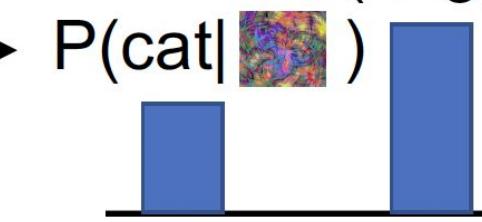


Generative Model:

Learn a probability distribution $p(x)$



$$P(\text{dog} | \text{colorful pattern})$$



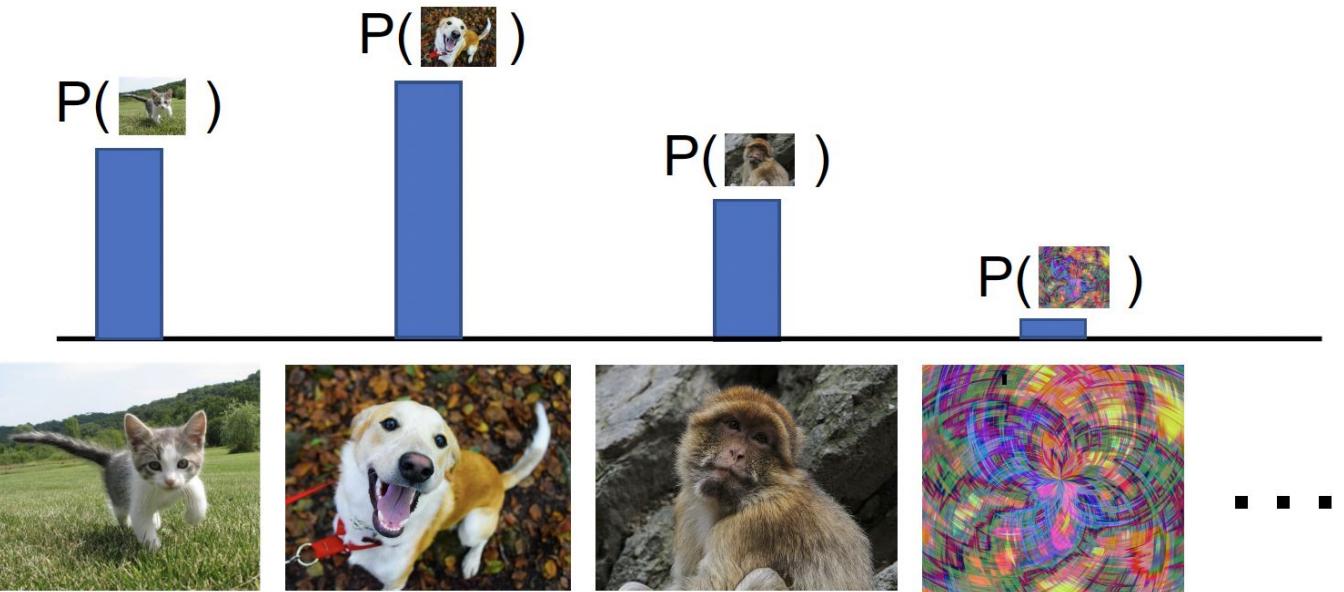
Discriminative model: No way for the model to handle unreasonable inputs; it must give label distributions for all images

Discriminative v/s Generative models

Discriminative Model:

Learn a probability distribution $p(y|x)$

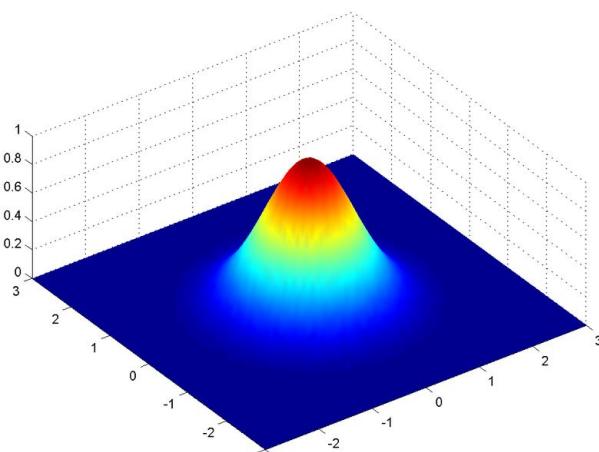
Generative Model:
Learn a probability distribution $p(x)$



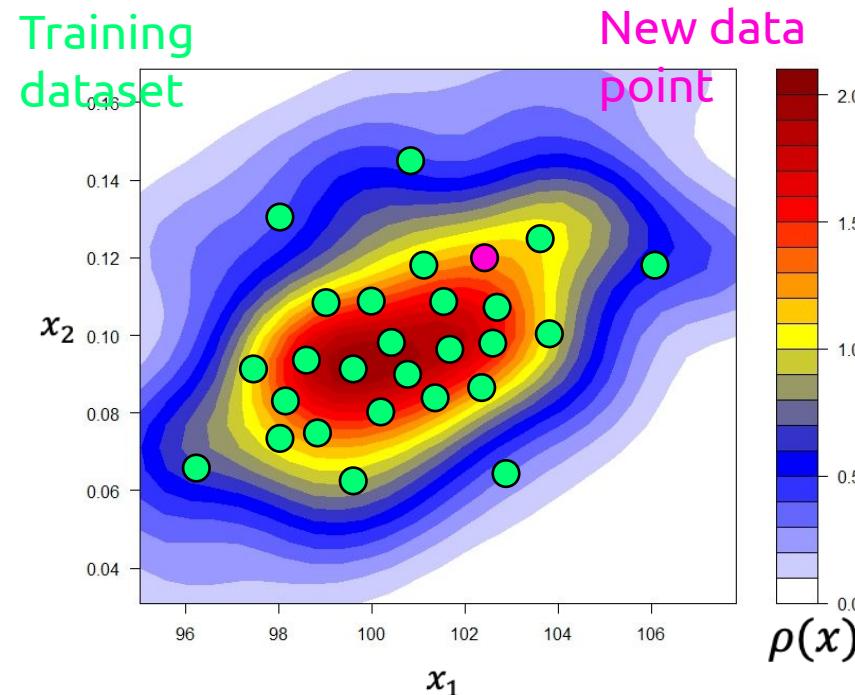
- Generative model: All possible images compete with each other for probability mass
- Intuition: Generation should require deep understanding! Is a dog more likely to sit or stand? How about 3-legged dog vs 3-armed monkey?
- Model can “reject” unreasonable inputs by assigning them small values

Generative Modeling Is:

1. A procedure for (approximately) *sampling* from the distribution from which a dataset was drawn



$$x' \sim p(\cdot)$$

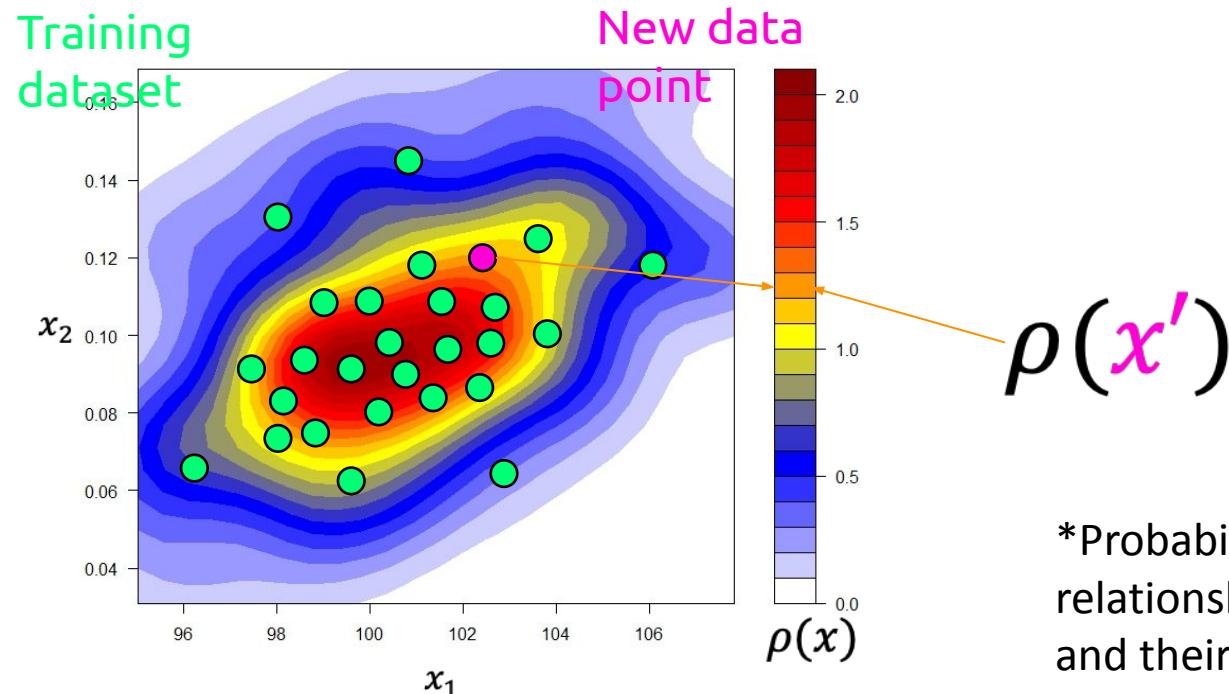


*Probability density is the relationship between observations and their probability.

Generative Modeling:

1. A procedure for (approximately) *sampling* from the distribution from which a dataset was drawn
2. A procedure for (approximately) *evaluating the probability density* of a datapoint under the distribution from which a dataset was drawn

$$x' \sim p(\cdot)$$



*Probability density is the relationship between observations and their probability.

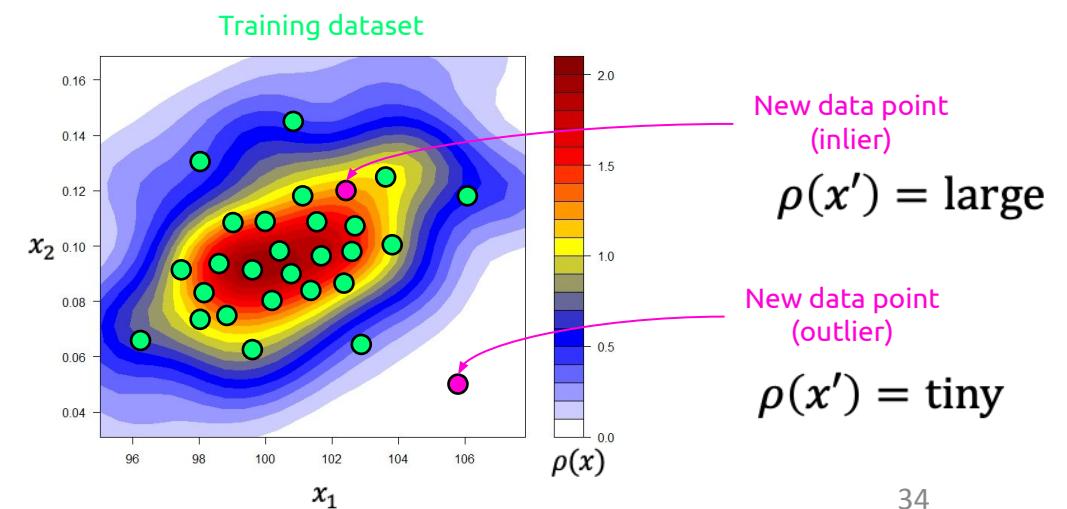
These two views are both useful

1. A procedure for (approximately) *sampling* from the distribution from which a dataset was drawn
2. A procedure for (approximately) *evaluating the probability density* of a datapoint under the distribution from which a dataset was drawn

Application: visual creativity



Application: outlier detection



These two views are both useful

1. A procedure for (approximately) *sampling* from the distribution from which a dataset was drawn
2. A procedure for (approximately) *evaluating the probability density* of a datapoint under the distribution from which a dataset was drawn

Application: things are getting more complicated!

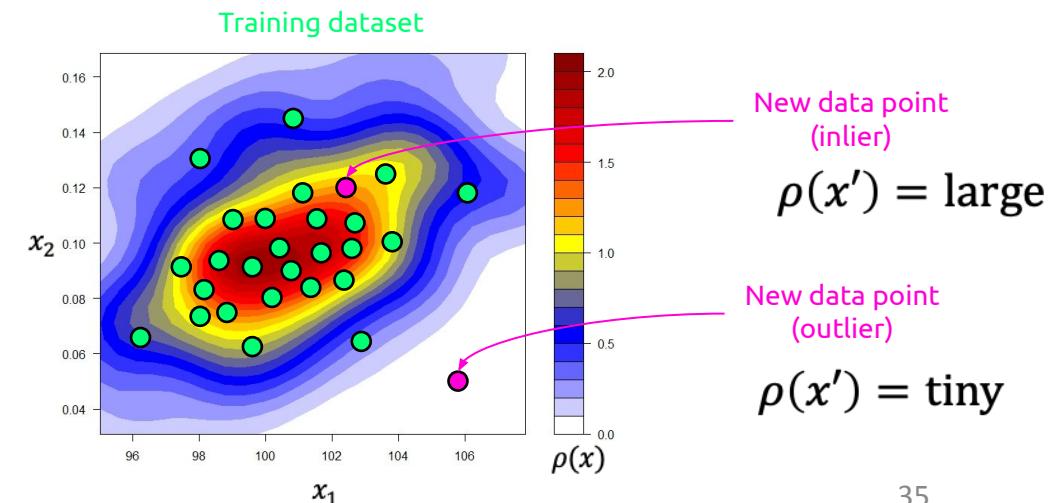
Computer-generated inclusivity: fashion turns to 'diverse' AI models

Fashion brands including Levi's and Calvin Klein are having custom AI models created to 'supplement' representation in size, skin tone and age



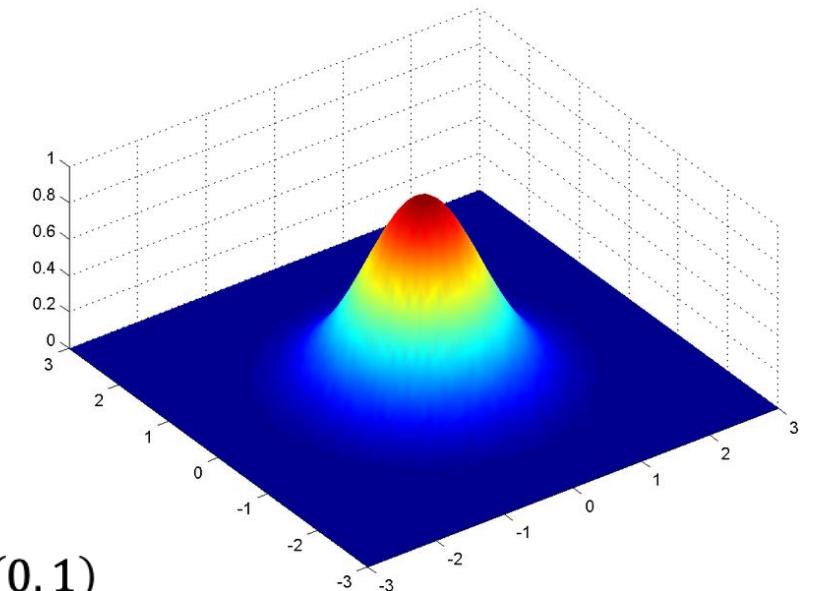
The Guardian

Application: outlier detection



What are some example generative models?

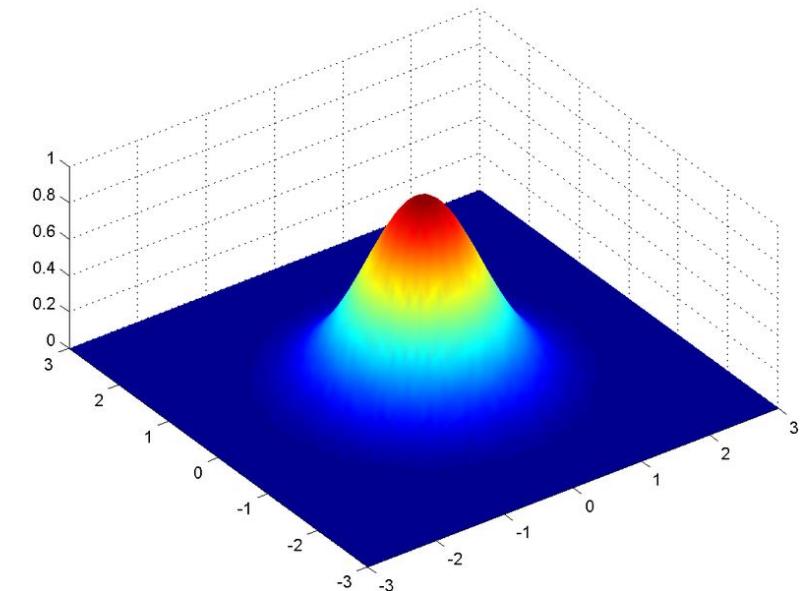
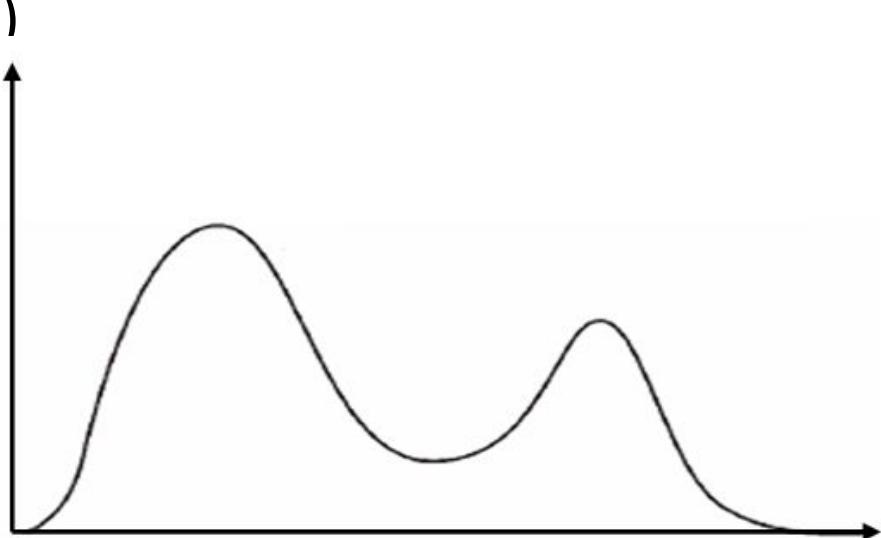
- Any probability distribution can be a generative model
- You already know some of these!
- E.g. The Gaussian Distribution
 - $p(x | \mu, \sigma) = \mathcal{N}(\mu, \sigma)(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$
 - Sampling:
 - [Sample from the unit normal distribution](#) $\rightarrow r \sim \mathcal{N}(0, 1)$
 - Return $\mu + r\sigma$



Disadvantages of Gaussian distribution

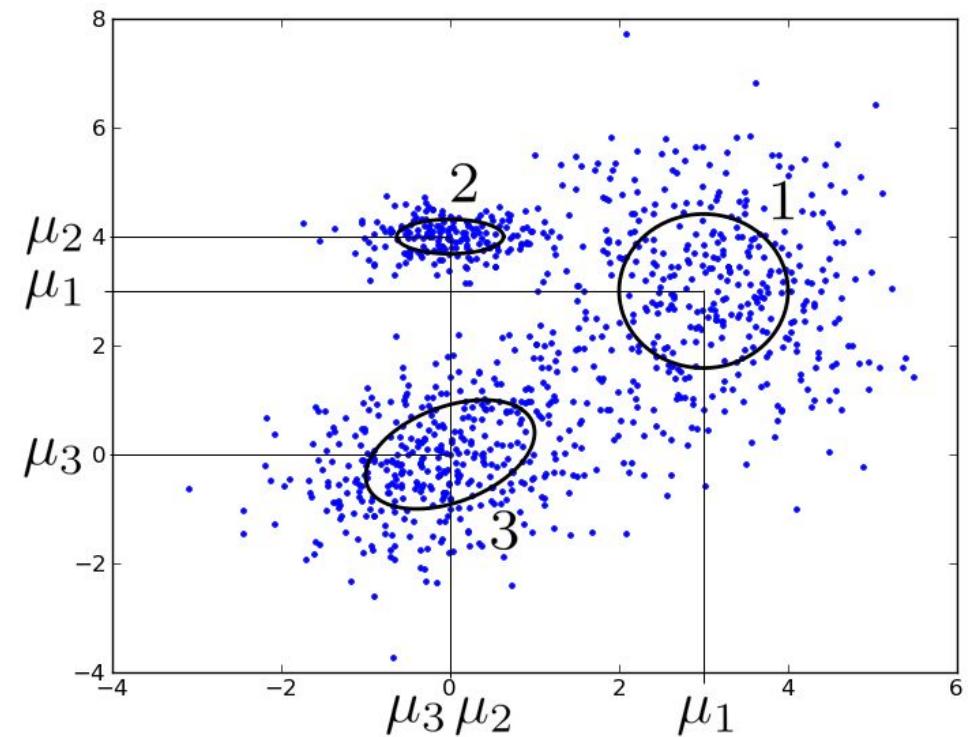
- Can only represent distributions with a single ***mode***

- What if the distribution has multiple “peaks?”
- E.g. book prices (concentrates around different price points if it’s hardcover, paperback, e-book, ...)

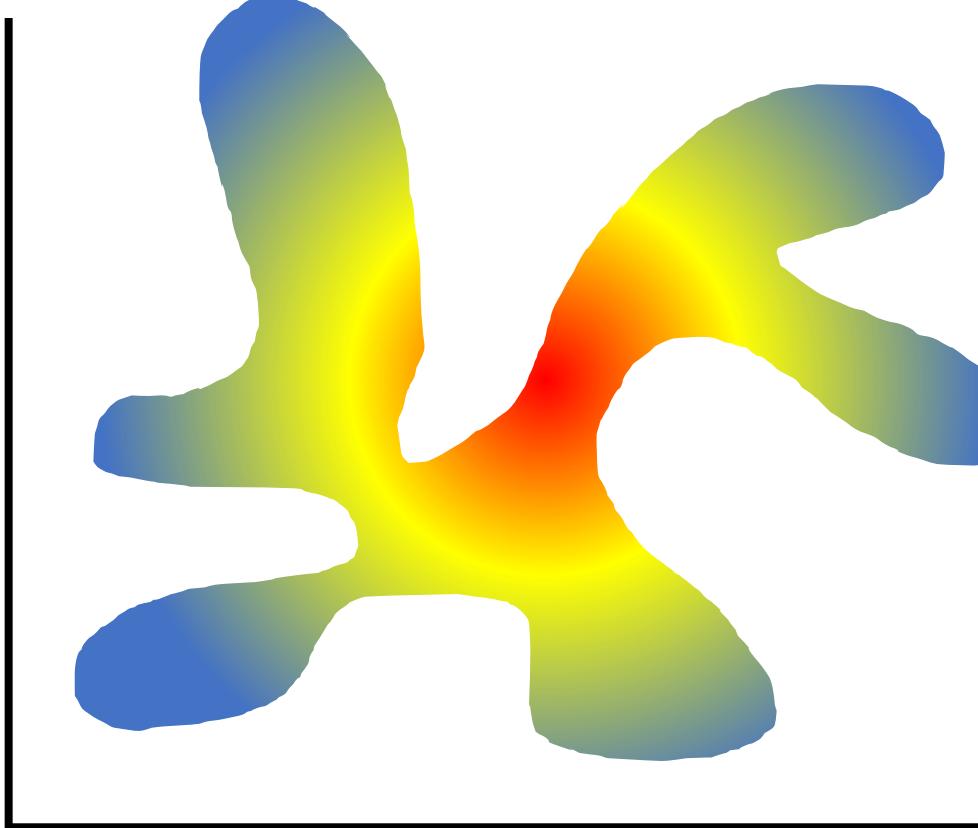


Better: *Mixture* of Gaussians

- A linear combination of multiple individual Gaussian distributions
 - $p(x | w, \mu, \sigma) = \sum_i w_i \mathcal{N}(\mu_i, \sigma_i)(x)$
 - Sampling:
 - Sample from the discrete weight distribution w to choose a Gaussian
 - Sample from that Gaussian as before



What about something like this?



- This doesn't look like a linear combination of Gaussians...
- ...but maybe it can be expressed as a *nonlinear* function of Gaussians?

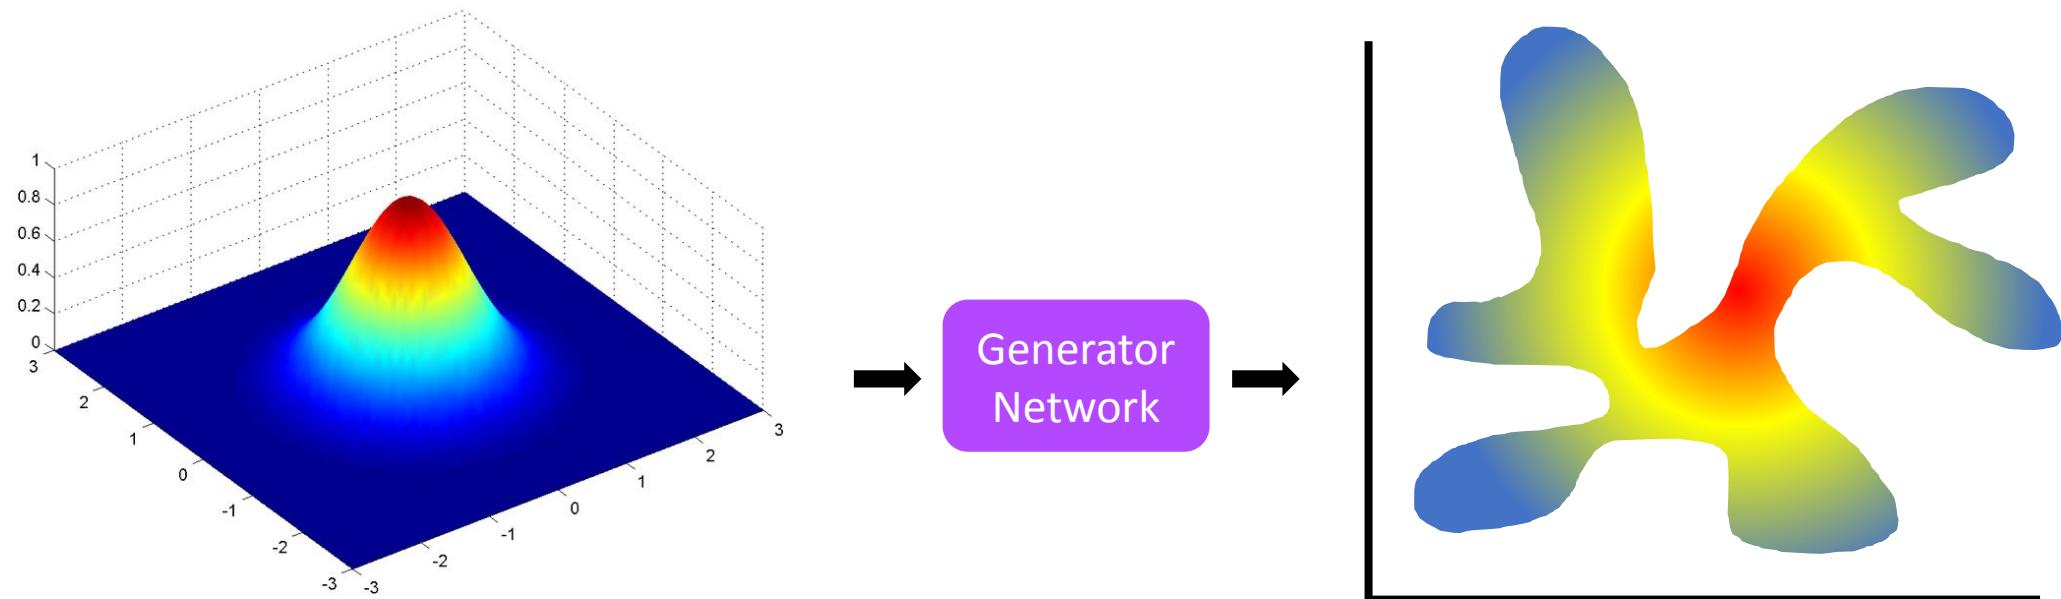
What can we do?

“I hear these neural nets are
pretty good at learning
non-linear functions” 😊



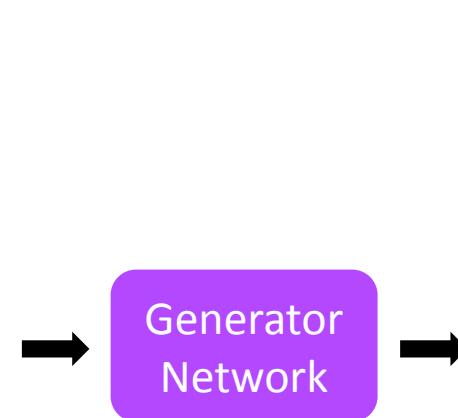
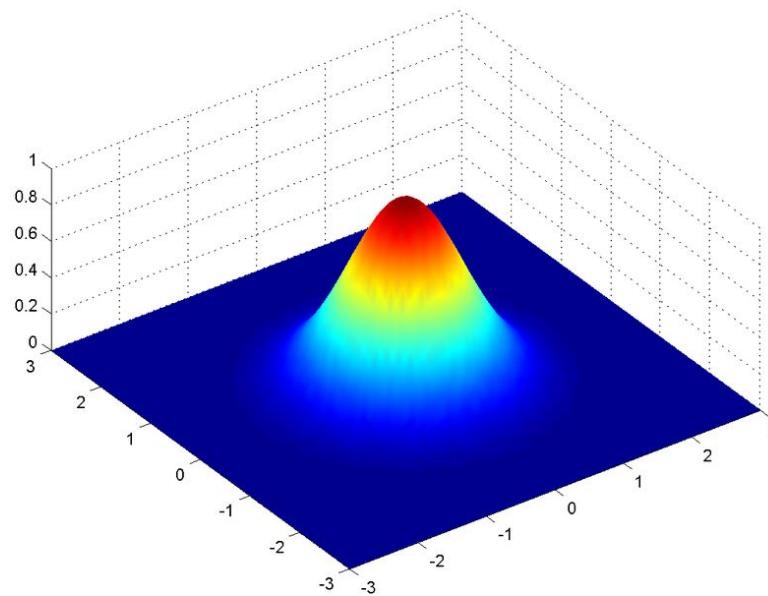
A Neural Generative Model

- Input: a point $z \in \mathbb{R}^n$ drawn from a normal distribution $\mathcal{N}(\mu, \sigma)$
- Output: a point $x \in \mathbb{R}^m$ distributed according to some more complex distribution



A Neural Generative Model

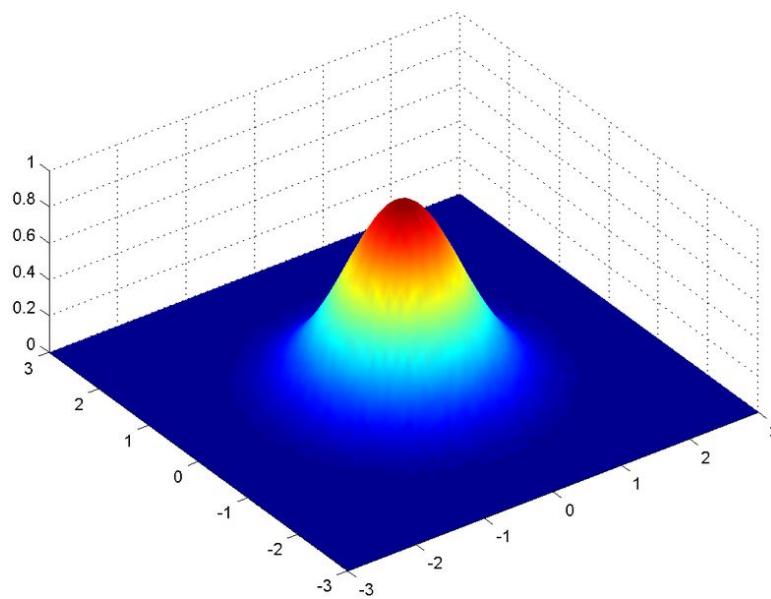
- Input: a point $z \in \mathbb{R}^n$ drawn from a normal distribution $\mathcal{N}(\mu, \sigma)$
- Output: a point $x \in \mathbb{R}^m$ distributed according to some more complex distribution



What are some
distributions that
look like this?

A Neural Generative Model

- Input: a point $z \in \mathbb{R}^n$ drawn from a normal distribution $\mathcal{N}(\mu, \sigma)$
- Output: a point $x \in \mathbb{R}^m$ distributed according to some more complex distribution

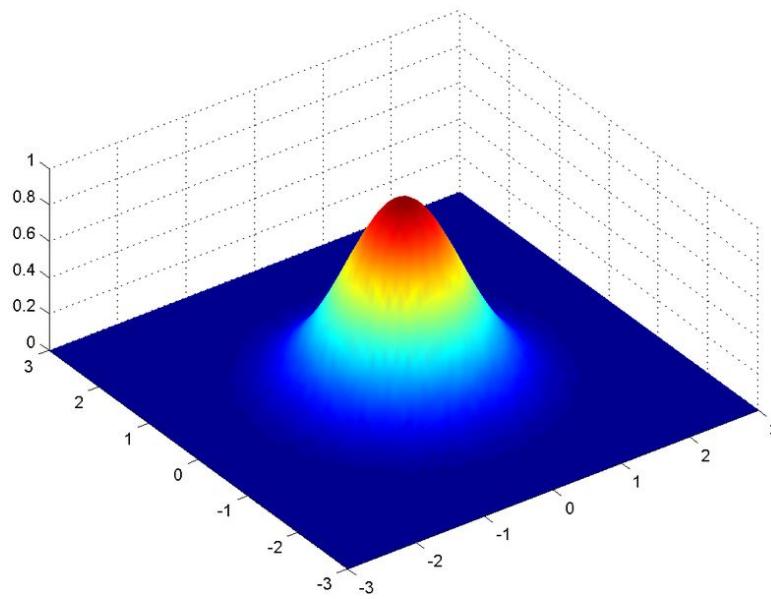


The distribution of
human faces



A Neural Generative Model

- Great! So...how do we train this thing?
 - Let's modify our autoencoder to achieve this

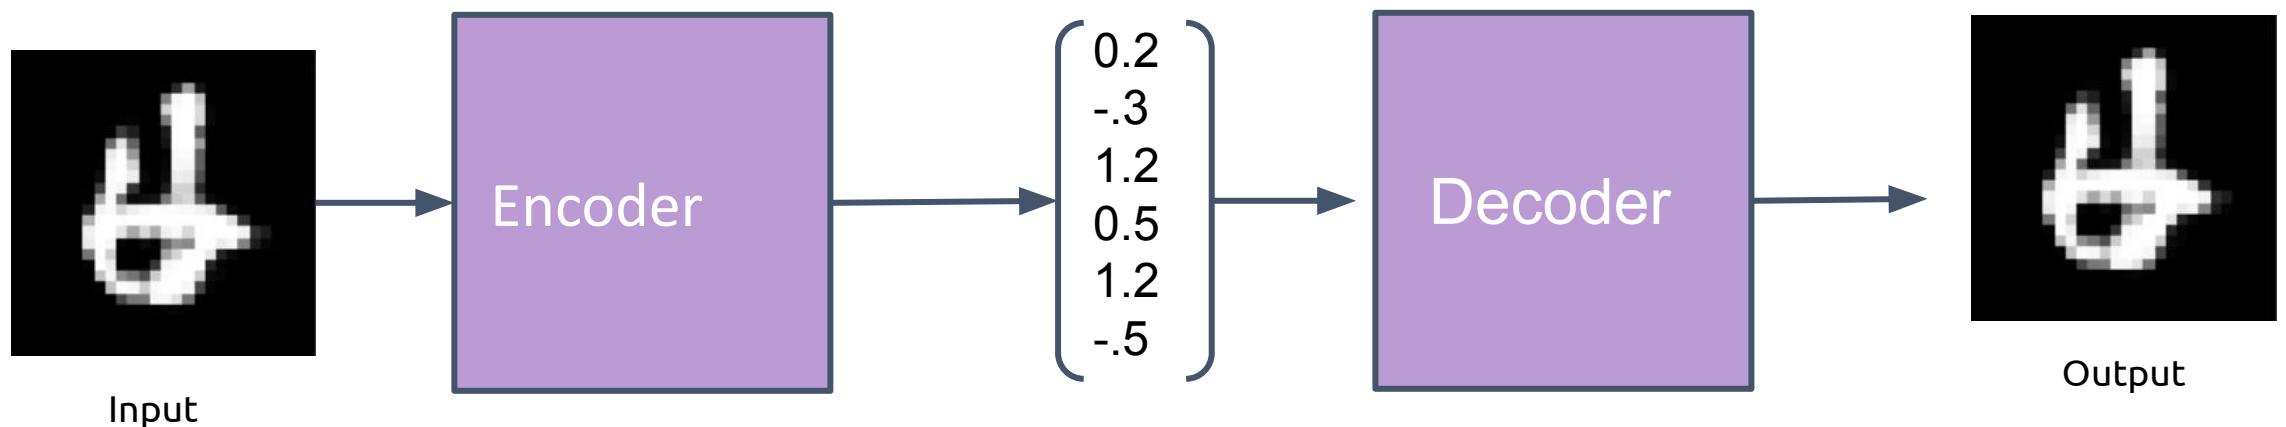


The distribution of human faces

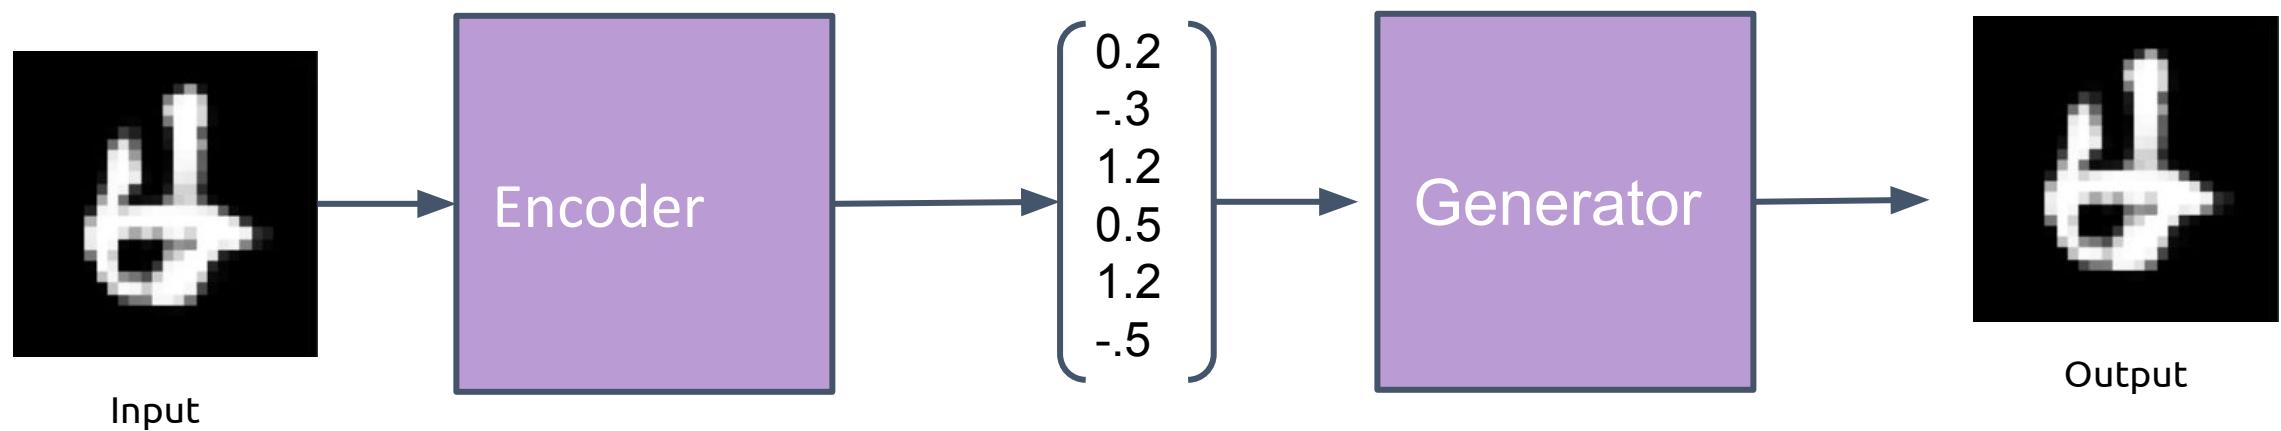


Autoencoder

Let's think for a bit – how to modify the autoencoder to make it a generative model?

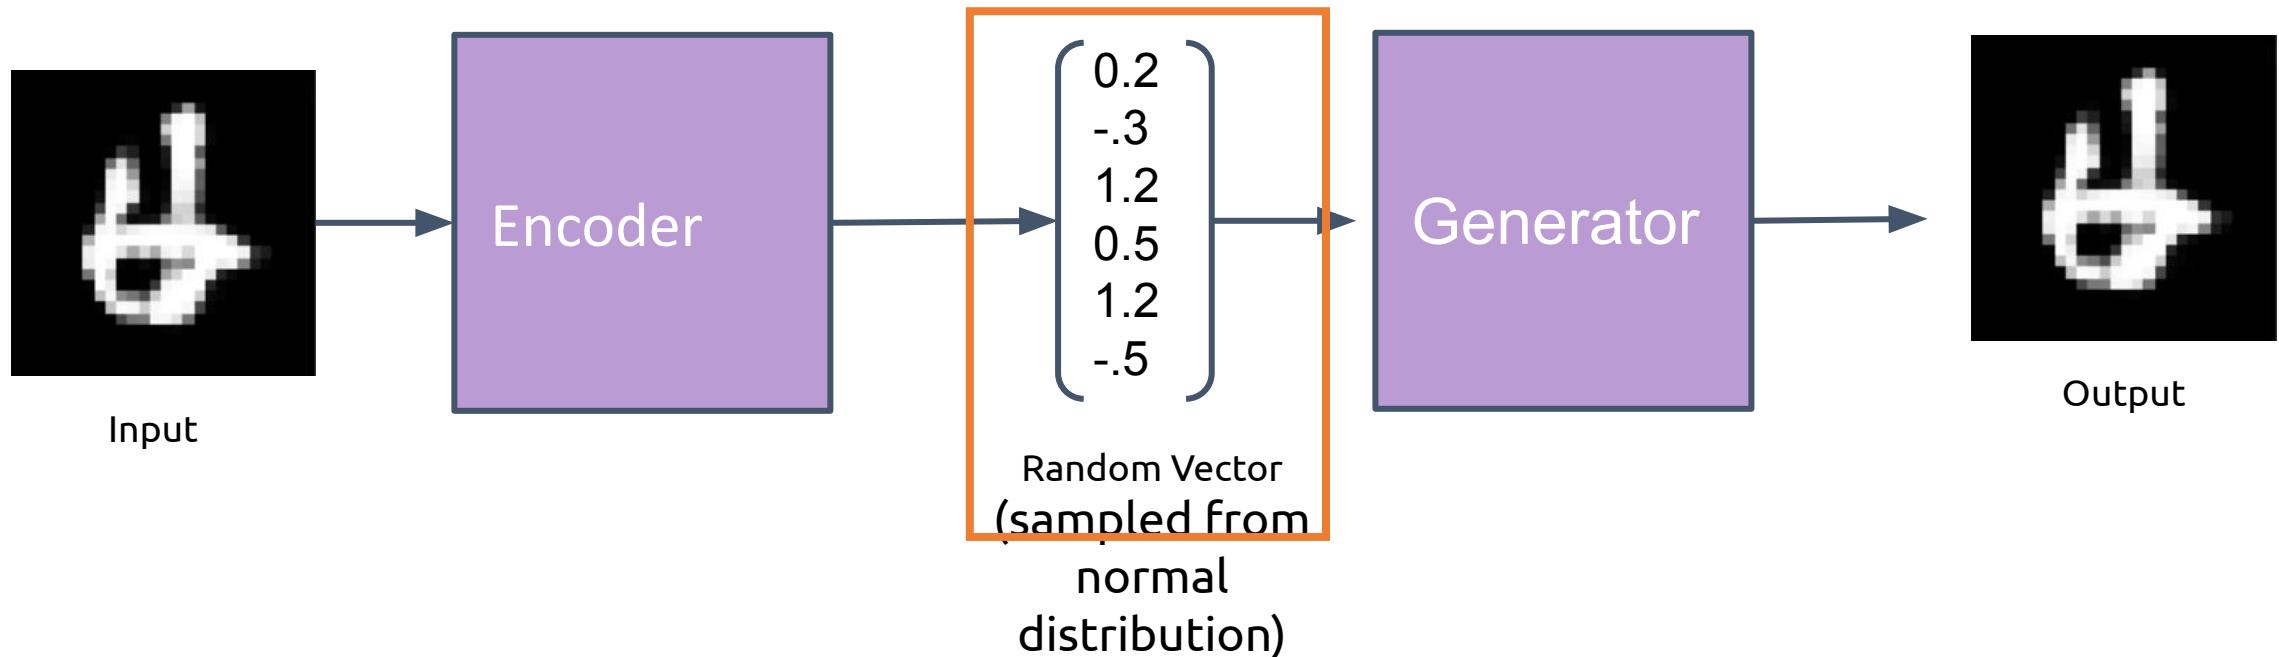


Variational Autoencoders (VAEs)



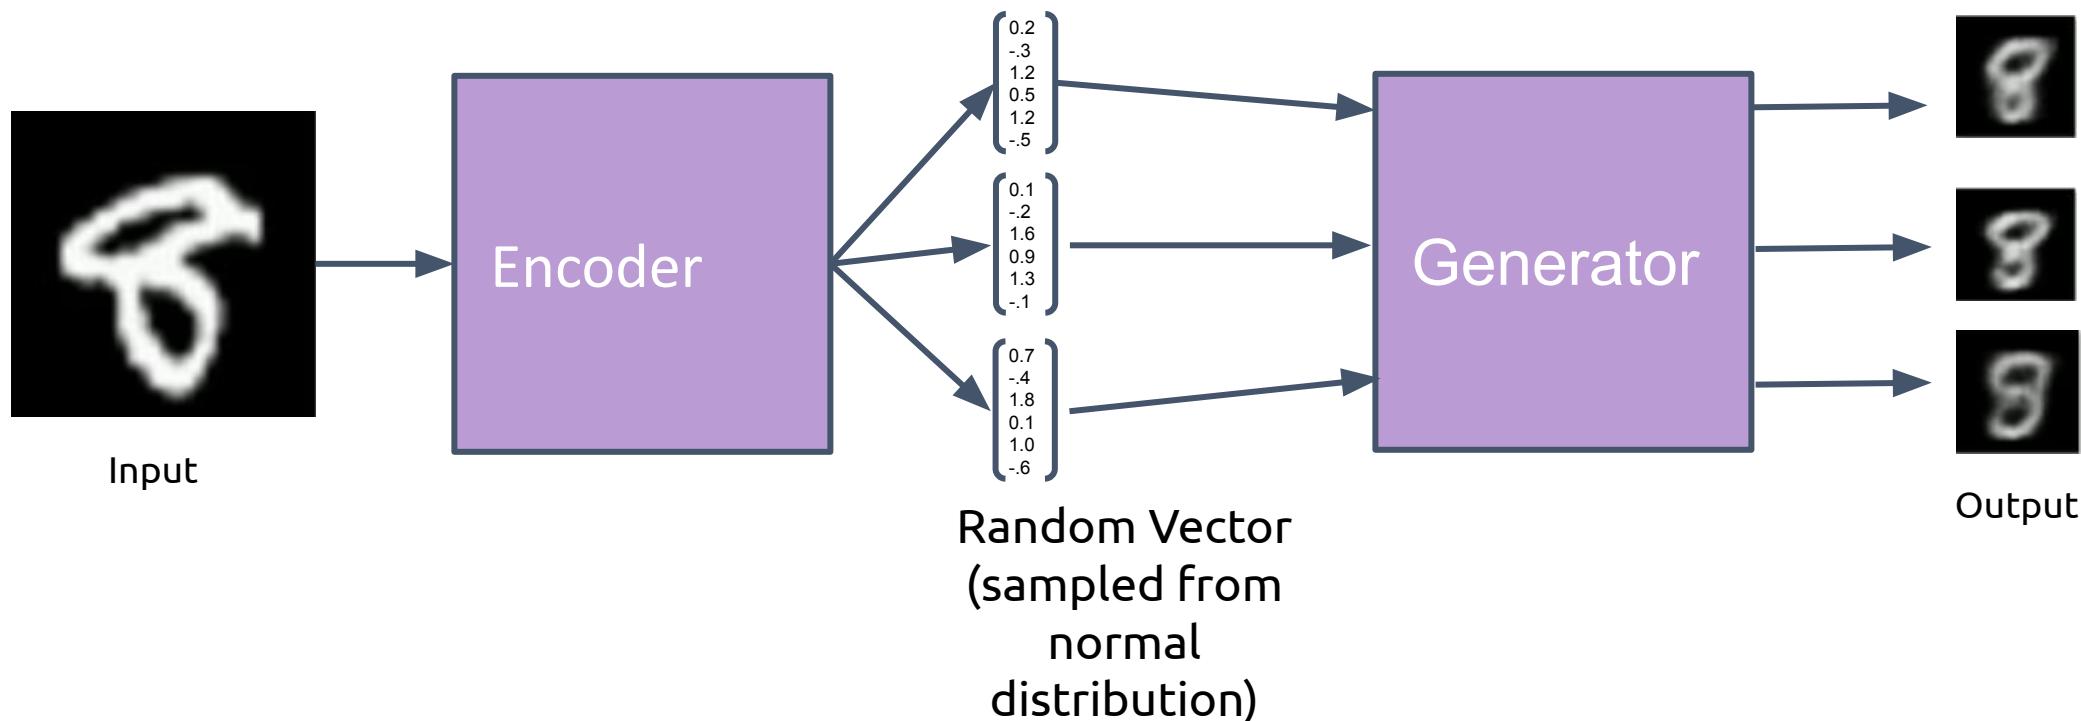
Variational Autoencoders (VAEs)

- This looks almost exactly like an autoencoder...
- ...except that this bottleneck vector is randomly sampled
 - We'll see how in a few slides



Variational Autoencoders (VAEs)

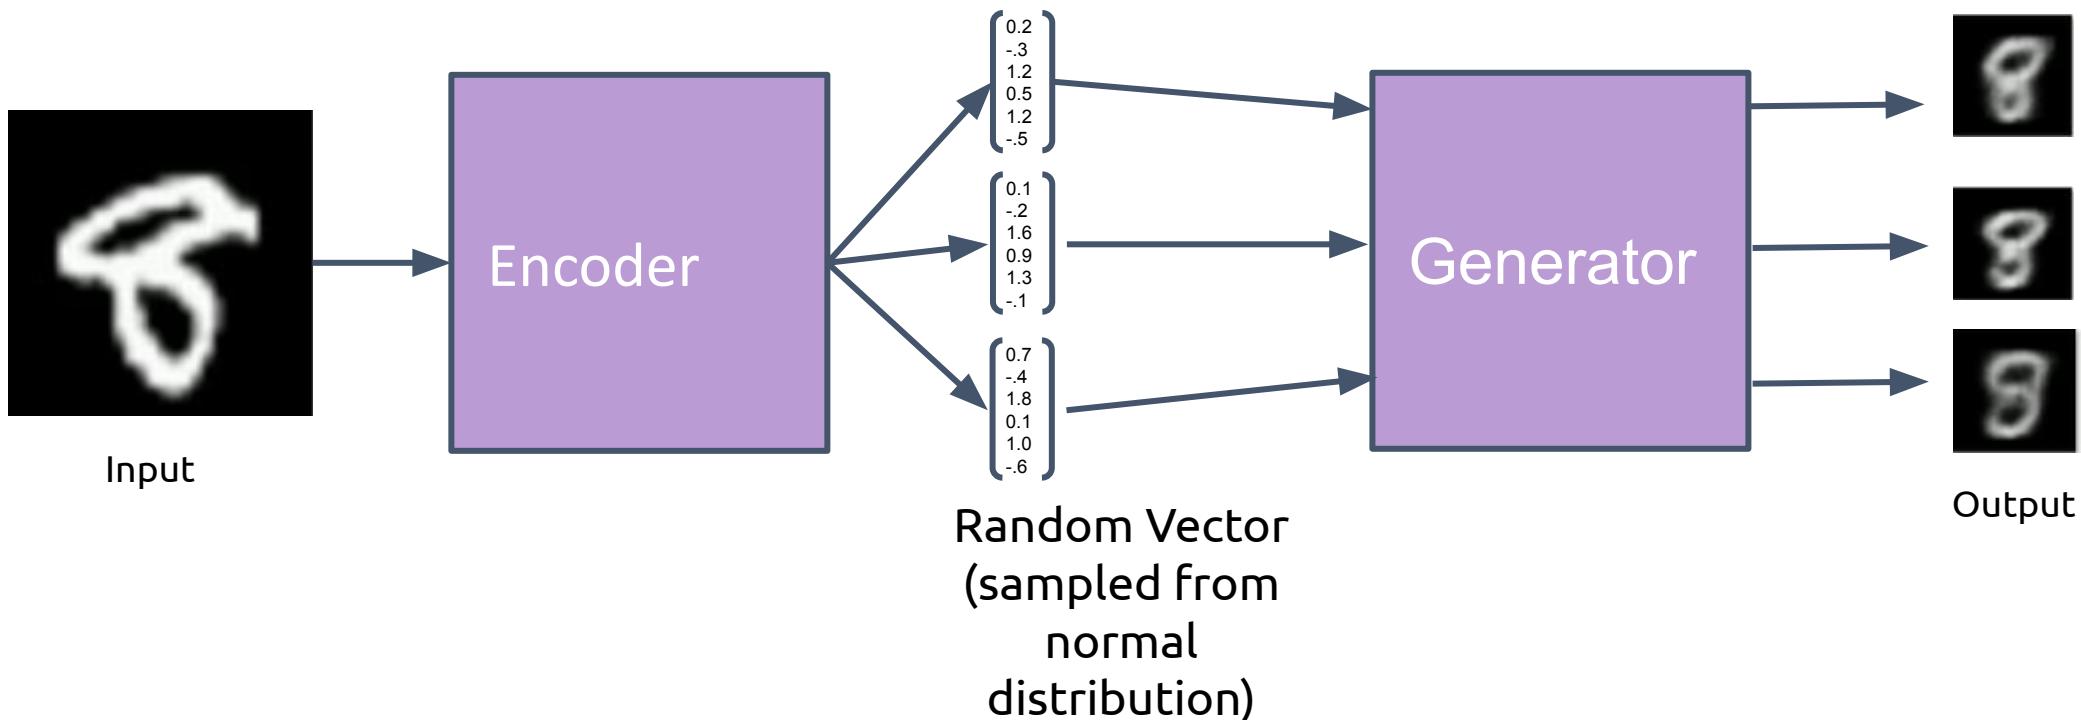
- In fact, the encoder can produce multiple different random vectors...
- ...which then lead to different outputs which are variants of the input



Variational Autoencoders (VAEs)

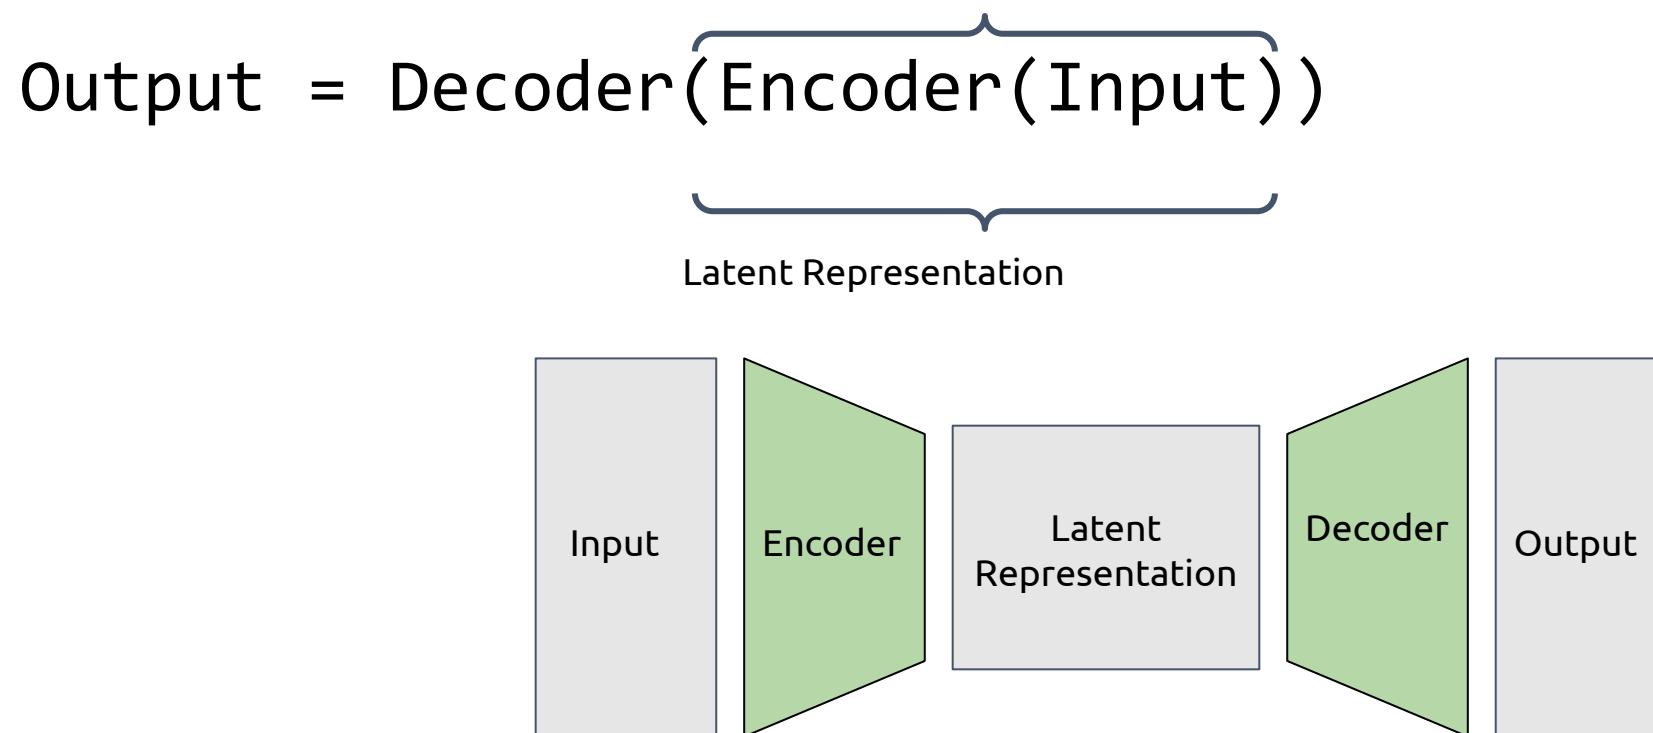
- ***Why do this?***

- We'll see shortly how this setup allows for a nice, stable learning algorithm
- (It's actually just a small modification to how autoencoders are trained)



Building up the VAE Architecture

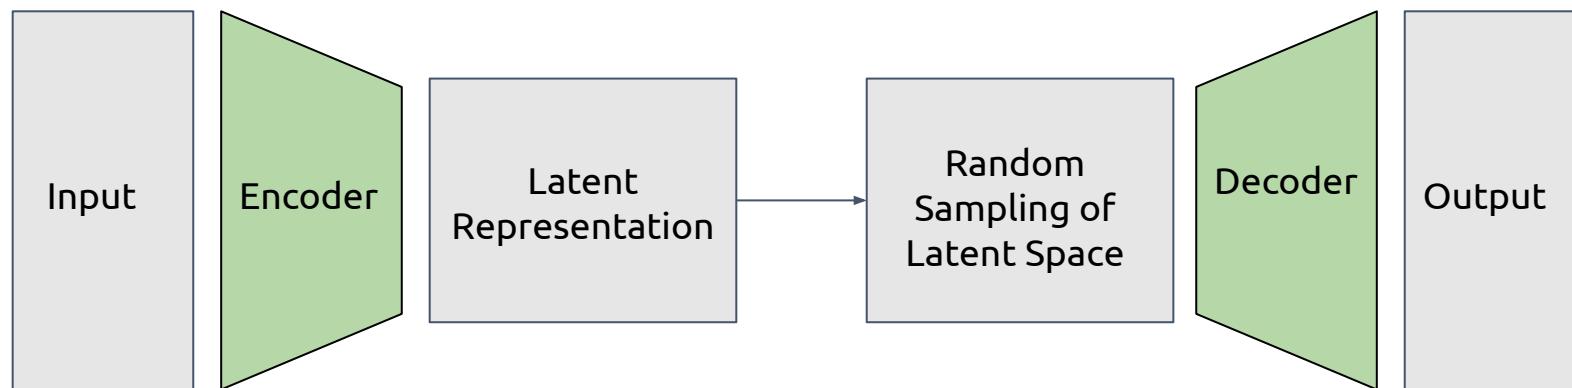
If we were to describe an autoencoder functionally:



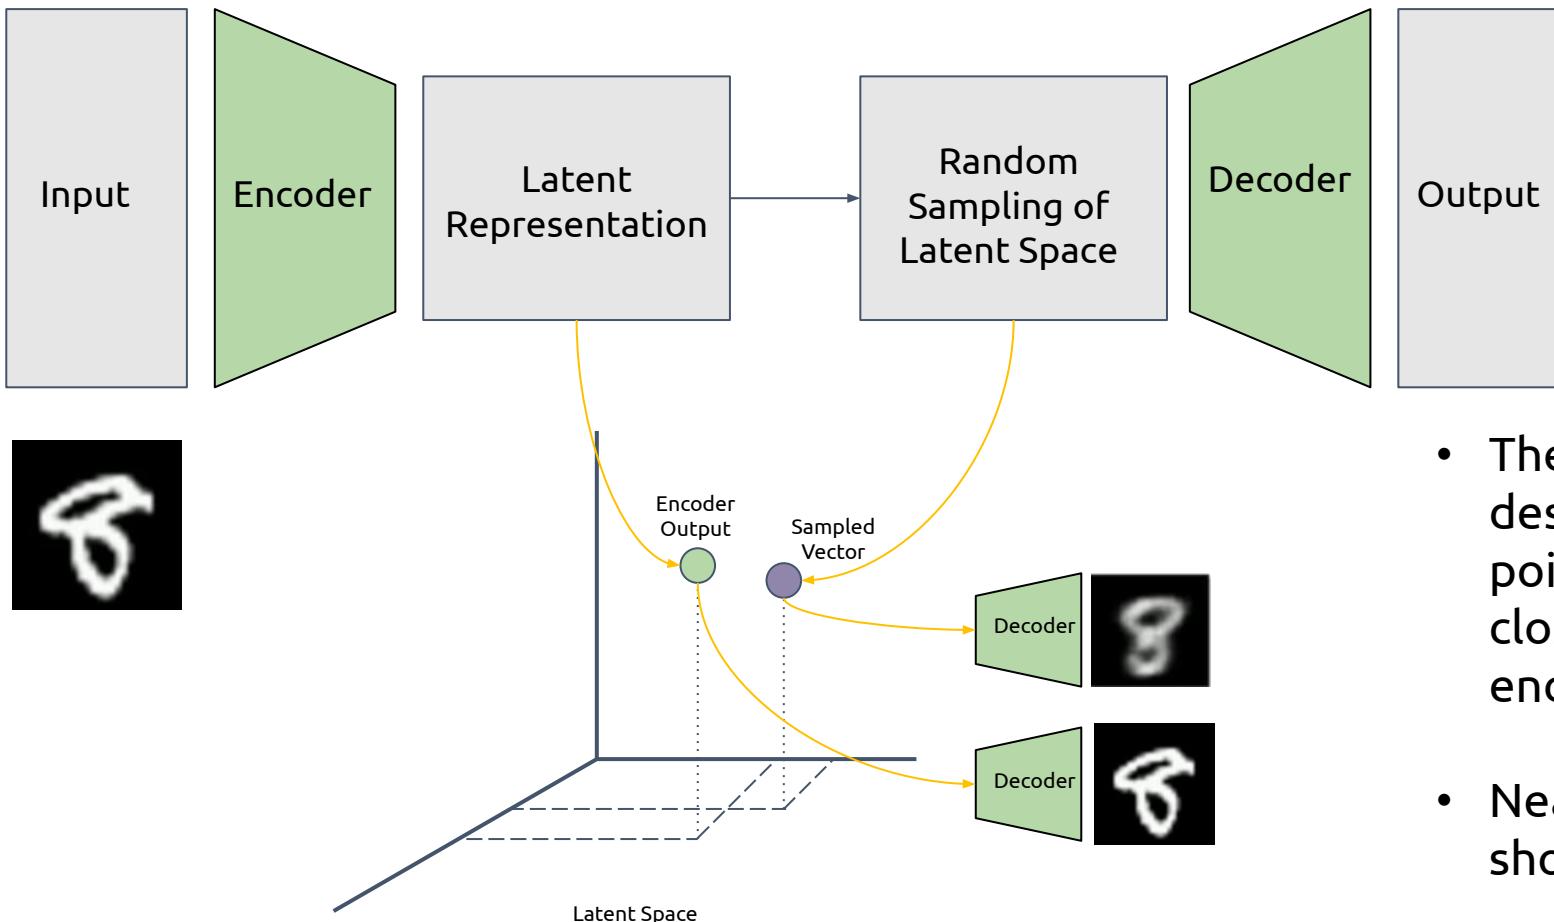
Building up the VAE Architecture

For variational autoencoders, we also do a random sampling operation at the bottleneck

```
Output = Decoder(random_sample(Encoder(Input)))
```



How does random sampling in latent space lead to variation?



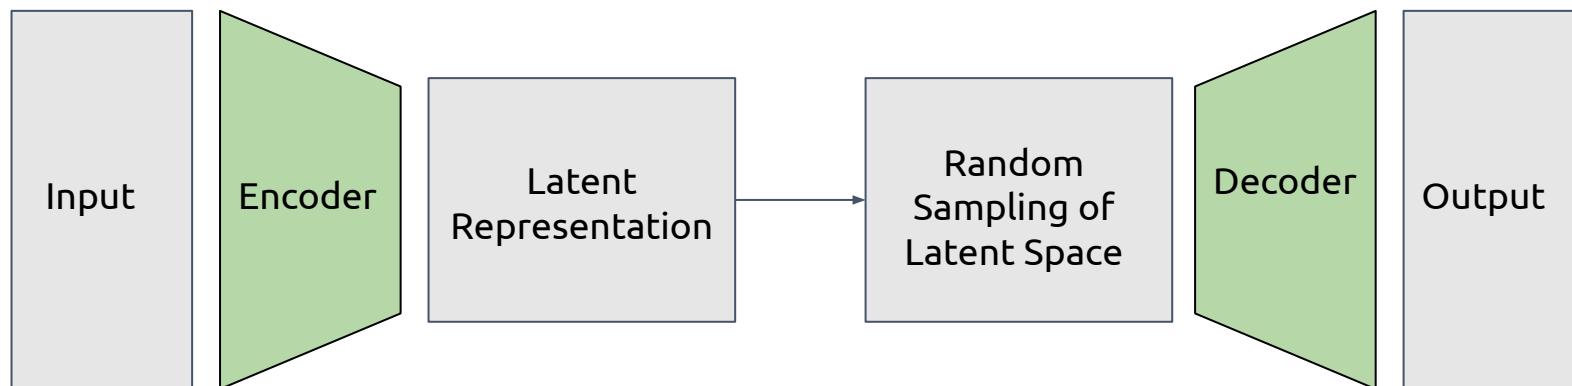
- The random sampling should be designed to produce random points in latent space that are close to the output of the encoder
- Nearby points in the latent space should decode to similar images

How should `random_sample` be defined?

```
Output = Decoder(random_sample(Encoder(Input)))
```

- We want the sample to be close to the encoder output
- One option: sample from a Gaussian centered at Encoder(Input)

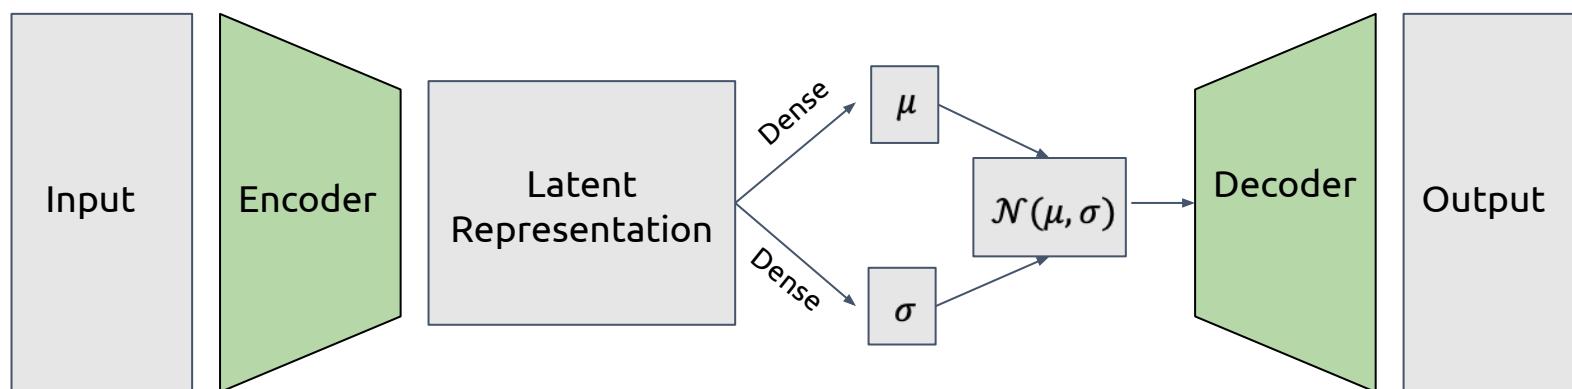
What can we modify?



How should `random_sample` be defined?

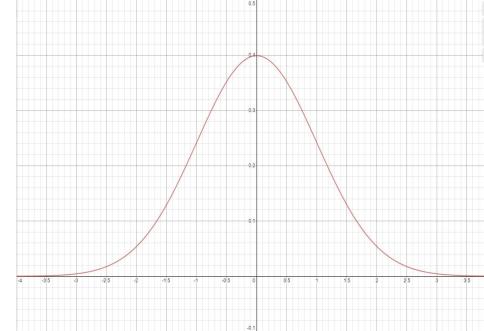
```
Output = Decoder(random_sample(Encoder(Input)))
```

- We want the sample to be close to the encoder output
- One option: sample from a Gaussian centered at Encoder(Input)
- Use two dense layers to convert the encoder output into the mean and standard deviation of the Gaussian

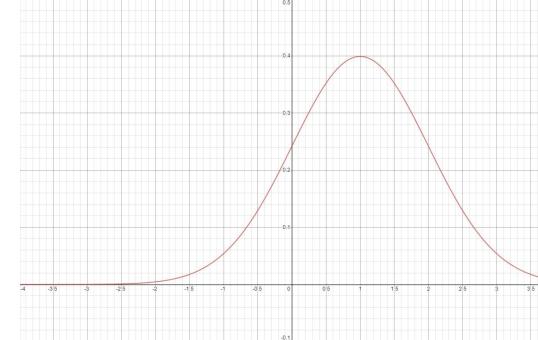




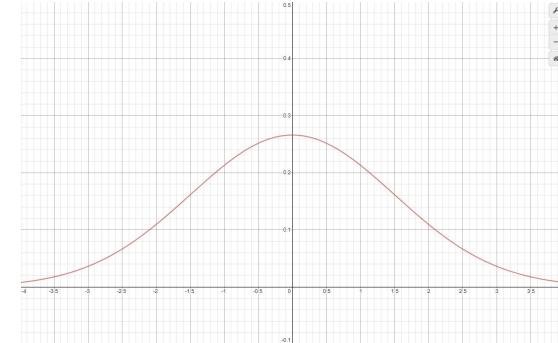
How should `random_sample` be defined?



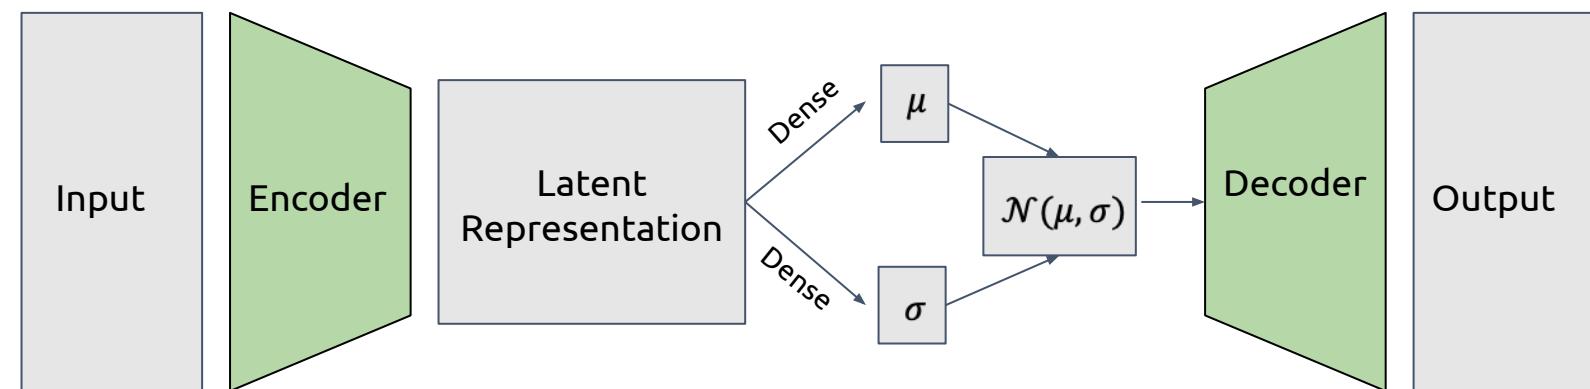
$$\begin{aligned}\mu &= 0 \\ \sigma &= 1\end{aligned}$$



$$\begin{aligned}\mu &= 1 \\ \sigma &= 1\end{aligned}$$



$$\begin{aligned}\mu &= 0 \\ \sigma &= 1.5\end{aligned}$$

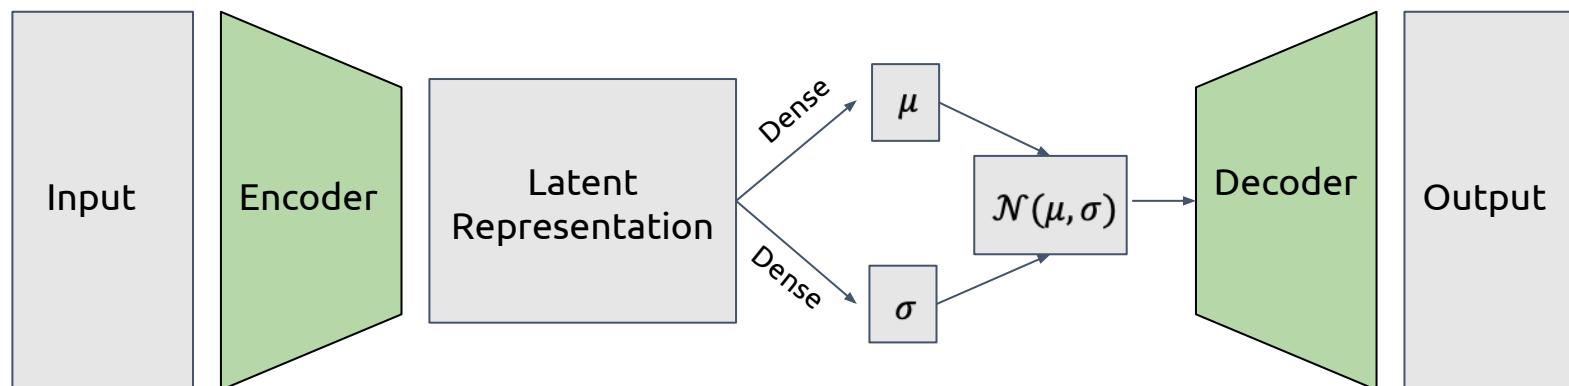


Training a VAE

Two goals:

1. Reproduce an output similar to the input (Input \approx Output)
2. Have some variation in our output (Input $\not\approx$ Output)

- Seems like two conflicting goals!
- How do we resolve these two goals?



Weighted Combination of Losses

L_1 = loss associated with producing output similar to input

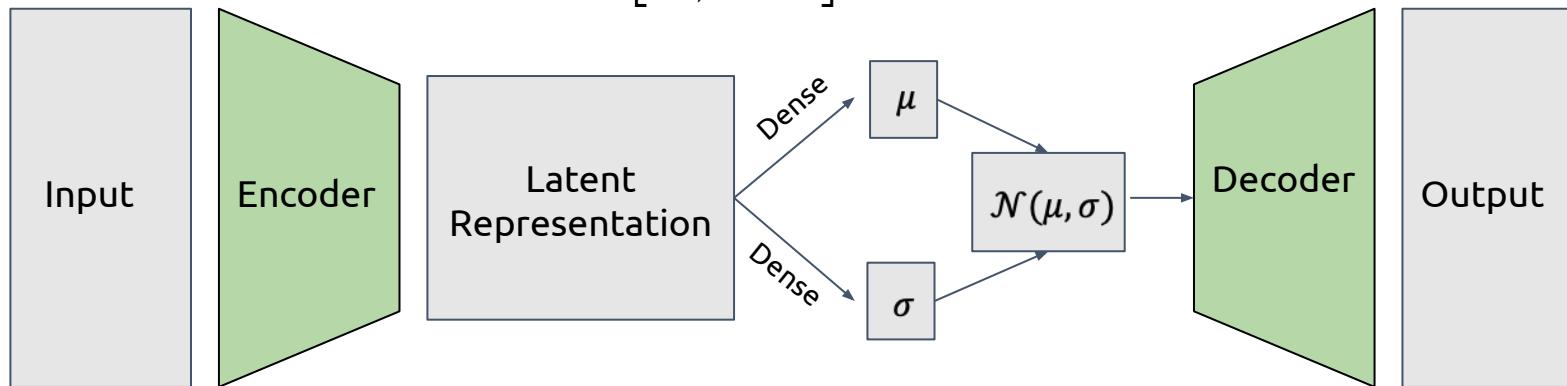
L_2 = loss associated with producing output with some variation to input

$$L = L_1 + \lambda L_2$$

Total Loss:

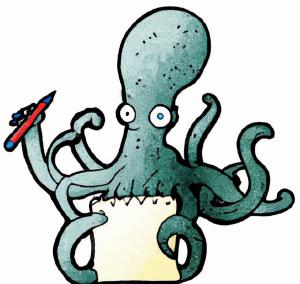


$$\lambda \in [0, \infty]$$



Recap

Generative
Modeling

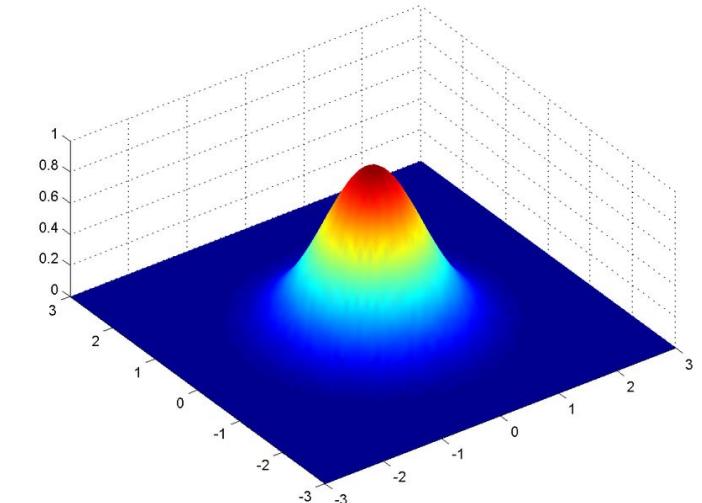


Variational
Autoencoders
(VAEs)

Convolutional AEs

Generative modeling – formulation
and applications

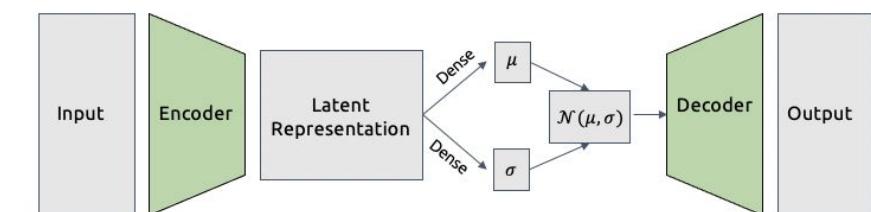
Probability distributions = generative models



Generative modeling for complex
distributions

Modifying AEs

VAE architecture

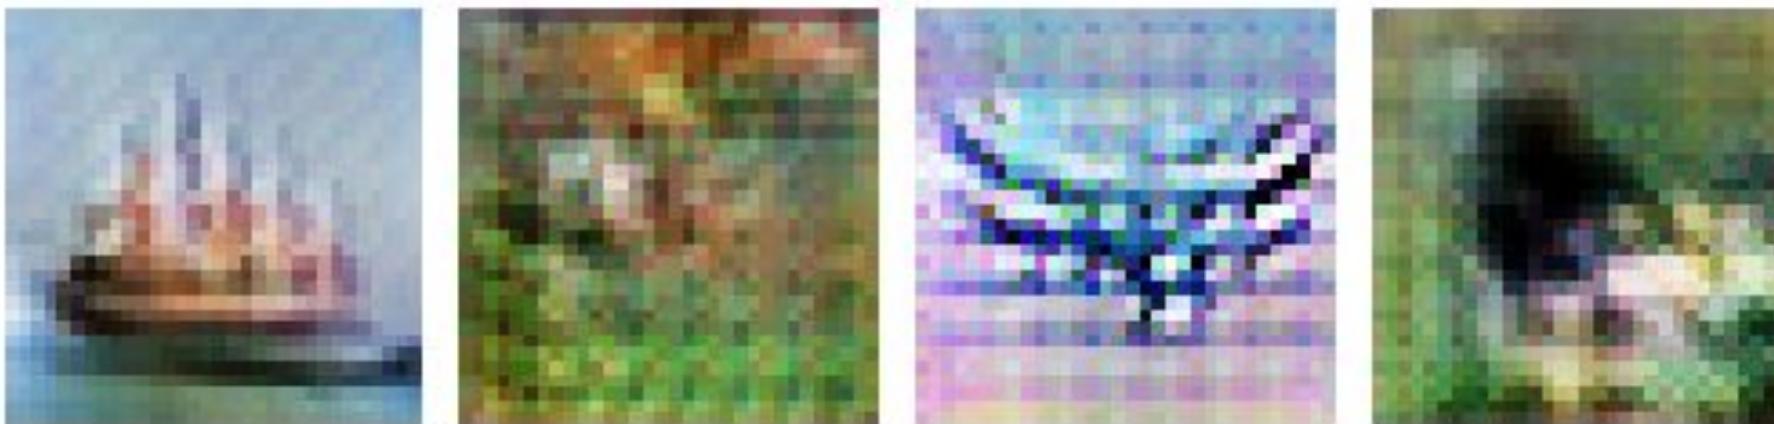


Extra material

[More reading on Transpose Convolution](#)

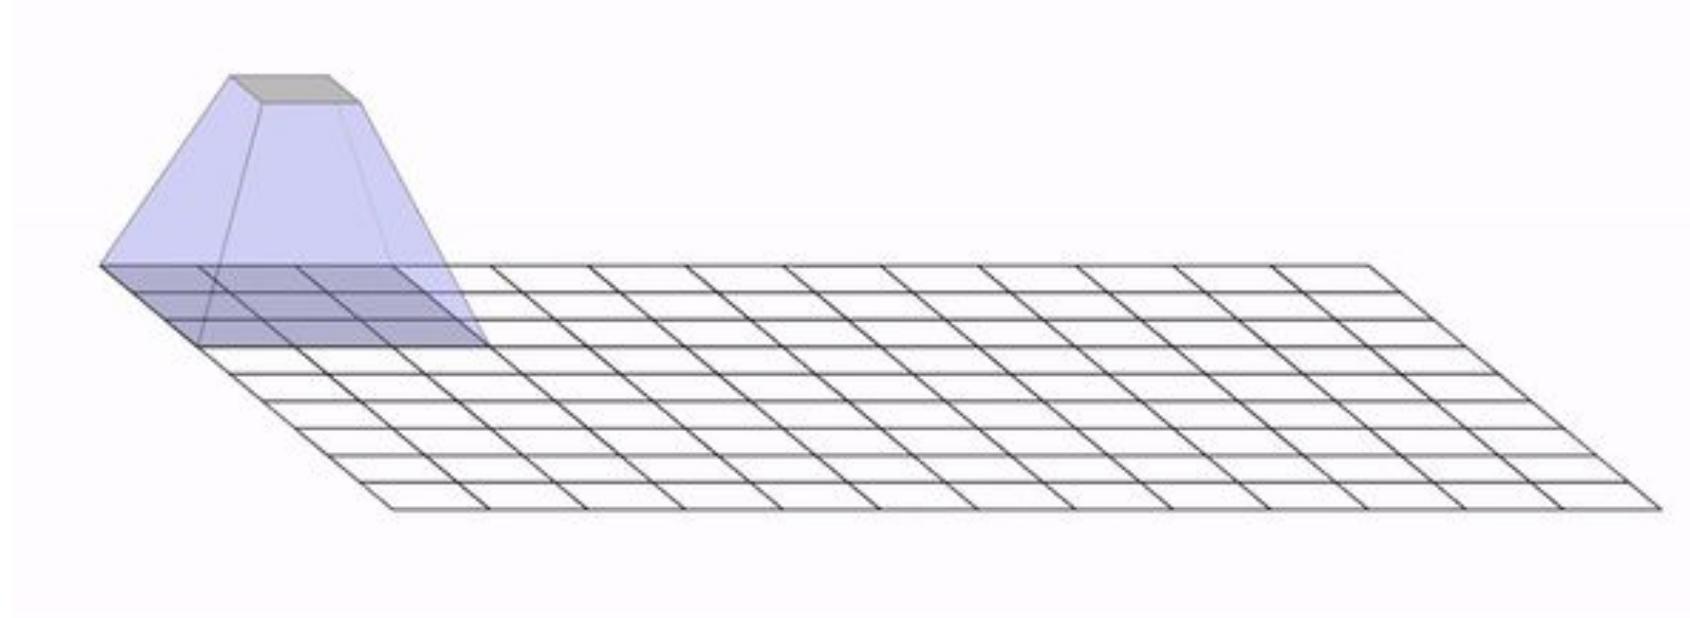
Caution: Checkerboard Artifacts

- Transpose convolution causes artifacts in output images



Caution: Checkerboard Artifacts

- Transpose convolution causes artifacts in output image
- Why? Some pixels get written to more often than others
- Is there a better way to upsample?



Eliminating checkerboard artifacts

Step 1: Upsample using nearest neighbor interpolation:

1	2
3	4

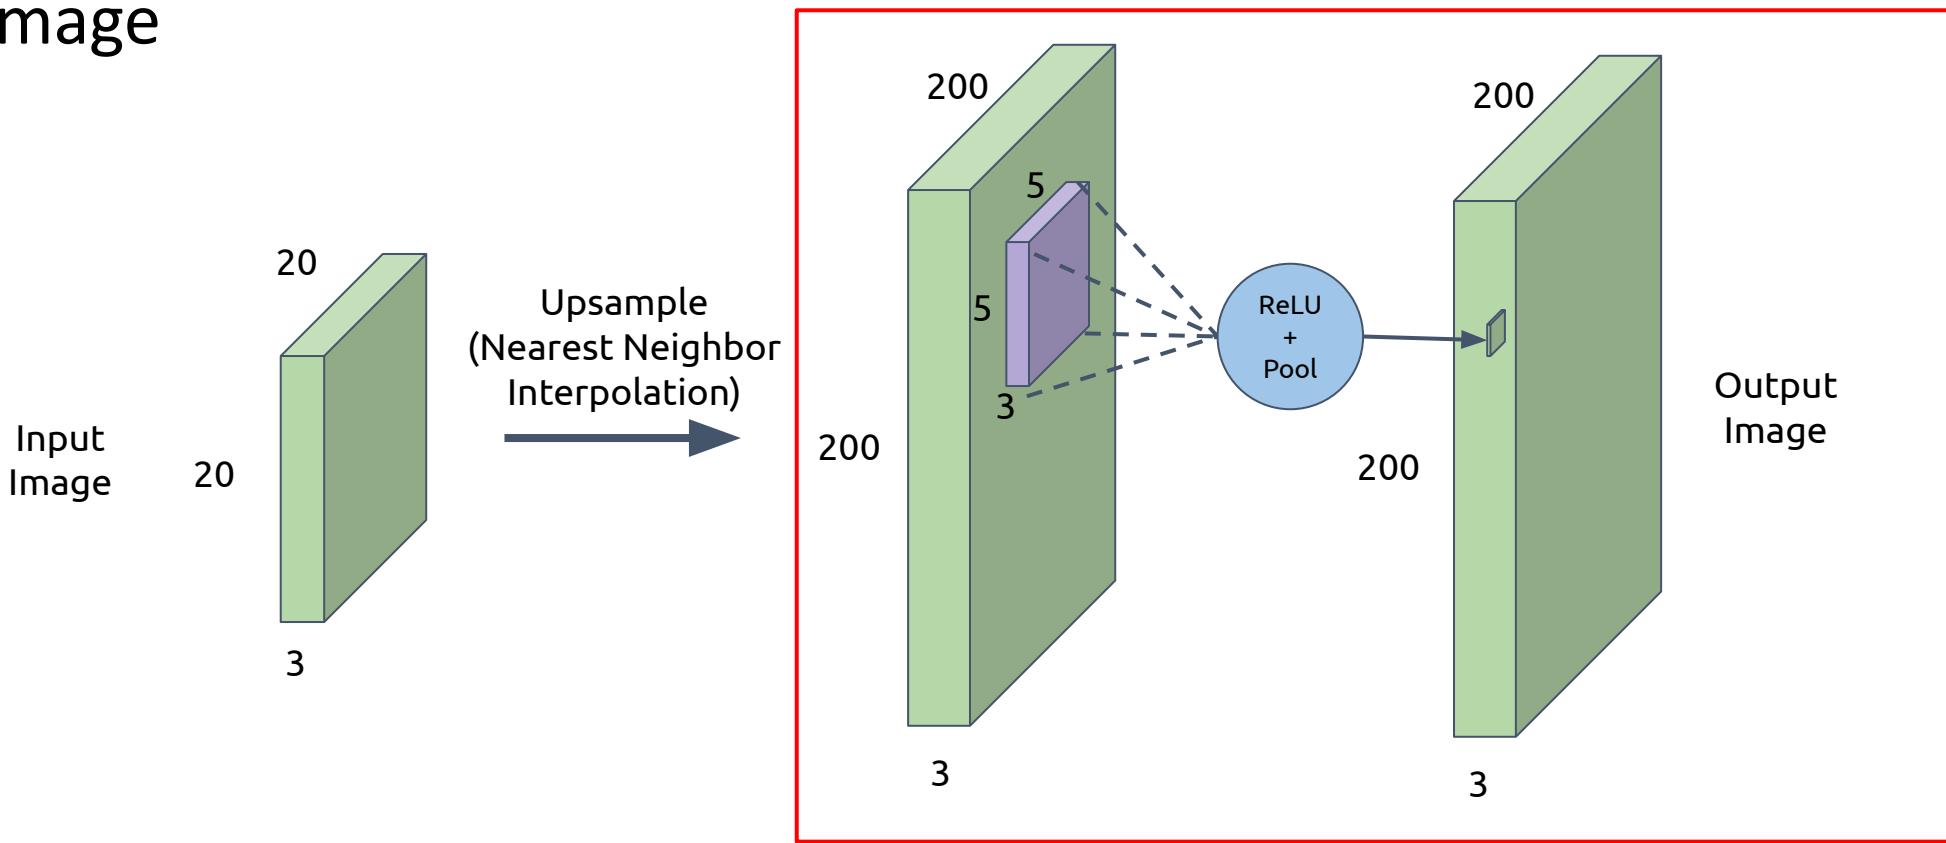
Pixels in upsampled
image are assigned pixel
value of CLOSEST pixel in
original image



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Eliminating checkerboard artifacts

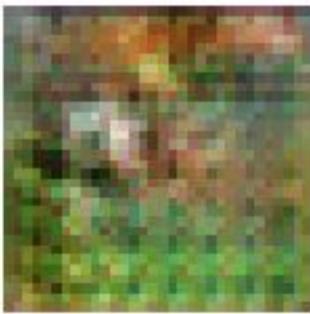
Step 2: Perform a convolution with SAME padding on the upsampled image



Dealing With it in Tensorflow

```
# Layer to upsample the image by a factor of 5 in x and y using nearest  
# neighbor interpolation  
  
tf.keras.layers.UpSampling2D(size=(5, 5), interpolation='nearest')  
  
# Do a convolutional layer on the result  
  
tf.keras.layers.Conv2D(filters = 1, kernel_size = (10,10), padding = "SAME")
```

Checkerboard Artifacts Resolved



With Transpose
Convolution



With Resize +
Convolution