

**DIFFUSION MODELS, WHAT IS THAT ALL ABOUT**



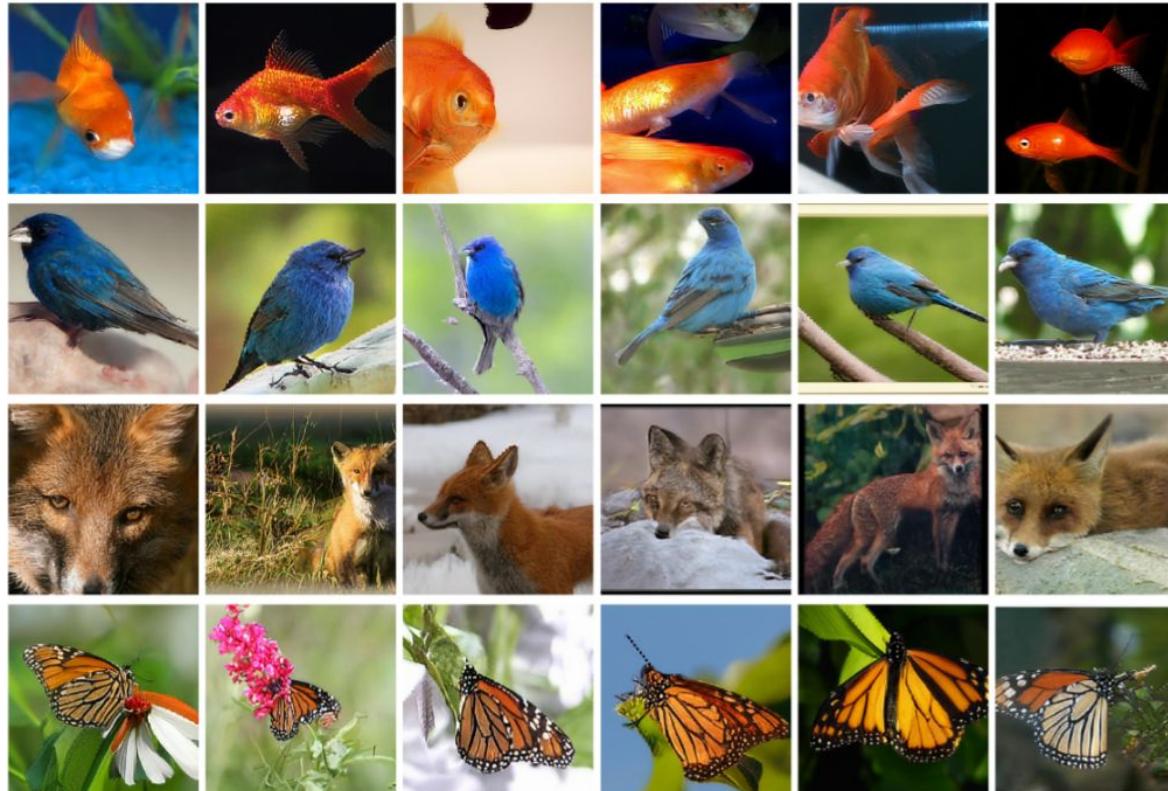
**IS IT GOOD? OR IS IT WACK?**

*Diffusion models are really good at learning conditional distributions.*

$$p(x \mid y)$$

# Use Case: Class-Conditioned Generation

$p(image \mid class\_label)$



source: [Image Super-Resolution via Iterative Refinement](#)

# Use Case: Text-to-Image Generation

$p(image \mid text\_caption)$

“a painting of a fox sitting in a field at sunrise in the style of Claude Monet”



Parti (but pretend it is ImageN)



StableDiffusion

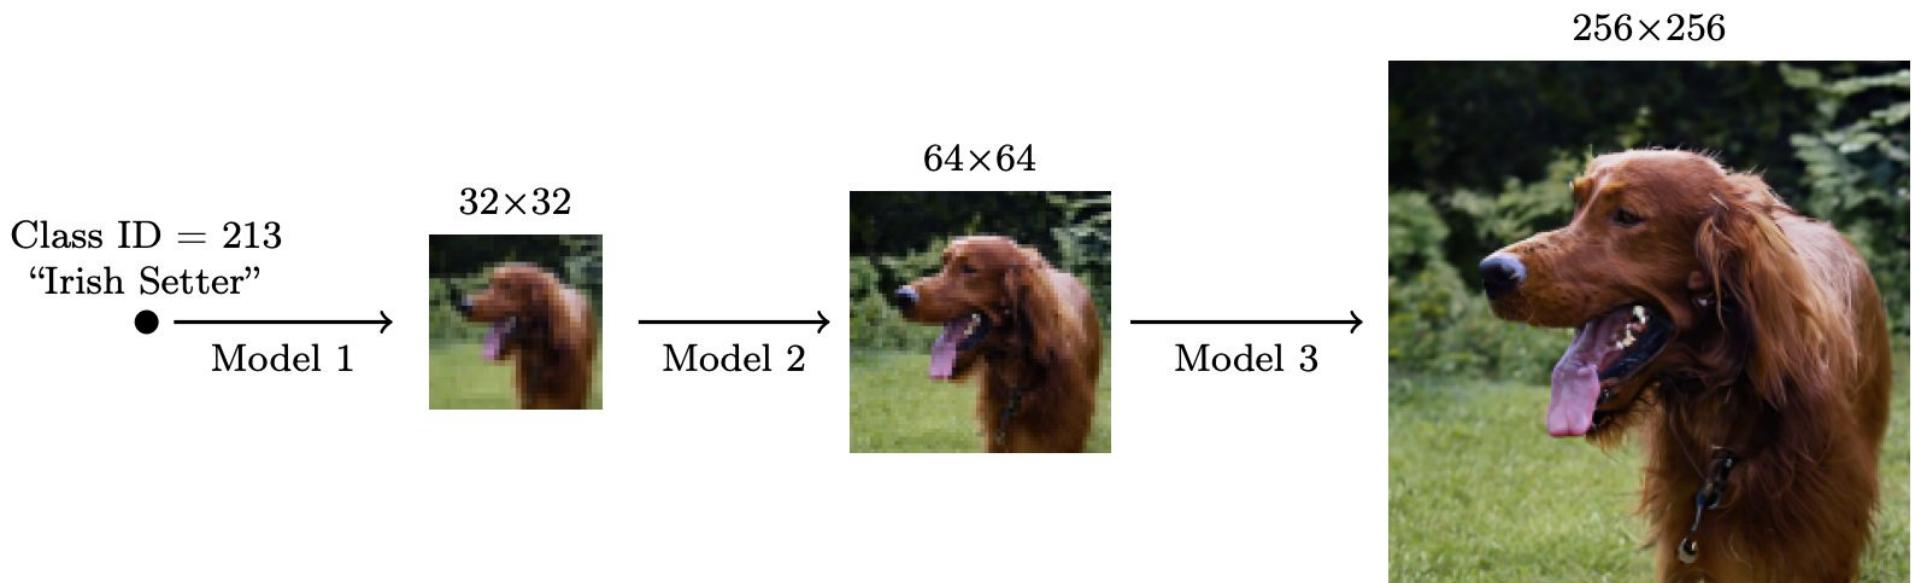


Dall-E 2.0

source: [ImageN](#), [StableDiffusion](#), [Dall-E 2.0](#)

# Use Case: Super Resolution

$p(image \mid low\_res)$



source: [Cascaded Diffusion Models](#)

# Recap: Generative Modeling

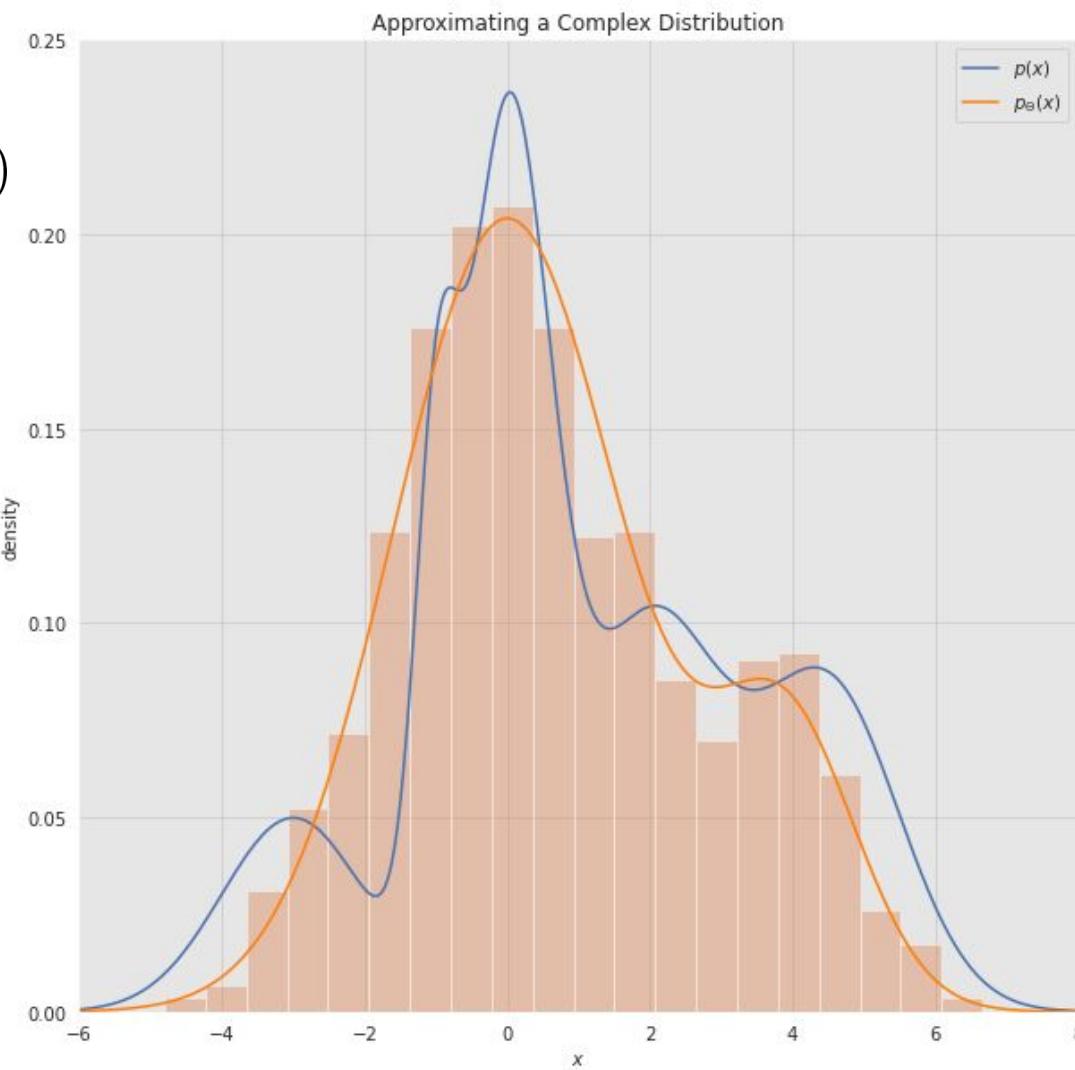
Recall the goal of generative modeling - learning a *model* of a distribution from which we can *generate* new samples.

Given  $\mathbf{x} \sim p(\mathbf{x})$  we might want to learn  $p_\theta(\mathbf{x}) \approx p(\mathbf{x})$  (*modeling*)

Then, we can generate new samples  $\mathbf{x}^* \sim p_\theta(\mathbf{x})$  (*generation*)

Why is this useful?

Given  $x \sim p(x)$



# Generative Modeling: Themes

What are some common themes of generative modeling?

- We want to learn some **complex** distribution  $p_\theta(\mathbf{x}) \approx p(\mathbf{x})$
- But we only have access to some **simple** distributions (such as Gaussians)



# Generative Modeling: Themes

What are some common themes of generative modeling?

- We want to learn some **complex** distribution  $p_\theta(\mathbf{x}) \approx p(\mathbf{x})$
- But we only have access to some **simple** distributions (such as Gaussians)

**Idea:** Let's learn a complex function (aka a neural network) to transform a simple distribution sample into a complex one!

- Gaussian Sample == (neural net) ==> Data Sample

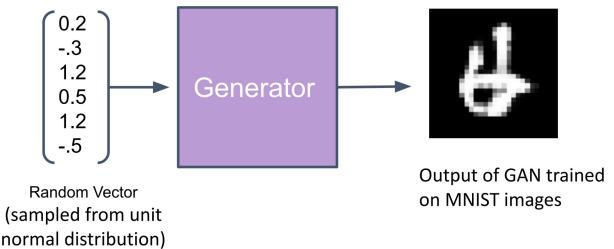
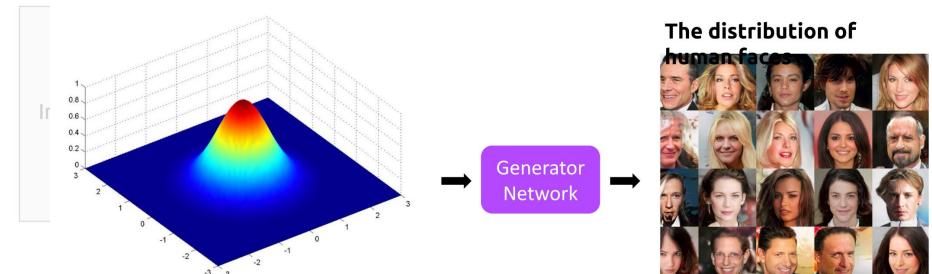


# Generative Modeling: Themes

**Idea:** Let's learn a complex function (aka a neural network) to transform a simple distribution sample into a complex one!

- Gaussian Sample == (neural net) ==> Data Sample

You have seen this before in:



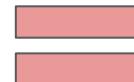
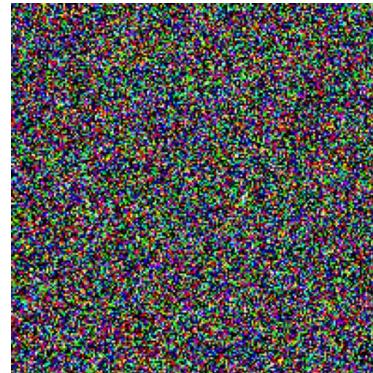
# Diffusion Models: a TLDR

**An observation:** adding steady amounts of Gaussian noise eventually corrupts an image into something indistinguishable from a random Gaussian sample.



# Diffusion Models: a TLDR

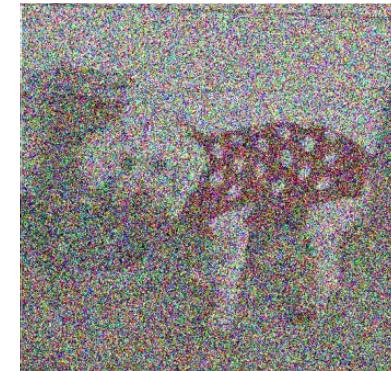
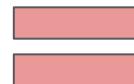
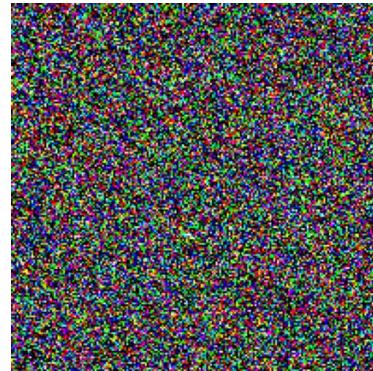
**An observation:** adding steady amounts of Gaussian noise eventually corrupts an image into something indistinguishable from a random Gaussian sample.



One Step

# Diffusion Models: a TLDR

**An observation:** adding steady amounts of Gaussian noise eventually corrupts an image into something indistinguishable from a random Gaussian sample.



Another Step

# Diffusion Models: a TLDR

**An observation:** adding steady amounts of Gaussian noise eventually corrupts an image into something indistinguishable from a random Gaussian sample.

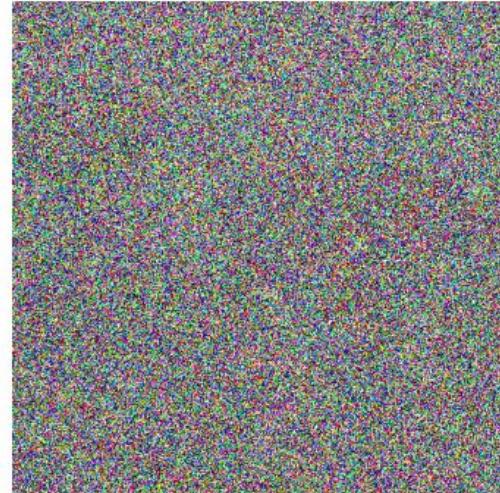


Many Steps

# Diffusion Models: a TLDR

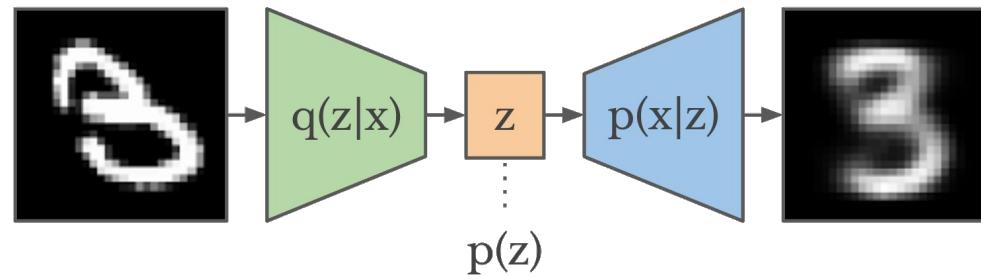
**An observation:** adding steady amounts of Gaussian noise eventually corrupts an image into something indistinguishable from a random Gaussian sample.

- Diffusion models simply learn to **reverse** this procedure over many timesteps



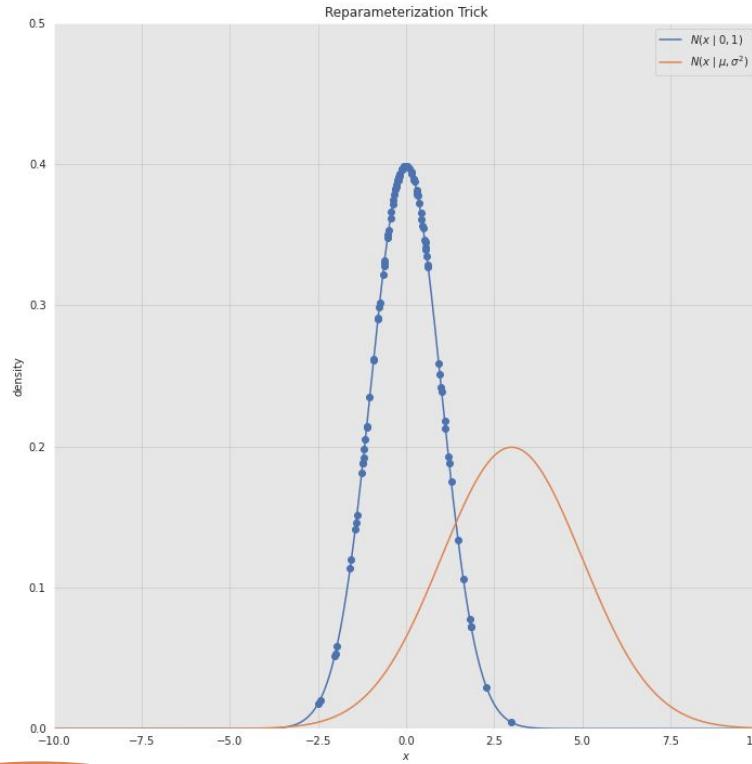
# Recap: Variational Autoencoders 🎩

Visually, we often see a VAE as:



How do we perform backpropagation through samples?

# Recap: Reparameterization Trick

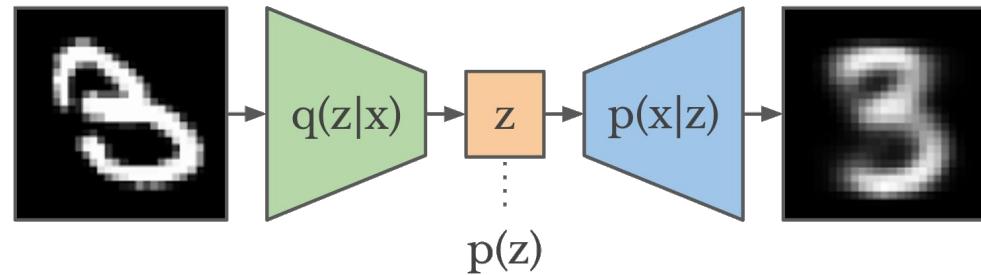


For  $x \sim \mathcal{N}(x | \mu, \sigma^2)$ ,

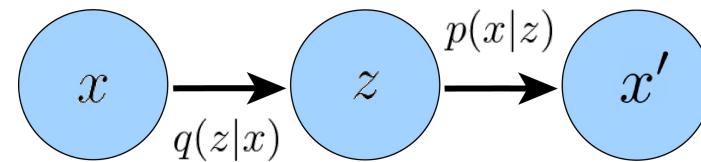
$\epsilon$ , where  $\epsilon \sim \mathcal{N}(x | 0, I)$

# Recap: Variational Autoencoders

Visually, we often see a VAE as:

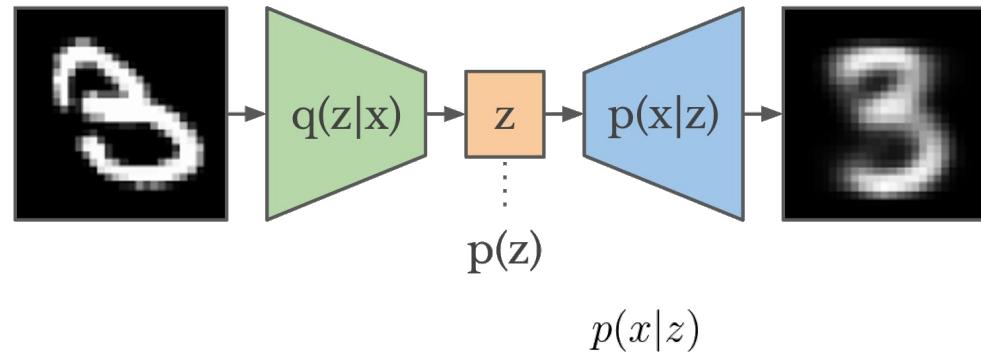


Or as:

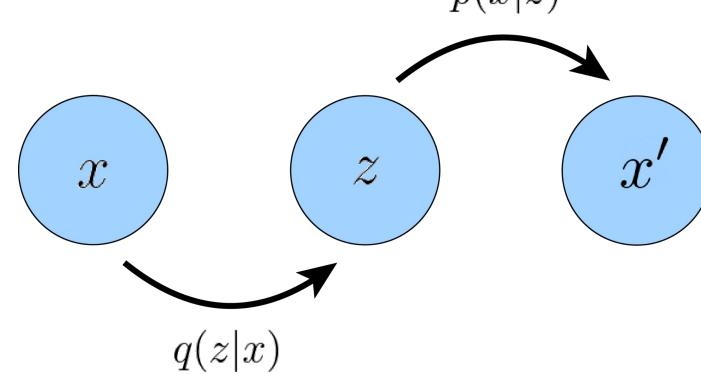


# Recap: Variational Autoencoders 🎩

Visually, we often see a VAE as:

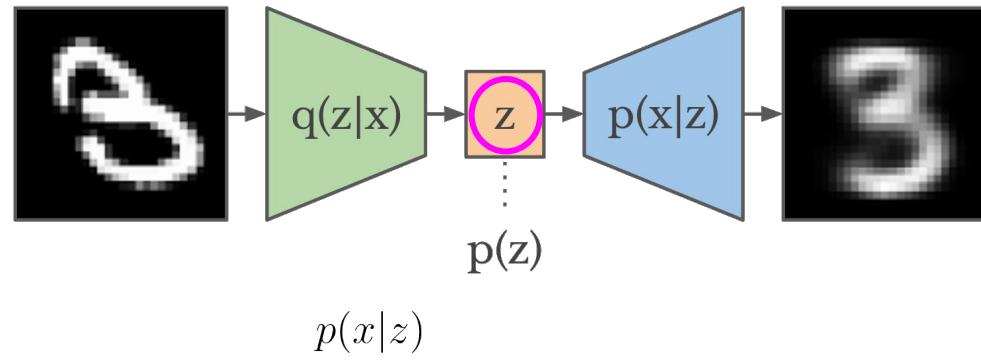


Or as:

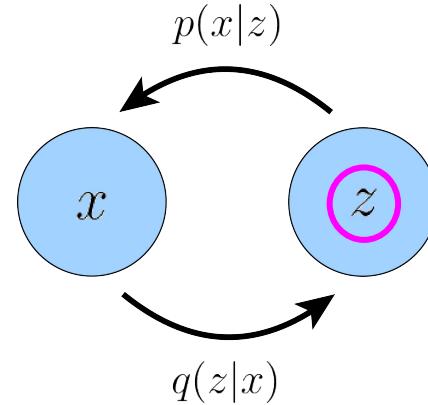


# Recap: Variational Autoencoders 🎓

Visually, we often see a VAE as:



Or as:

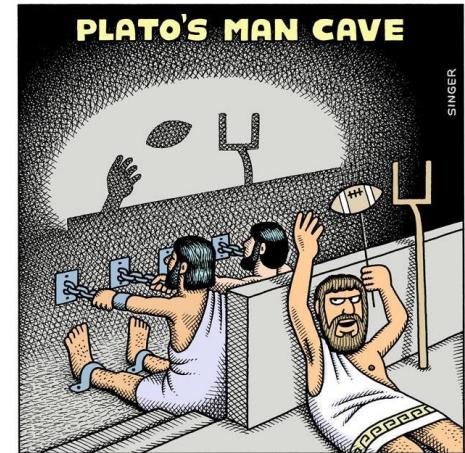


...but what's the intuition behind what is learned?

# Generative Modeling with Latent Variables

Given  $\mathbf{x} \sim p(\mathbf{x})$  we might want to learn  $p_\theta(\mathbf{x}) \approx p(\mathbf{x})$  (*modeling*)

What if we assume latent variables  $\mathbf{z}$  exist?



Elon has an idea...



Elon has an idea...



Elon has an idea...



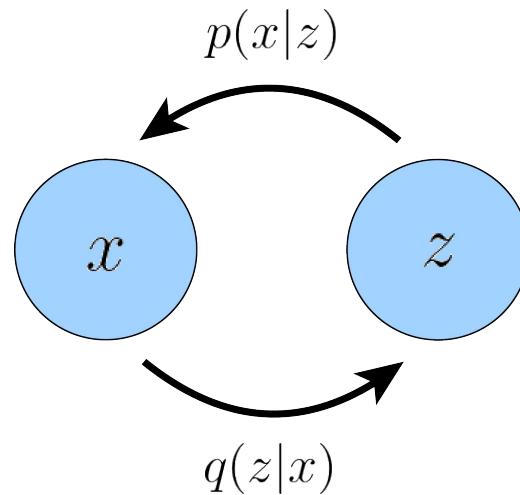
# Elon has an idea...



# Hierarchical VAEs

Generalize VAEs by enabling a hierarchy of latents  $z = z_1, \dots, z_T$

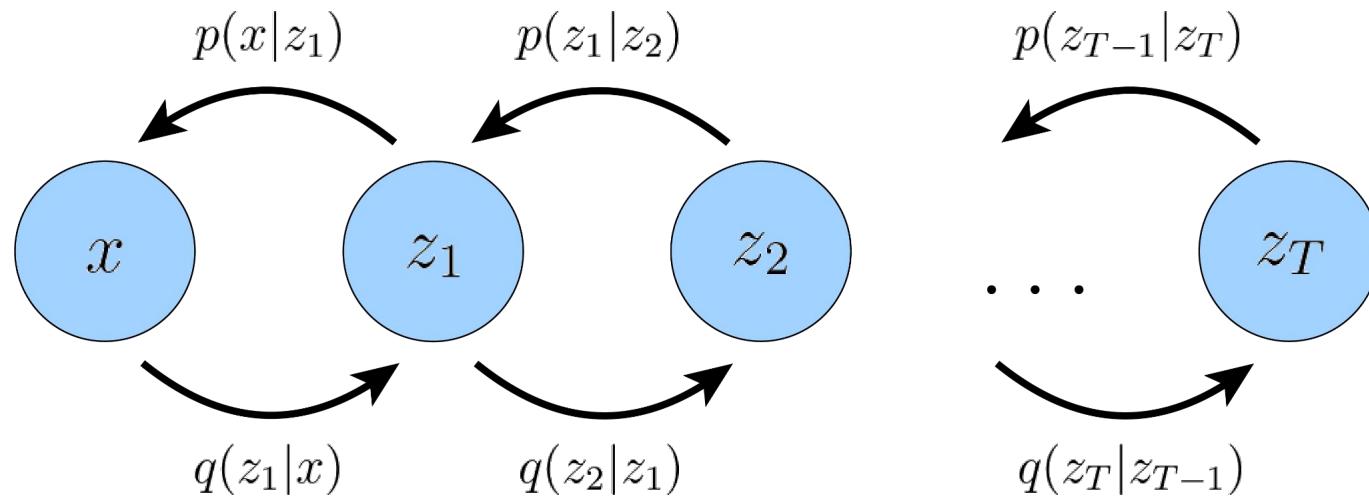
This is essentially learning a bunch of stacked VAEs



# Hierarchical VAEs

Generalize VAEs by enabling a hierarchy of latents  $z = z_1, \dots, z_T$

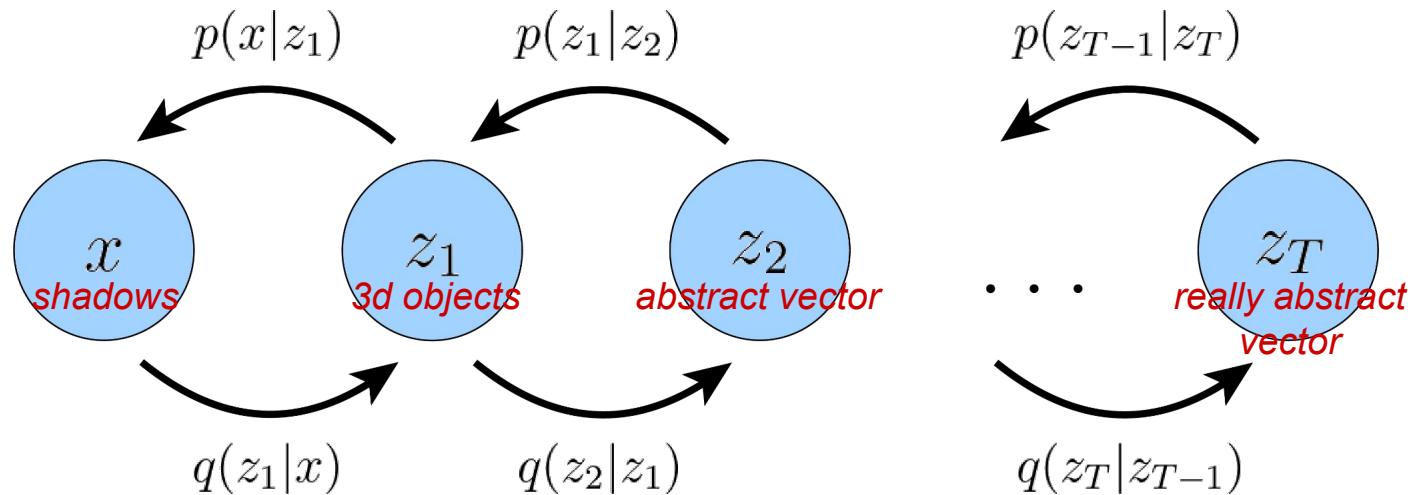
This is essentially learning a bunch of stacked VAEs



*Disclaimer: Elon did not actually come up with this idea.*

# Hierarchical VAEs

Let's think like a caveman...

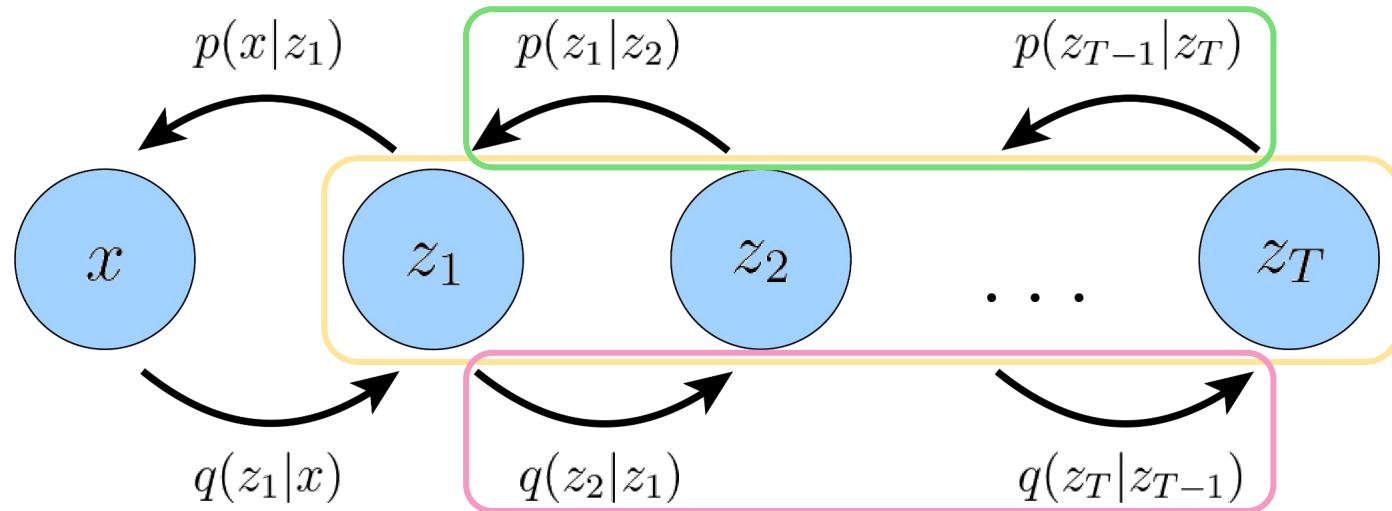


# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume all latent dimensions are the same?*

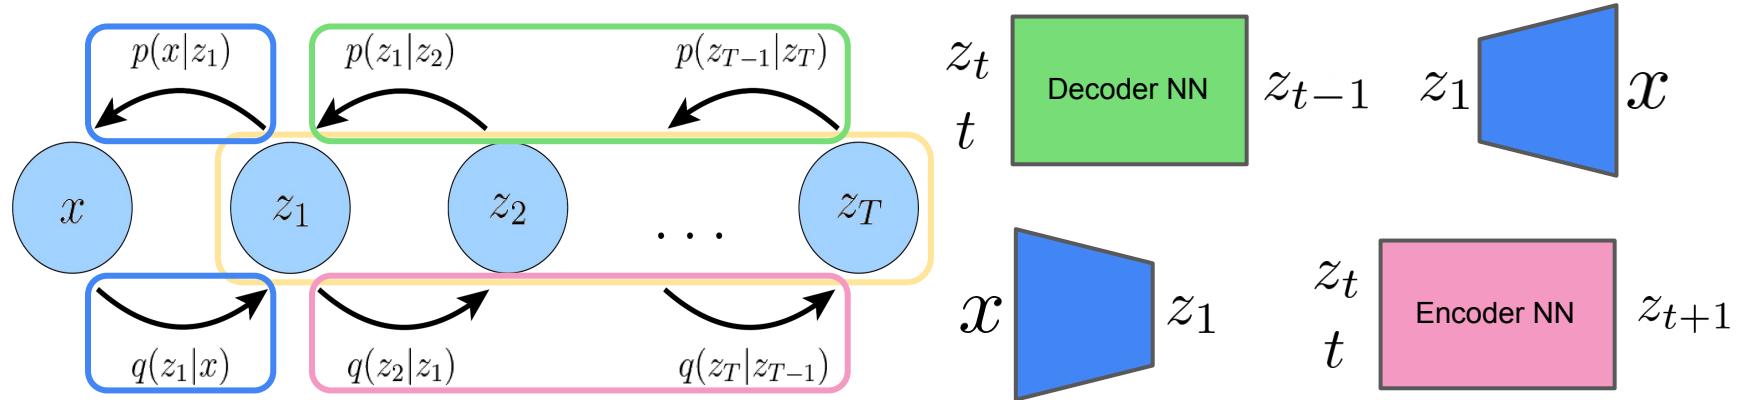


# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume all latent dimensions are the same?*

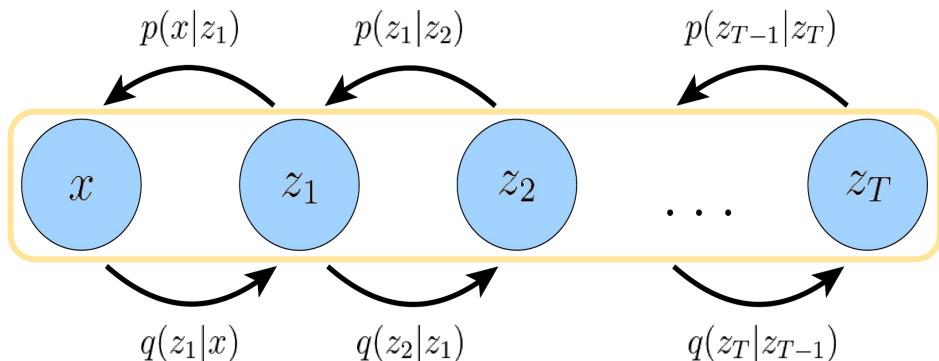


# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume **all** dimensions are the same?*

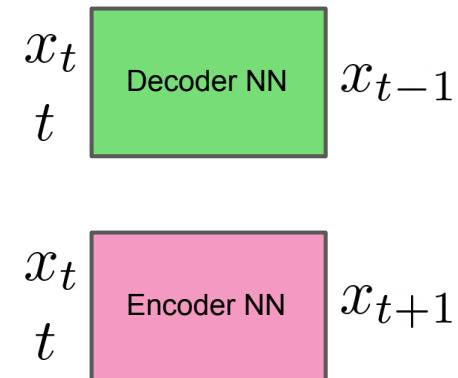
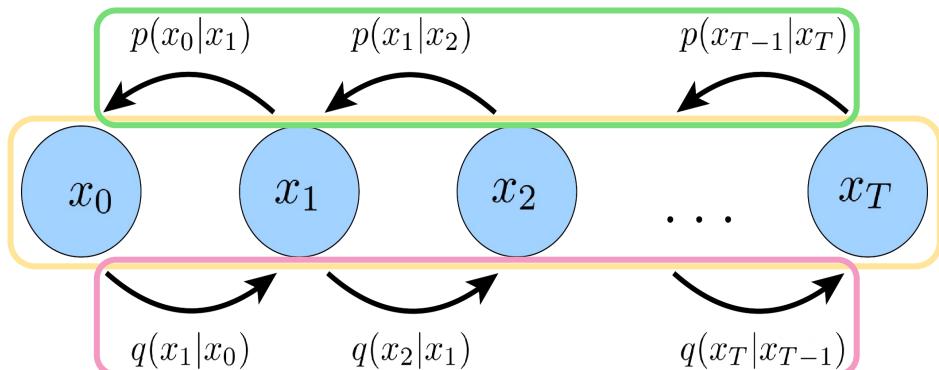


# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume **all** dimensions are the same?*

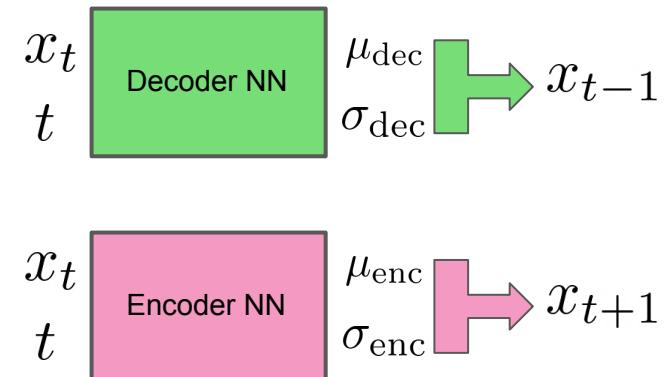
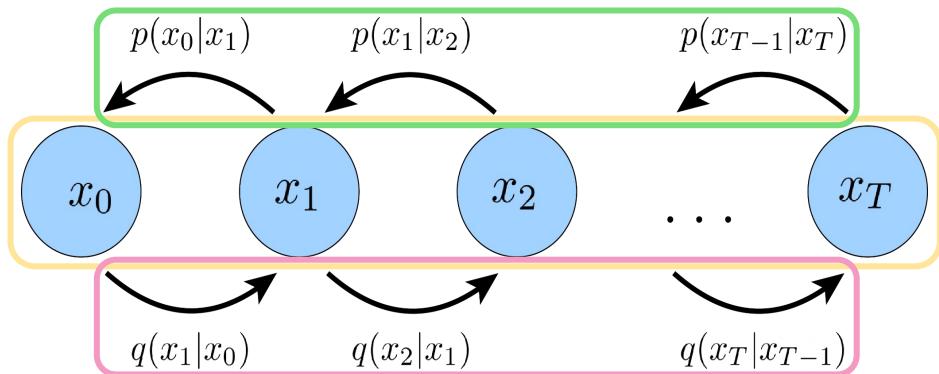


# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume **all** dimensions are the same?*



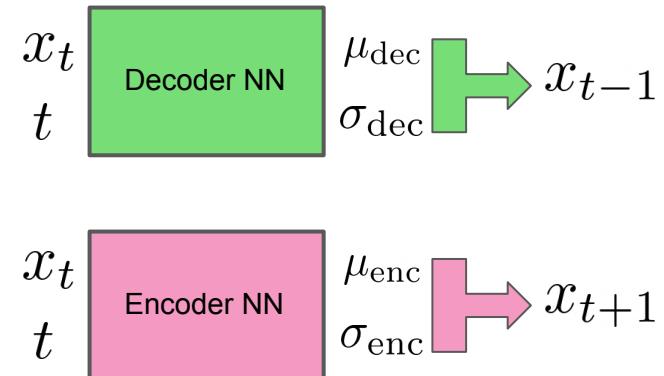
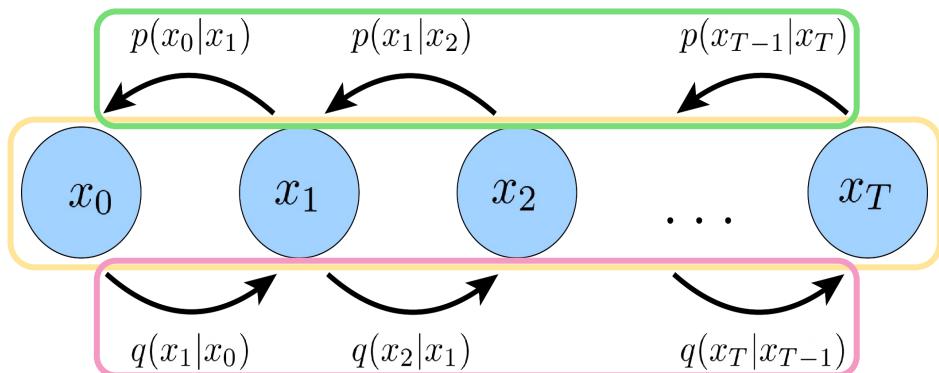
# Hierarchical VAEs

Question:

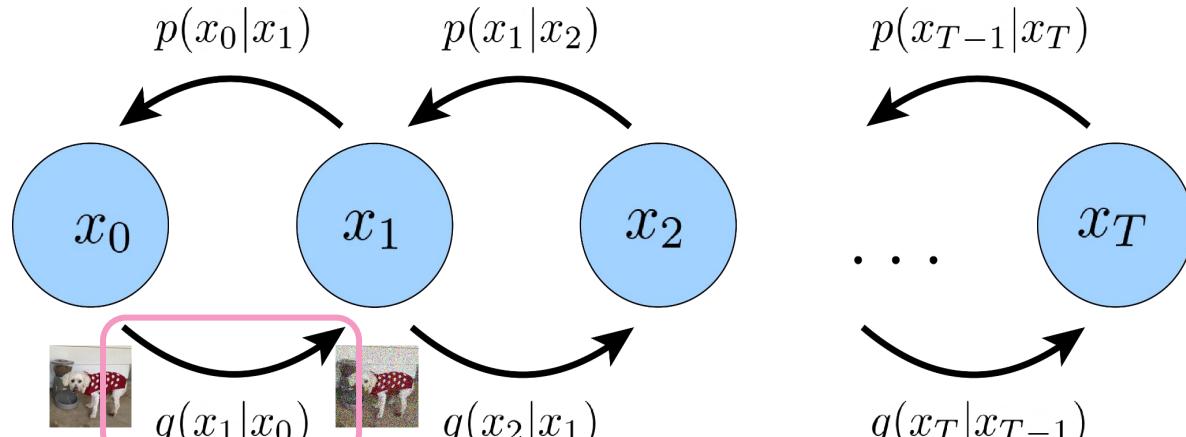
- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume **all** dimensions are the same?*

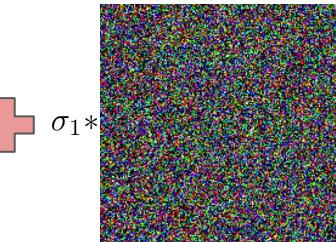
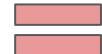
*...what if we assume all encoder transitions are known Gaussians centered around their previous input?*



# Let's take a look at one encoding

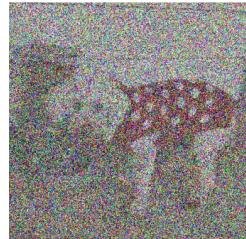
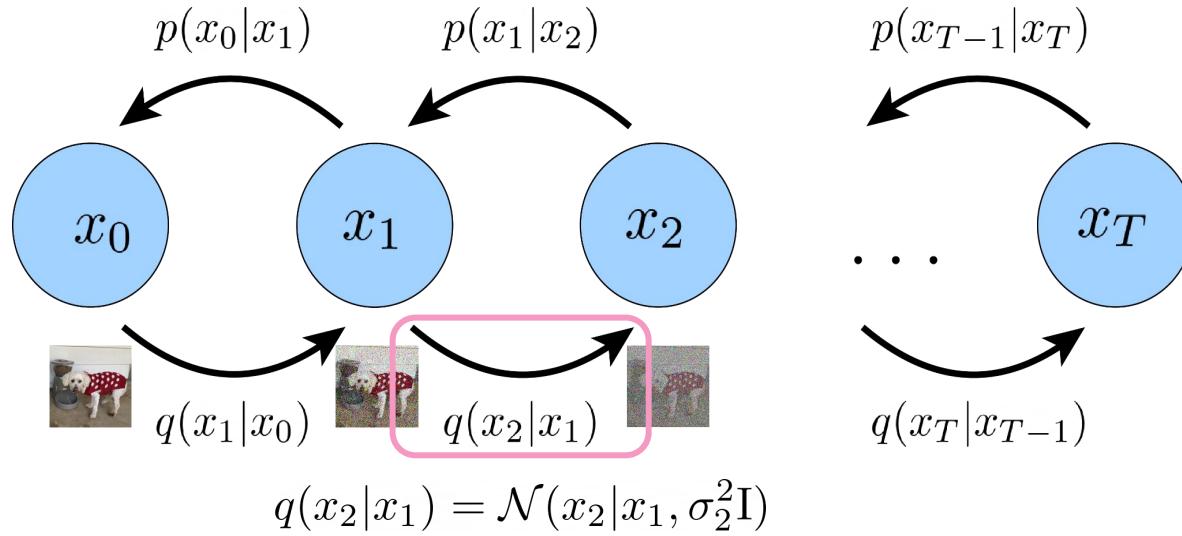


$$q(x_1|x_0) = \mathcal{N}(x_1|x_0, \sigma_1^2 \mathbf{I})$$



reparam. trick!

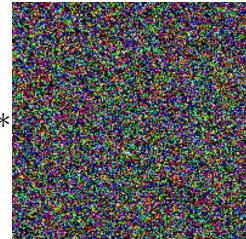
# Let's take a look at one encoding



$$=$$

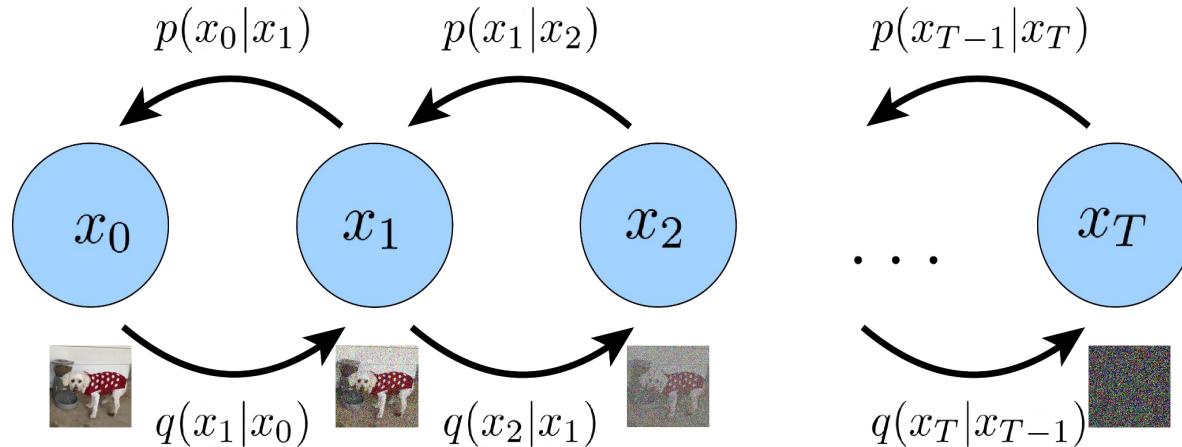


$$+$$

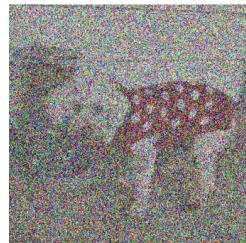
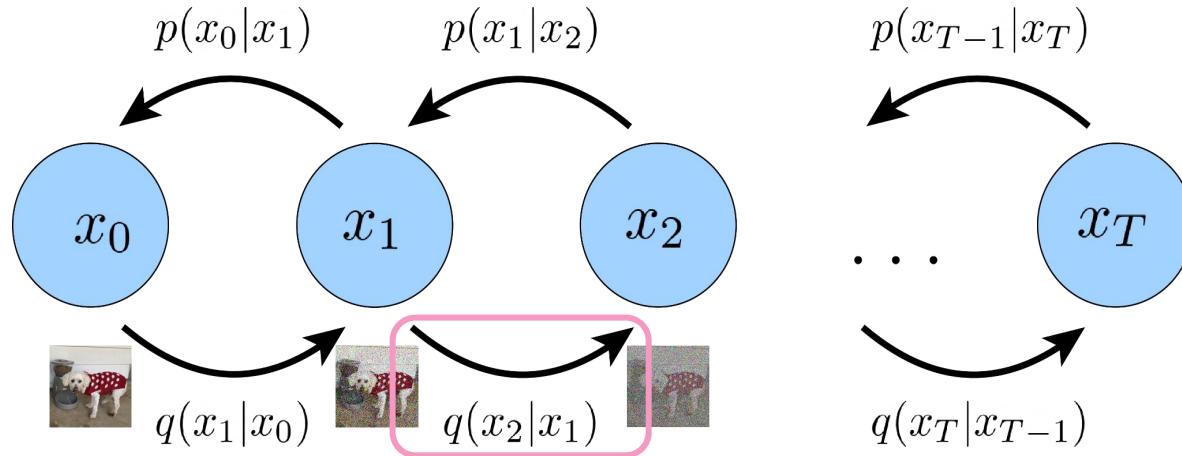


reparam. trick!

# Let's take a look at one encoding



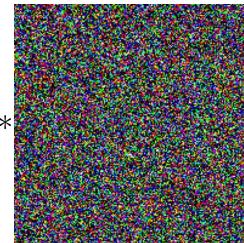
# Let's take a look at one encoding



$$=$$



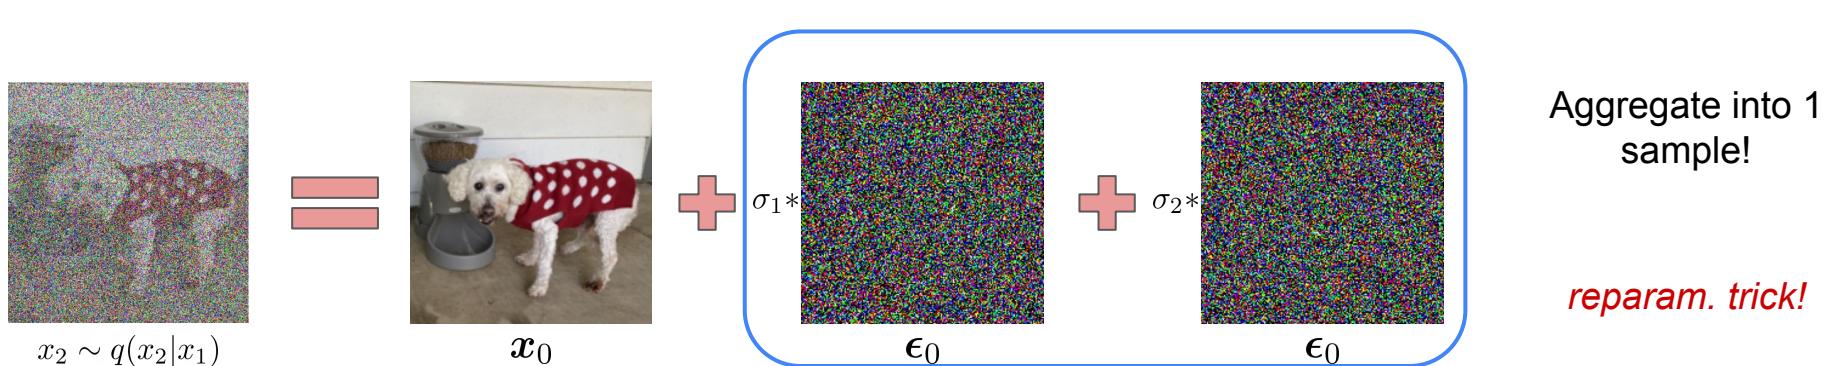
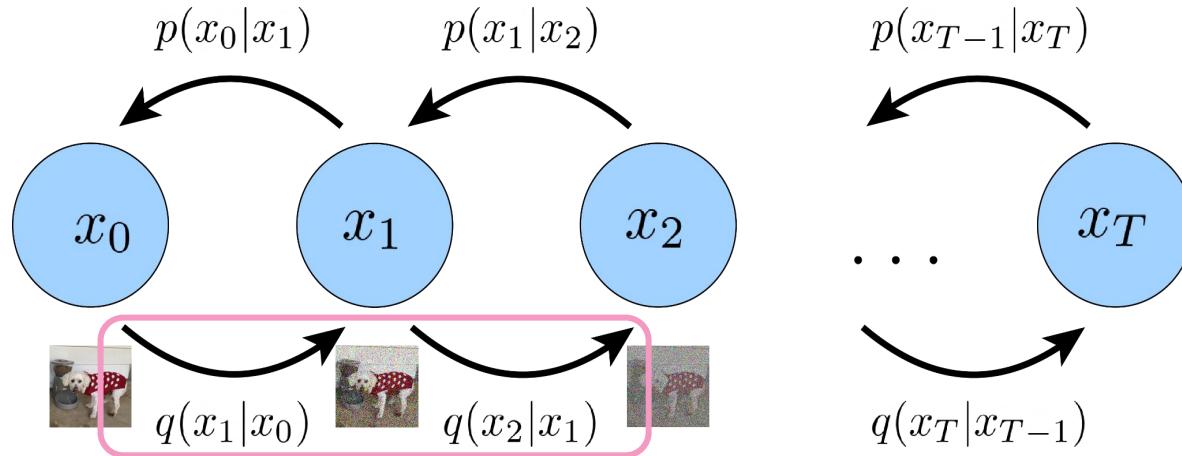
$$+$$



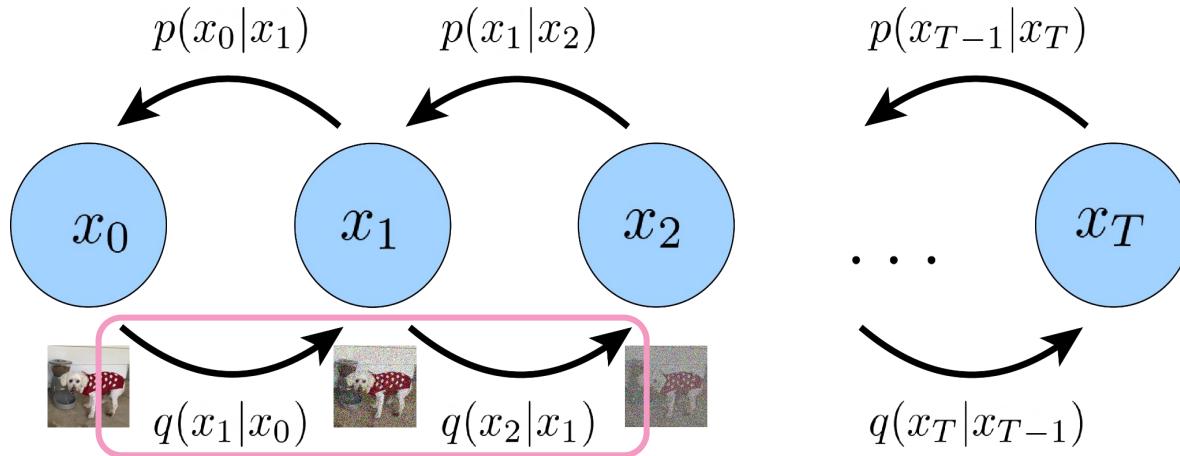
$$\sigma_2^*$$

reparam. trick!

# Let's take a look at one encoding



# Let's take a look at one encoding



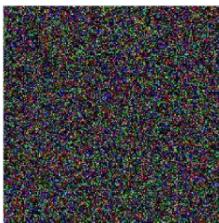
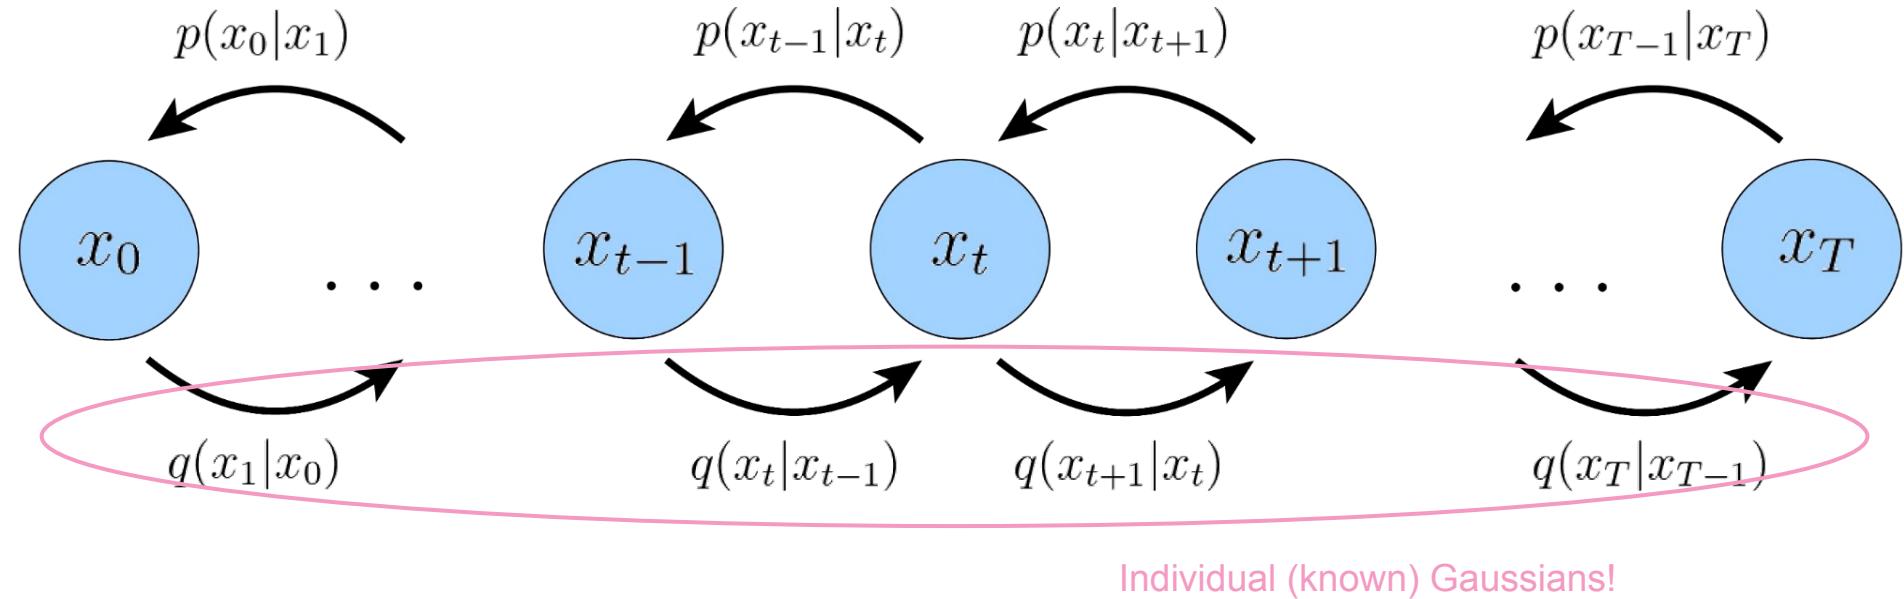
where,

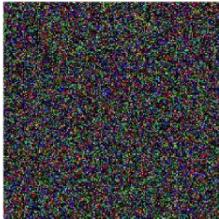
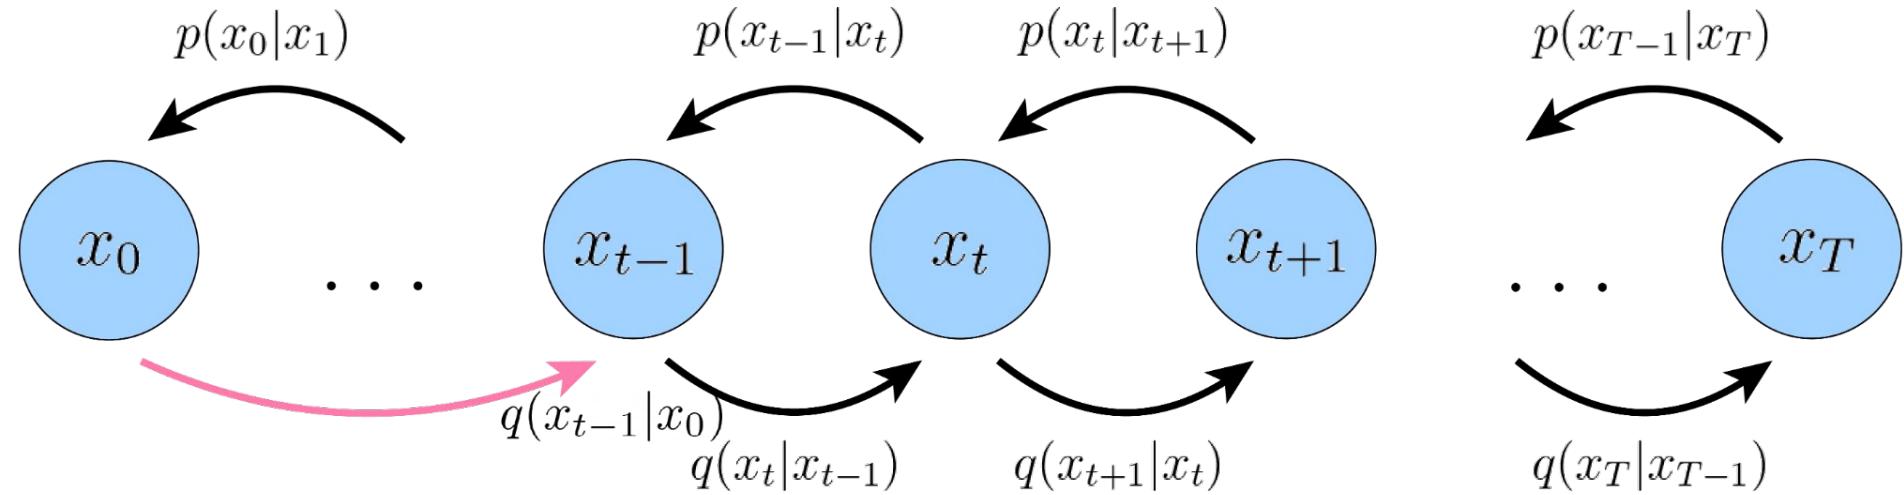
$$\alpha_2 = \sqrt{\sigma_1^2 + \sigma_2^2}$$

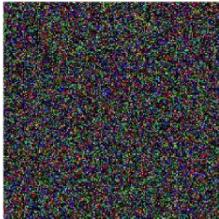
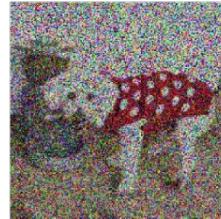
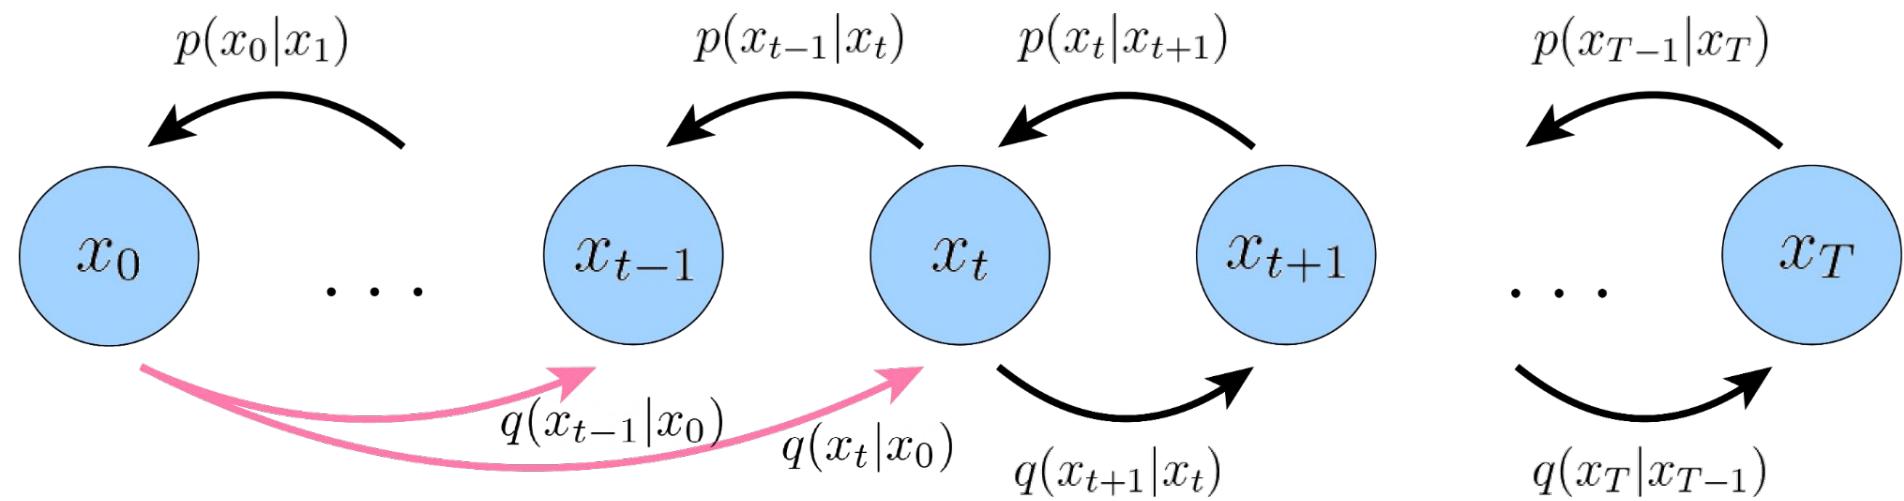
and,

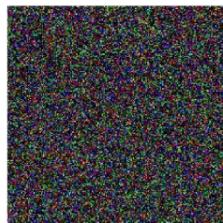
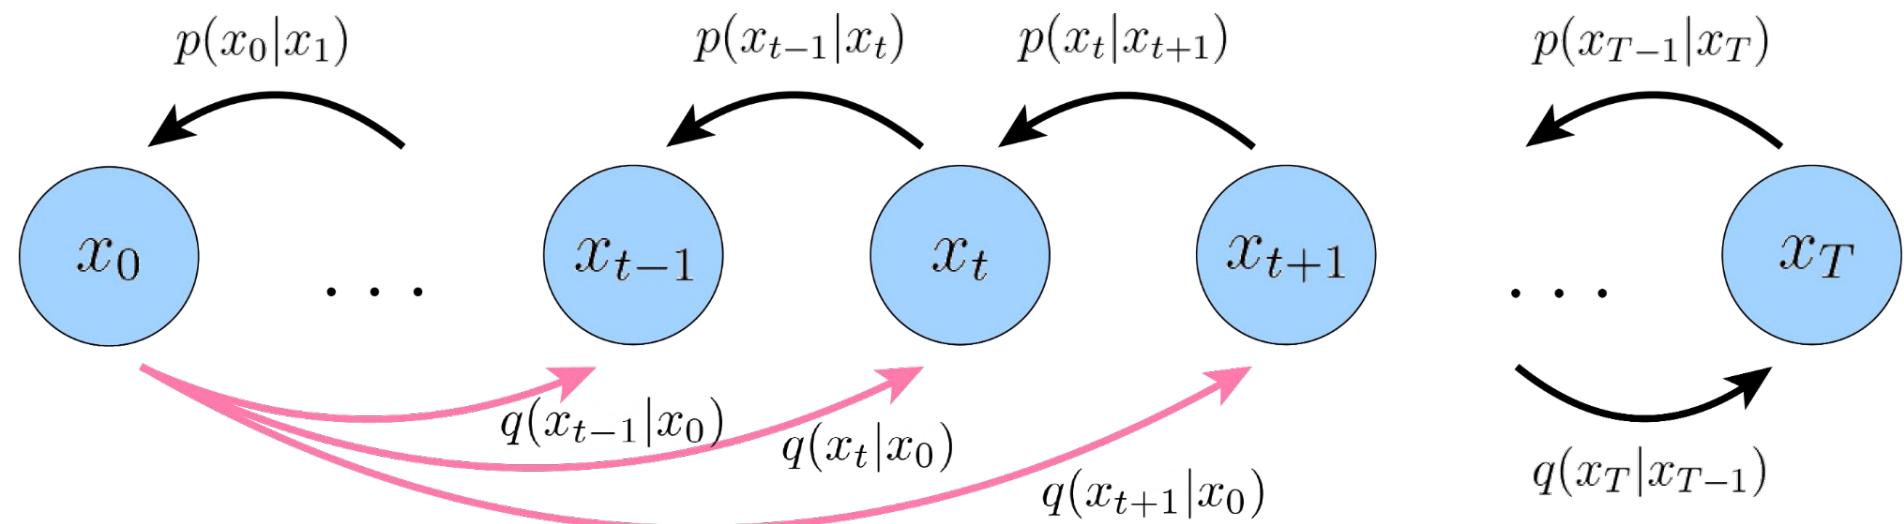
$$q(x_2|x_0) = \mathcal{N}(x_2|x_0, \alpha_2^2)$$

Aggregate into 1 sample!  
reparam. trick!



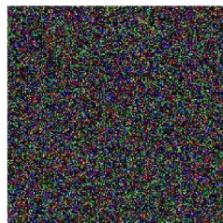
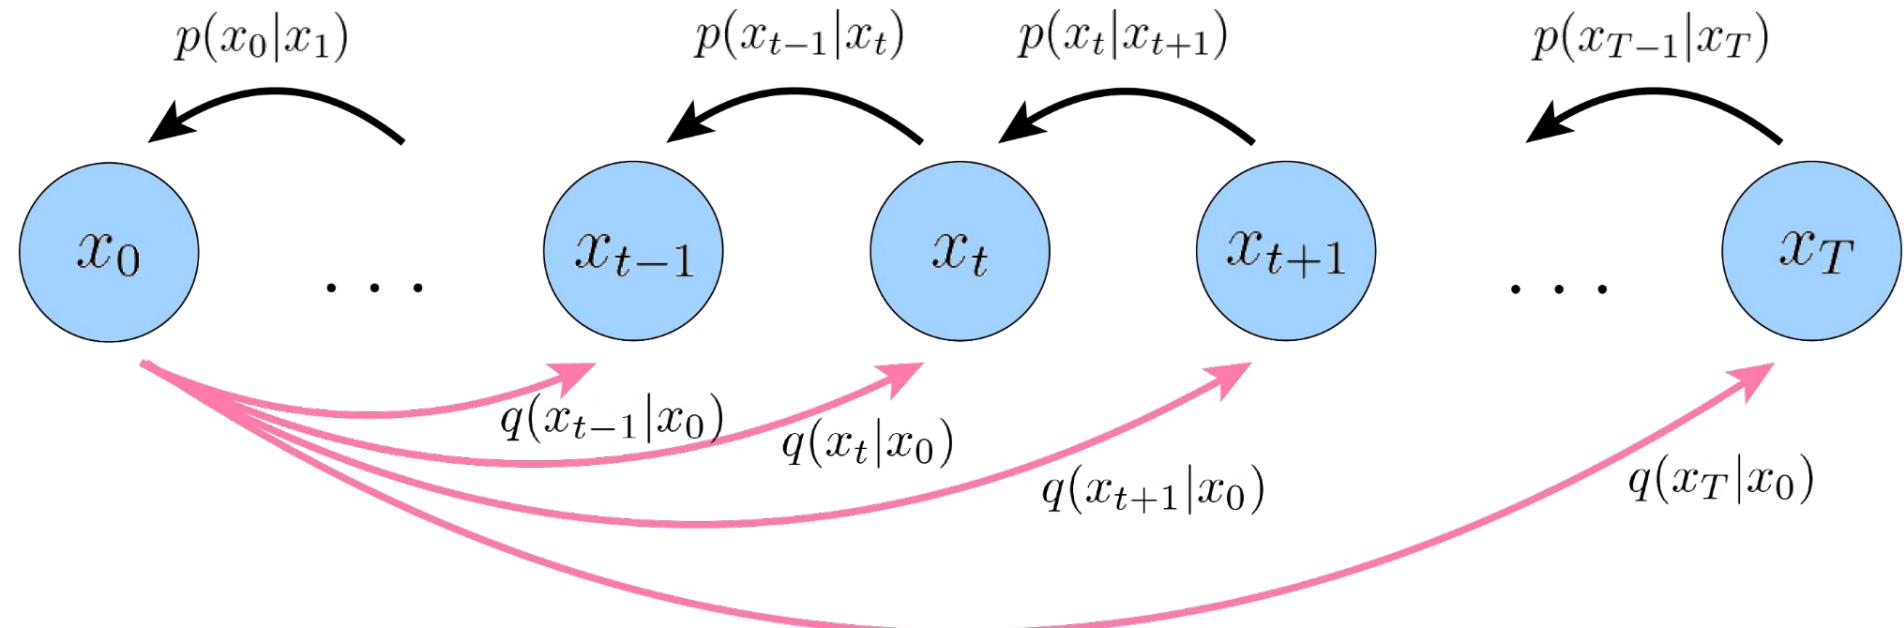






$q(x_t|x_0)$  is a Gaussian, for arbitrary  $t$ !

$q(x_t|x_0) = \mathcal{N}(x_t|x_0, \alpha_t^2 \mathbf{I})$ , where  $\alpha_0, \alpha_1, \dots, \alpha_T$  are all known/fixed.



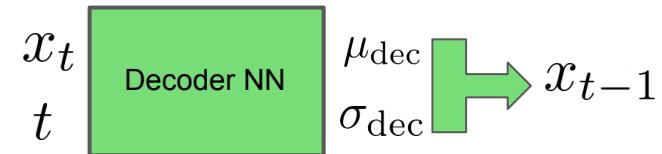
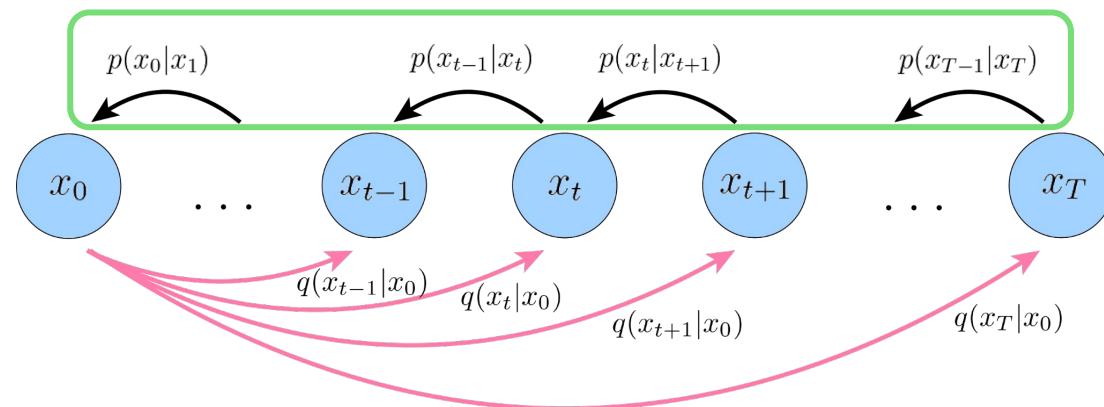
# Hierarchical VAEs

Question:

- In a VAE we learn two networks: an encoder and a decoder.
- How many do we need to learn for a Hierarchical VAE?

*...what if we assume **all** dimensions are the same?*

*...what if we assume all encoder transitions are known Gaussians centered around their previous input?*



*...then we can aggregate and simplify the distribution of each intermediate “latent”!*

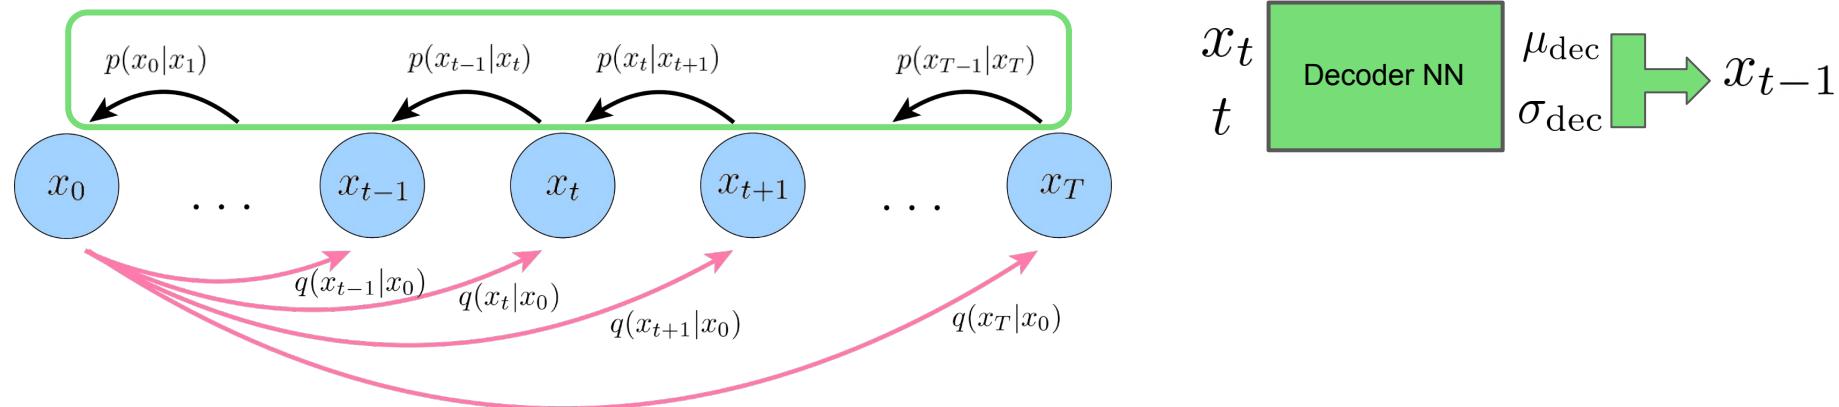
# Diffusion Models

It turns out, that this is exactly what a diffusion model is!

- A Hierarchical VAE with these assumptions:

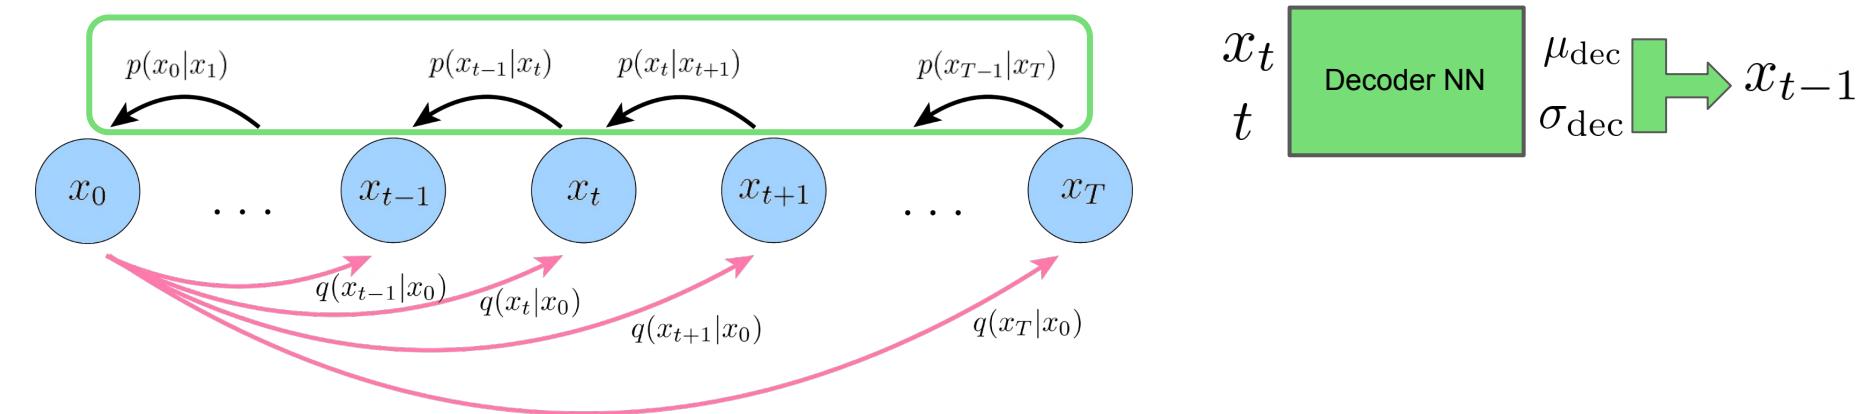
*...what if we assume **all** dimensions are the same?*

*...what if we assume all encoder transitions are known Gaussians centered around their previous input?*



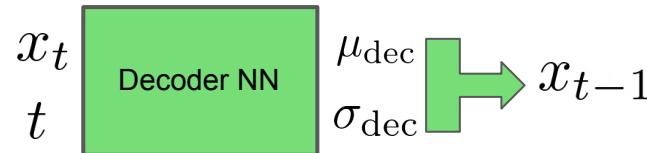
# Diffusion Models

A diffusion model is implemented as a single neural network (the decoder)



# Optimization?

We want to learn a denoising decoder:

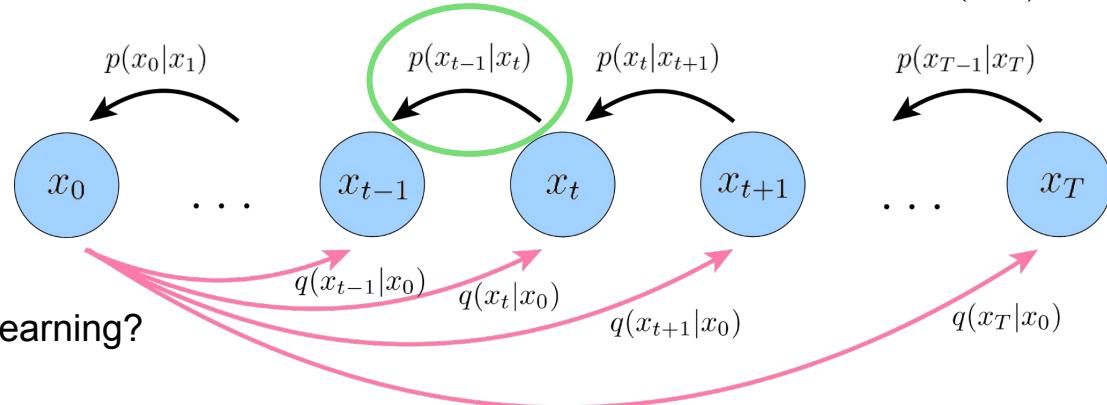


$$\hat{x}_{t-1} = \mu_{\text{dec}} + \sigma_{\text{dec}} * \epsilon$$

*reparam. trick!*

$$\epsilon \sim \mathcal{N}(0, I)$$

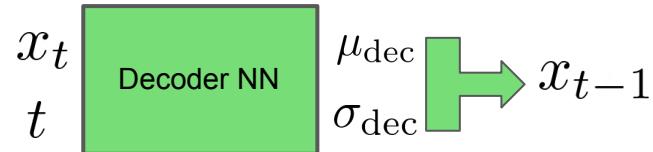
But what is the form of  $x_{t-1}$ ?



...can we formulate this as supervised learning?

# Optimization?

We want to learn a denoising decoder:



$$\hat{x}_{t-1} = \mu_{\text{dec}} + \sigma_{\text{dec}} * \epsilon$$

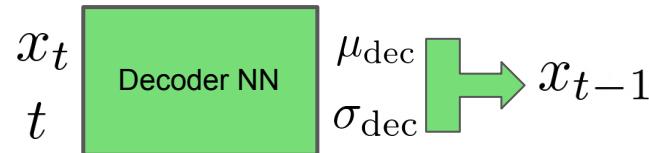
*reparam. trick!*

$$\epsilon \sim \mathcal{N}(0, I)$$

But what is the form of  $x_{t-1}$ ?

# Optimization?

We want to learn a denoising decoder:



$$\hat{x}_{t-1} = \mu_{\text{dec}} + \sigma_{\text{dec}} * \epsilon$$

*reparam. trick!*

$$\epsilon \sim \mathcal{N}(0, I)$$

But what is the form of  $x_{t-1}$ ?

Recall that:

$$q(x_{t-1}|x_0) = \mathcal{N}(x_{t-1}|x_0, \alpha_{t-1}^2 I)$$
$$\therefore x_{t-1} = x_0 + \alpha_{t-1} * \epsilon$$

*reparam. trick!*

$$\epsilon \sim \mathcal{N}(0, I)$$

*Do we really need to predict  $\sigma_{\text{dec}}$ ?*

*What is the ground truth signal for  $\mu_{\text{dec}}$ ?*

# Optimization?

We want to learn a denoising decoder:



$$\hat{x}_{t-1} = \hat{x}_0 + \alpha_{t-1} * \epsilon$$

*reparam. trick!*

$$\epsilon \sim \mathcal{N}(0, I)$$

But what is the form of  $x_{t-1}$ ?

Recall that:

$$q(x_{t-1}|x_0) = \mathcal{N}(x_{t-1}|x_0, \alpha_{t-1}^2 I)$$

$$\therefore x_{t-1} = x_0 + \alpha_{t-1} * \epsilon$$

*reparam. trick!*

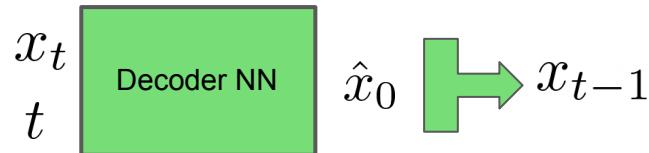
$$\epsilon \sim \mathcal{N}(0, I)$$

*Do we really need to predict  $\sigma_{dec}$ ?*

*What is the ground truth signal for  $\mu_{dec}$ ?*

# Optimization?

We want to learn a denoising decoder:



$$\hat{x}_{t-1} = \hat{x}_0 + \alpha_{t-1} * \epsilon$$

reparam. trick!

$$\epsilon \sim \mathcal{N}(0, I)$$

But what is the form of  $x_{t-1}$ ?

Recall that:

$$q(x_{t-1}|x_0) = \mathcal{N}(x_{t-1}|x_0, \alpha_{t-1}^2 I)$$

$$\therefore x_{t-1} = x_0 + \alpha_{t-1} * \epsilon$$

reparam. trick!

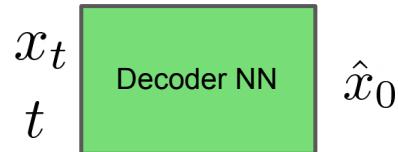
$$\epsilon \sim \mathcal{N}(0, I)$$

*Do we really need to predict  $\sigma_{dec}$ ?*

*What is the ground truth signal for  $\mu_{dec}$ ?*

# Optimization?

We want to learn a denoising decoder:

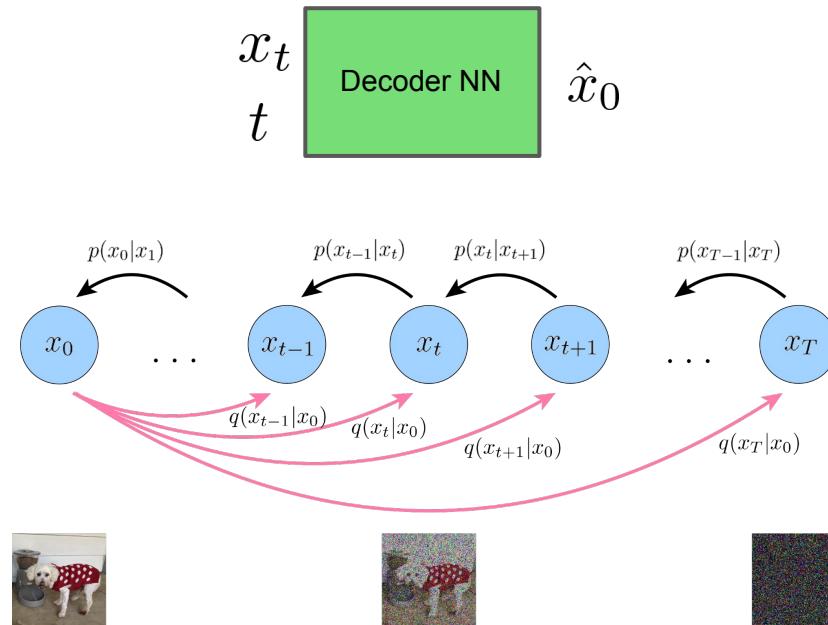


So in the end, a diffusion model is simply one Neural Network that predicts a clean image  $x_0$  from arbitrary noisified image  $x_t$ .

# Diffusion Models: A Summary

A Diffusion Model is:

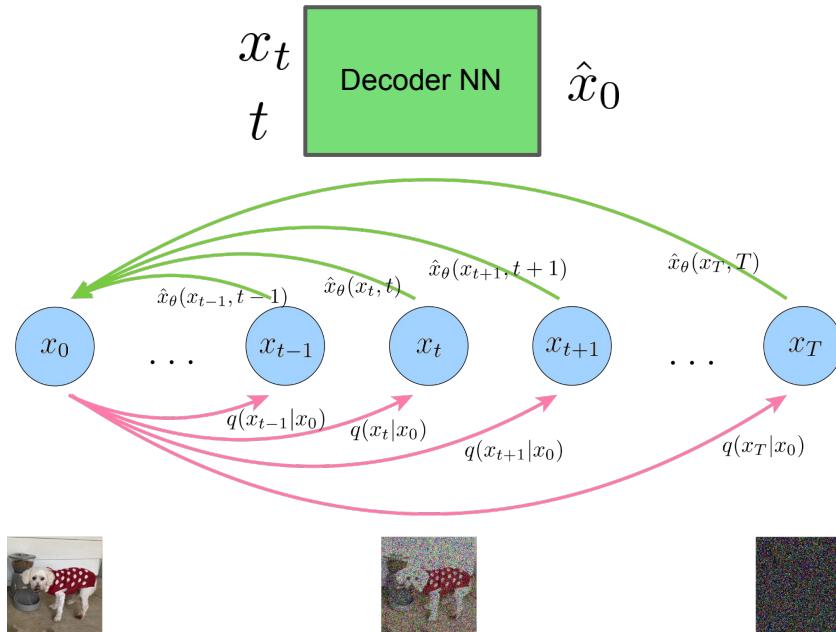
- One NN that predicts a clean image from a noisy version of the image



# Diffusion Models: A Summary

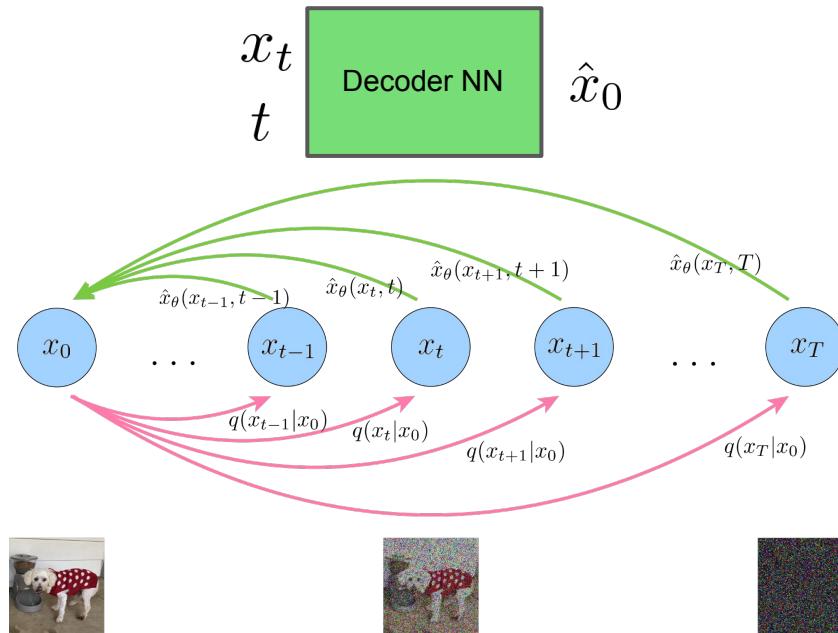
A Diffusion Model is:

- One NN that predicts a clean image from a noisy version of the image

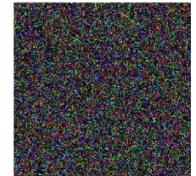
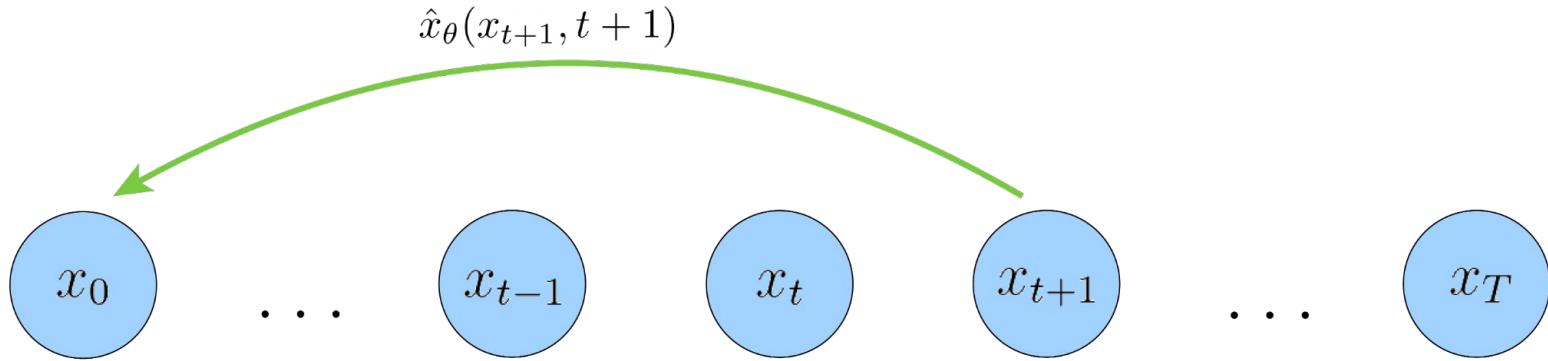


# Diffusion Models: A Summary

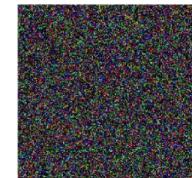
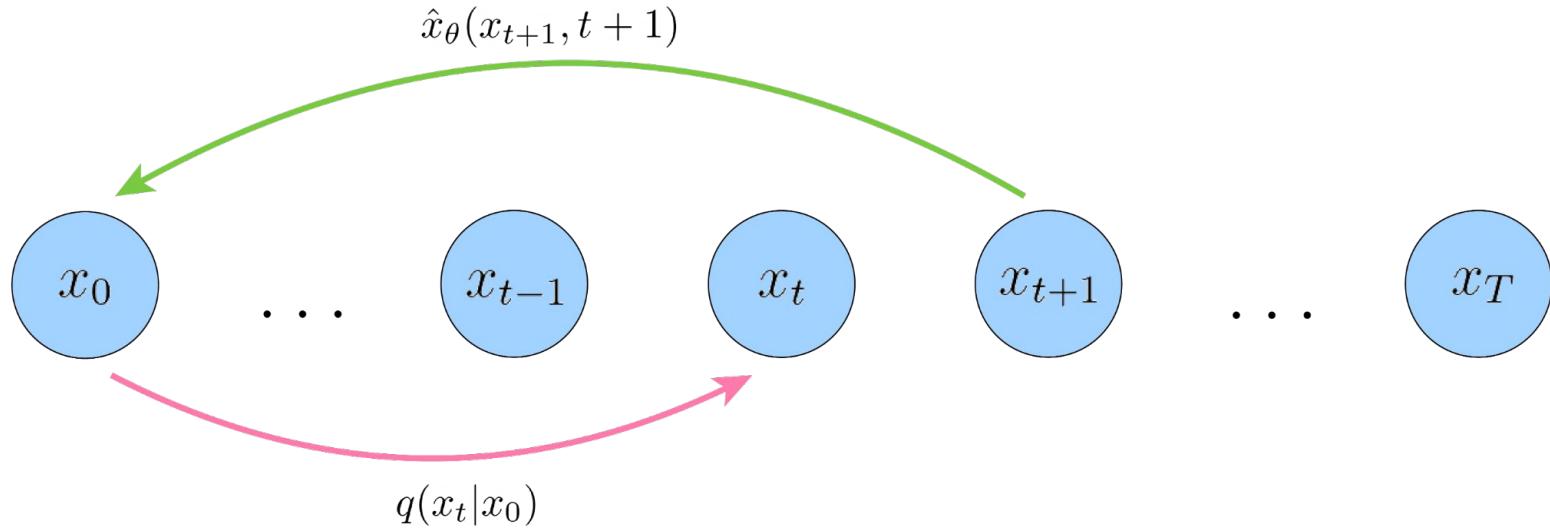
How do we perform sampling?



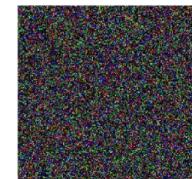
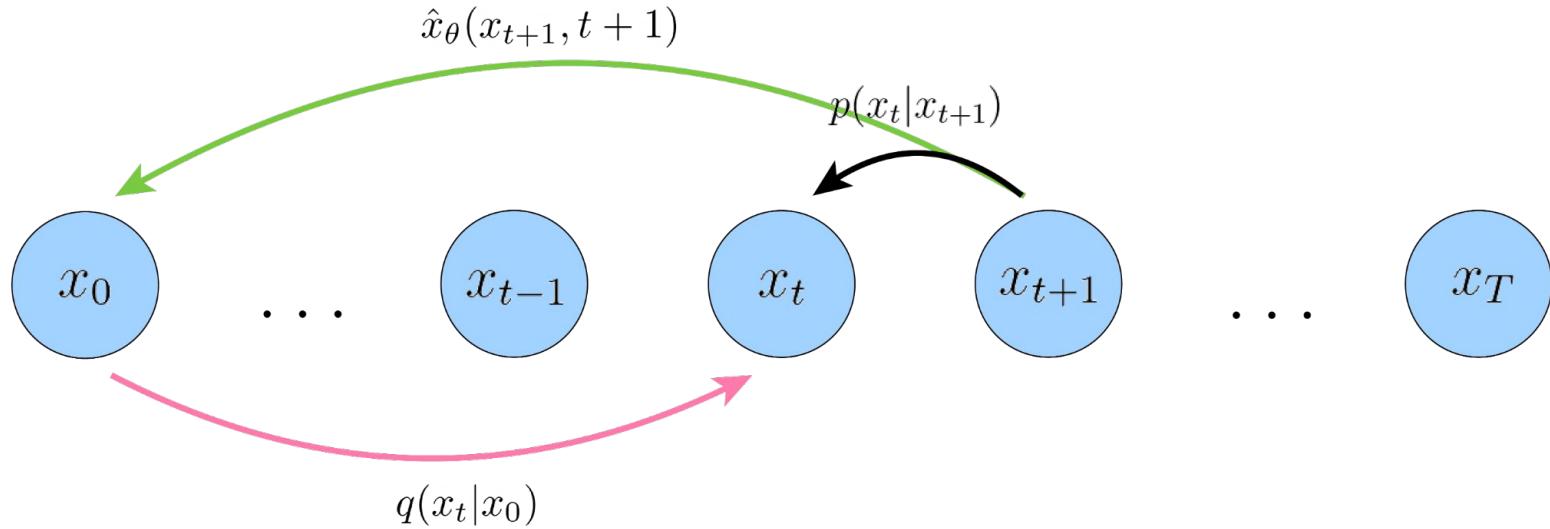
# Sampling



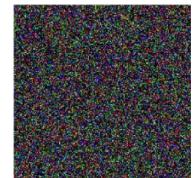
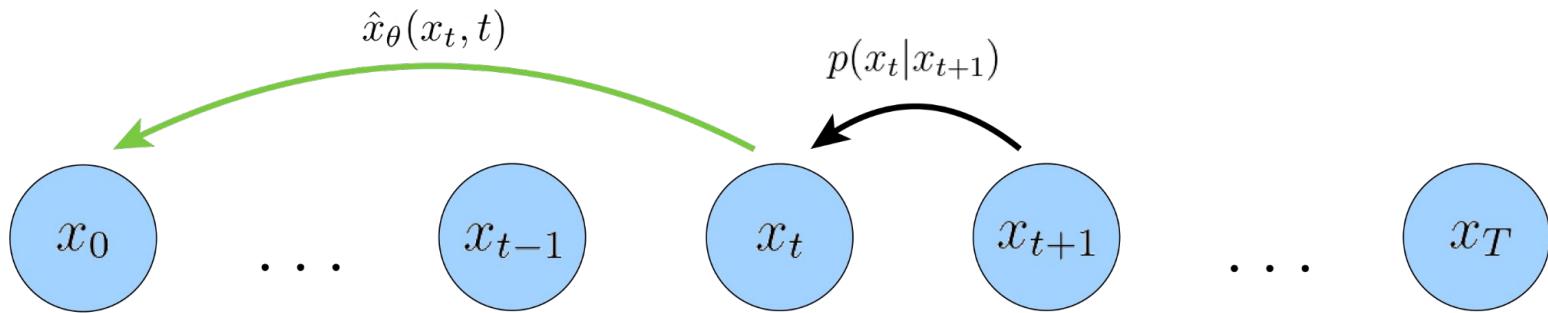
# Sampling



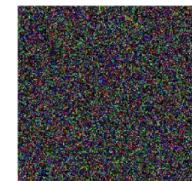
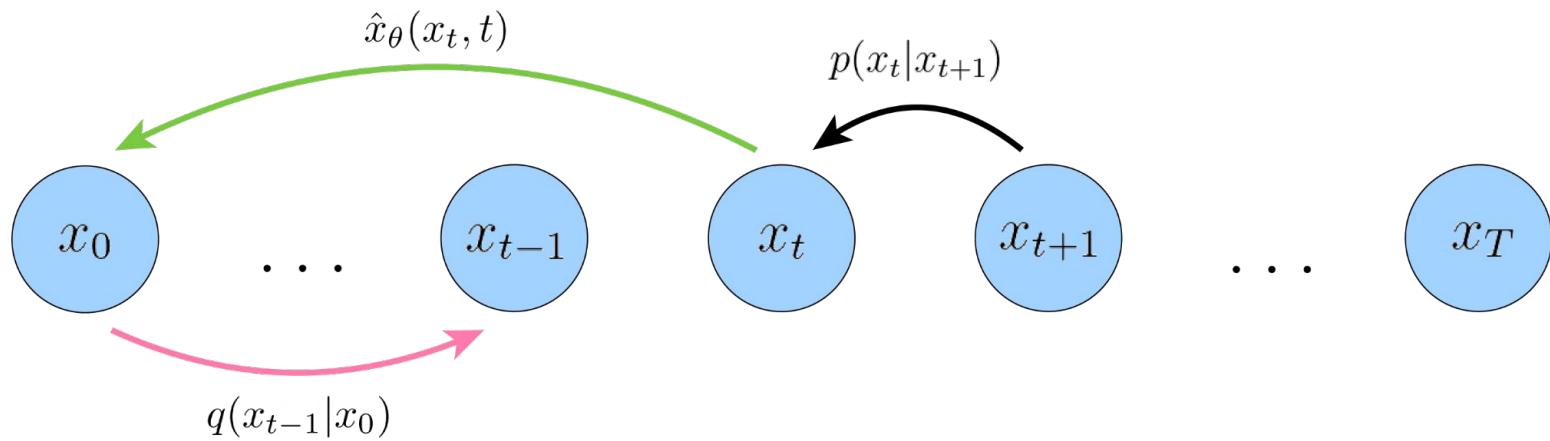
# Sampling



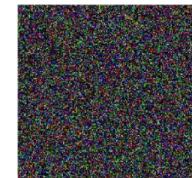
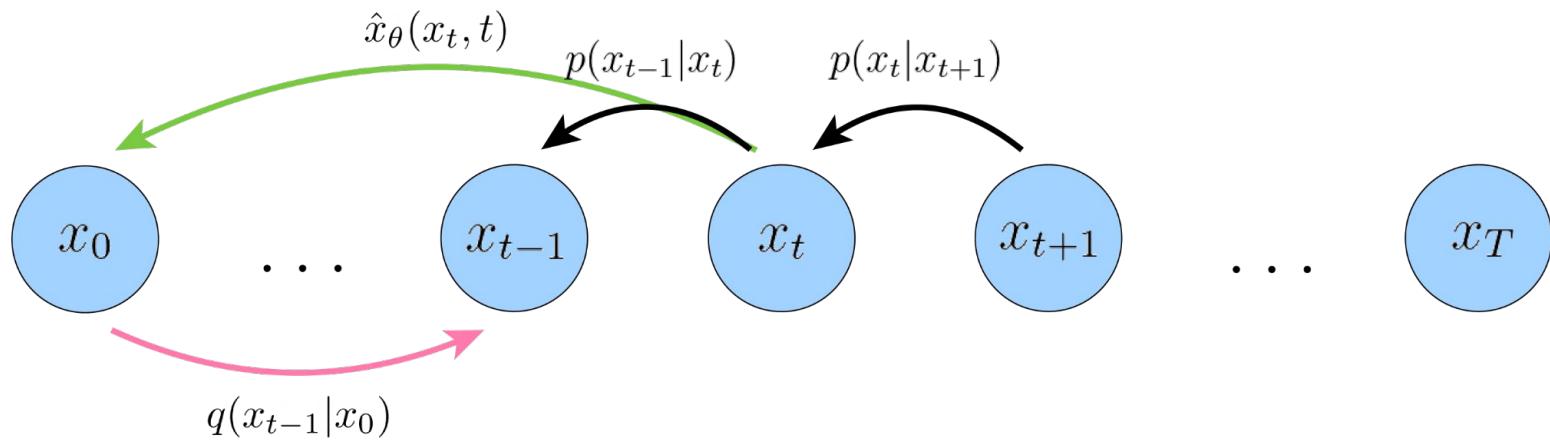
# Sampling



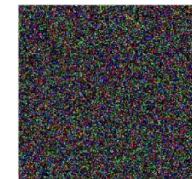
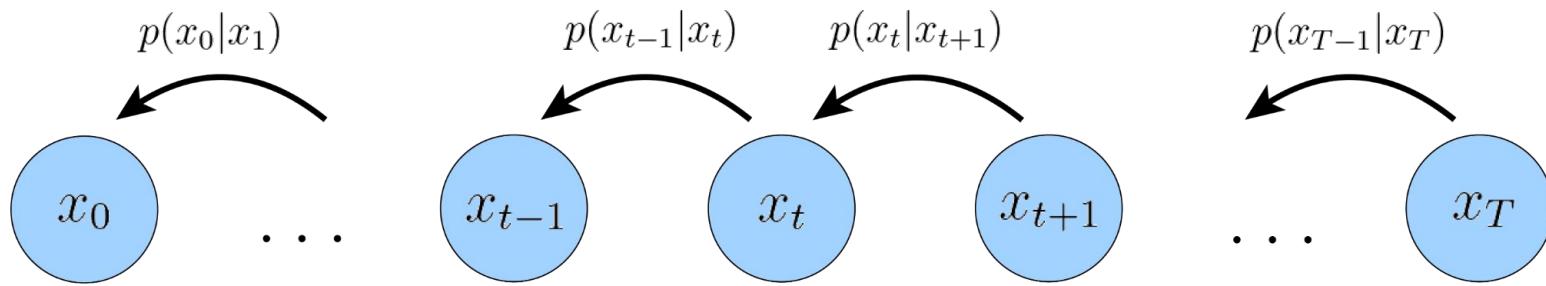
# Sampling



# Sampling



# Sampling



# Pseudocode

---

## Algorithm 1 Training

---

```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(1, \dots, T)$ 
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
        $\nabla_{\theta} \|\mathbf{x}_0 - \hat{\mathbf{x}}_{\theta}(\mathbf{x}_0 + \alpha_t \boldsymbol{\epsilon}, t)\|^2$ 
7: until converged
```

---

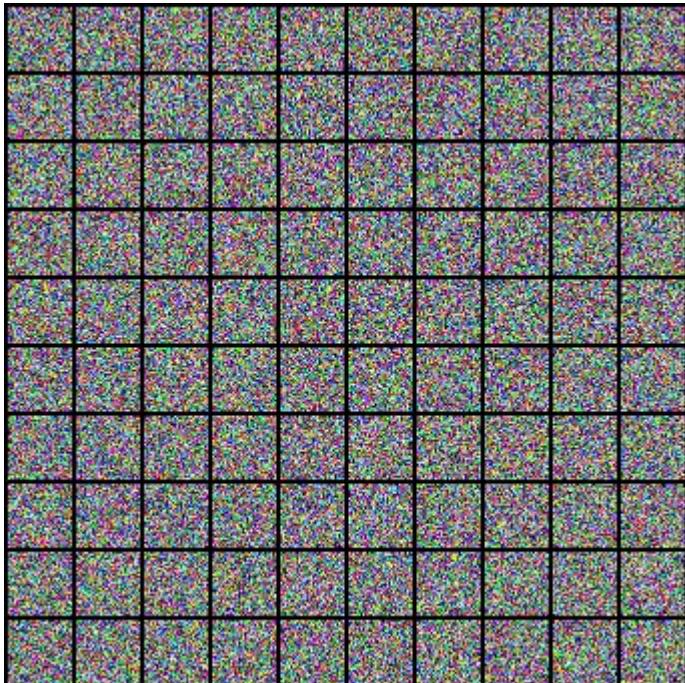
## Algorithm 2 Sampling

---

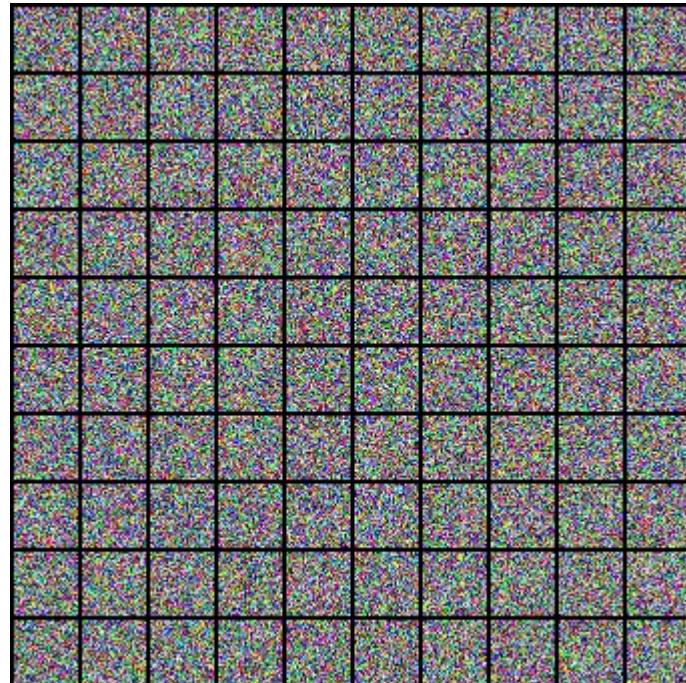
```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$ :
3:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\boldsymbol{\epsilon} = 0$ 
4:    $\mathbf{x}_{t-1} = \hat{\mathbf{x}}_{\theta}(\mathbf{x}_t, t) + \alpha_{t-1} \boldsymbol{\epsilon}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

---

# Examples!



Celeb-A



CIFAR-10

# Examples!



1024x1024 samples

source: [Generative Modeling by Estimating Gradients of the Data Distribution](#)

# Three Different Interpretations

It turns out, training a VDM can be done using three different interpretations:

- Predicting original image  (we just did this)
- Predicting noise  (coming up!)
- Predicting score function  (coming up!)

# VDM as a Noise Predictor

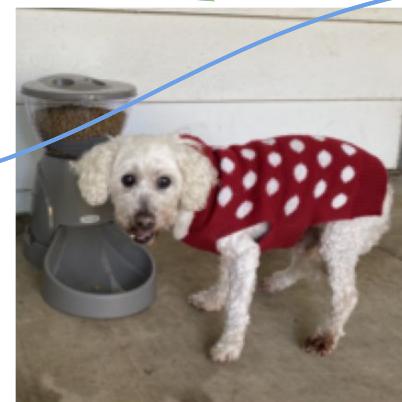
Recall that our objective is to predict  $\hat{x}_{\theta}(x_t, t) \approx x_0$

Image  and Noise ? They are the same!

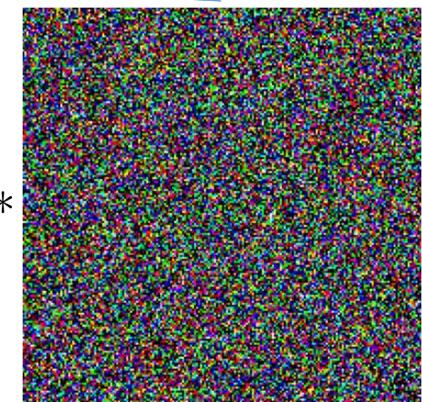
What does it mean intuitively?

For arbitrary  $x_t \sim q(x_t | x_0)$ , we can rewrite it as  $x_t = x_0 + \alpha_t \epsilon_0$

Predicting  $x_0$  determines  $\epsilon_0$  and vice-versa, since they sum to the same thing!



$\alpha_t *$



$x_t \sim q(x_t | x_0)$

$x_0$

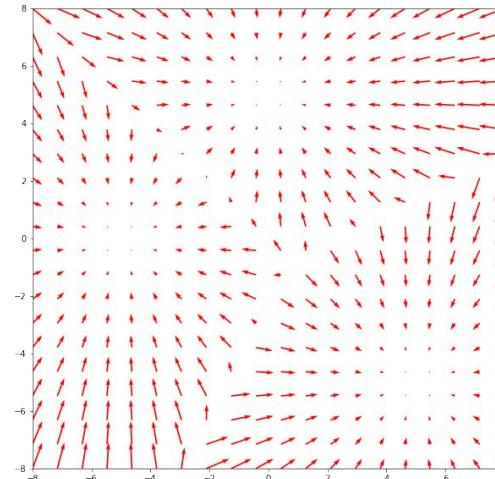
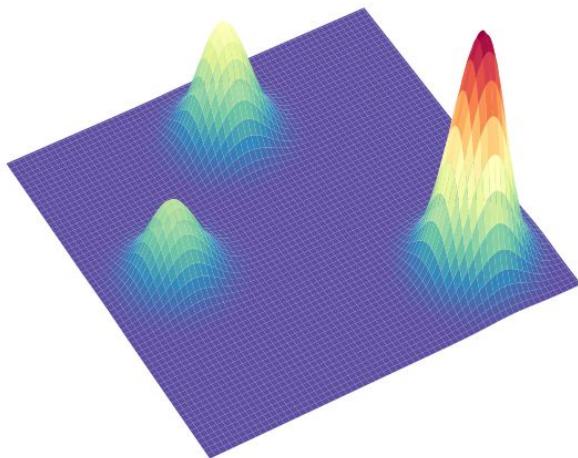
$\epsilon_0$

# Score Functions 100

What are score functions?

$$\nabla_{\mathbf{x}} \log p(\mathbf{x})$$

Intuitively, they describe how to move in data space to improve the (log) likelihood.



# Tweedie's Formula

Mathematically, for a Gaussian variable  $\mathbf{z} \sim \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_z, \boldsymbol{\Sigma}_z)$  Tweedie's formula states:

$$\mathbb{E} [\boldsymbol{\mu}_z \mid \mathbf{z}] = \mathbf{z} + \boldsymbol{\Sigma}_z \nabla_{\mathbf{z}} \log p(\mathbf{z})$$

Then, since we have previously shown that:

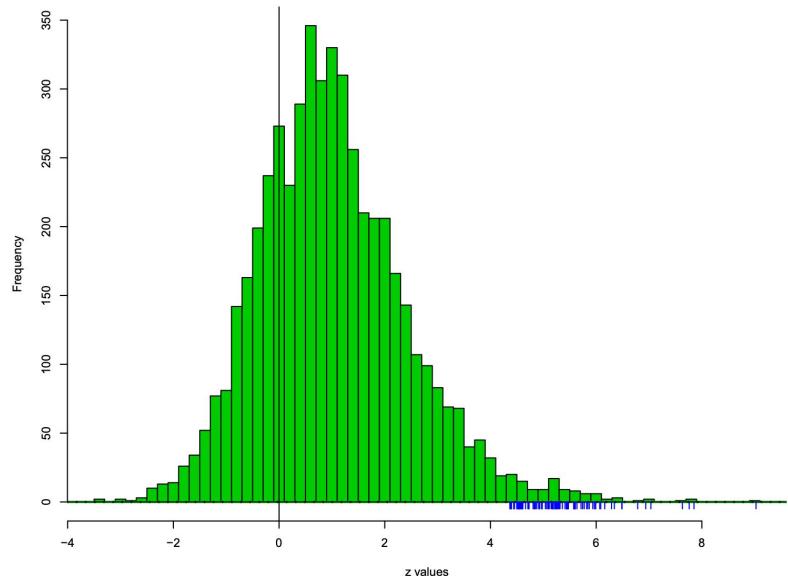
$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \mathbf{x}_0, \alpha_t^2 \mathbf{I})$$

By Tweedie's Formula, we derive:

$$\mathbb{E} [\boldsymbol{\mu}_{\mathbf{x}_t} \mid \mathbf{x}_t] = \mathbf{x}_t + \alpha_t^2 \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t)$$

The best estimate for the true mean is  $\boldsymbol{\mu}_{\mathbf{x}_t} = \mathbf{x}_0$

$$\mathbf{x}_0 = \mathbf{x}_t + \alpha_t^2 \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t)$$



# Tweedie's Formula

There exists a mathematical formula that states that:

$$\boldsymbol{x}_0 \approx \boldsymbol{x}_t + \alpha_t^2 \nabla_{\boldsymbol{x}_t} \log p(\boldsymbol{x}_t)$$

Due to the fact that the distribution is Gaussian:

$$q(\boldsymbol{x}_t | \boldsymbol{x}_0) = \mathcal{N}(\boldsymbol{x}_t | \boldsymbol{x}_0, \alpha_t^2 \mathbf{I})$$



# VDM as a Score Predictor 100

Recall that our objective is to predict  $\hat{x}_{\theta}(x_t, t) \approx x_0$

## Score and Noise ?

There is a relationship between the score and the noise, which we can derive by equating Tweedie's formula with the Reparameterization Trick.

$$\begin{aligned}\mathbf{x}_0 &= \mathbf{x}_t + \alpha_t^2 \nabla \log p(\mathbf{x}_t) = \mathbf{x}_t - \alpha_t \boldsymbol{\epsilon}_0 \\ \therefore \alpha_t^2 \nabla \log p(\mathbf{x}_t) &= -\alpha_t \boldsymbol{\epsilon}_0 \\ \nabla \log p(\mathbf{x}_t) &= -\frac{1}{\alpha_t} \boldsymbol{\epsilon}_0\end{aligned}$$

Intuitively, the direction to move in data space towards a natural image is the negative noise term that was added.

# Three Different Interpretations

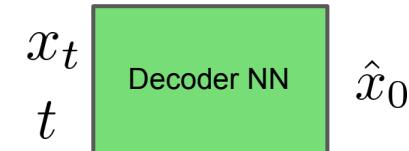
It turns out, training a VDM can be implemented as a neural net that:

-  Predicts original image  $\hat{x}_\theta(x_t, t) \approx x_0$
-  Predicts noise epsilon  $\hat{\epsilon}_\theta(x_t, t) \approx \epsilon_0$   
? ? ?
-  Predicts score function  $s_\theta(x_t, t) \approx \nabla_{x_t} \log p(x_t)$

# A Summary

We have learned that a diffusion model is simply one neural network that predicts a clean image from a noisy image.

Objective:  $\arg \min_{\theta} \|x_0 - \hat{x}_{\theta}(x_t, t)\|^2$



Sampling:

